

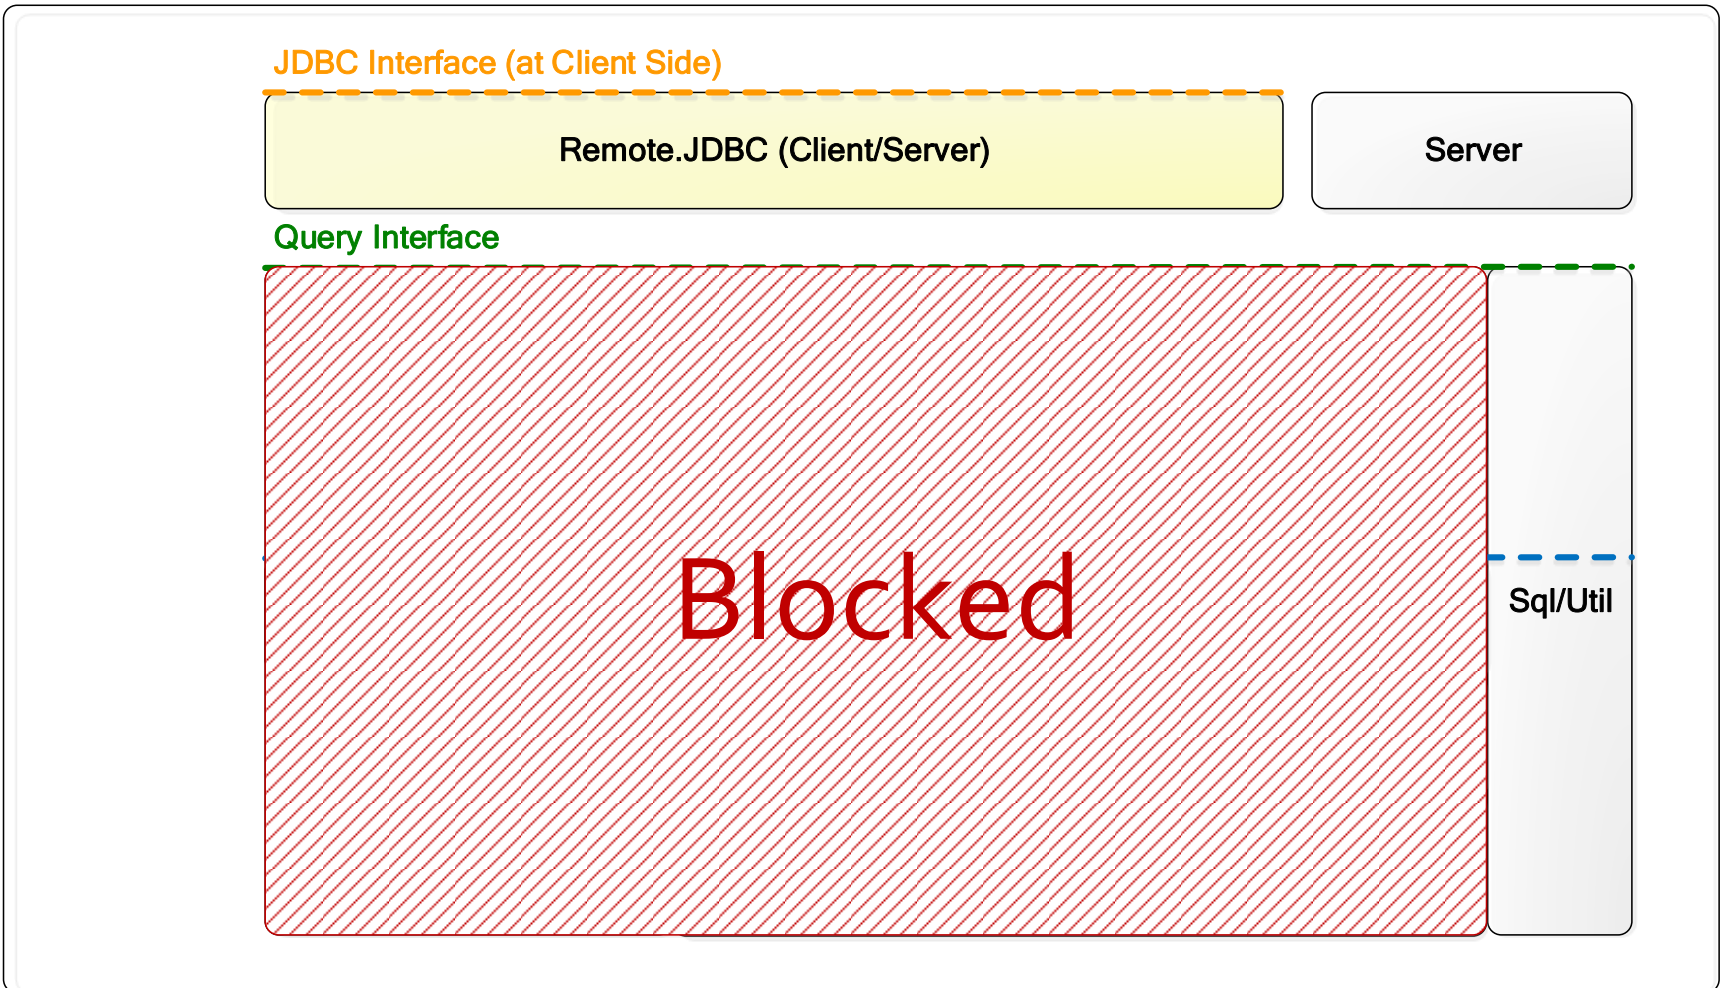
VanillaCore Walkthrough

Part 2

Database systems
DataLab, CS, NTHU
Spring 2018

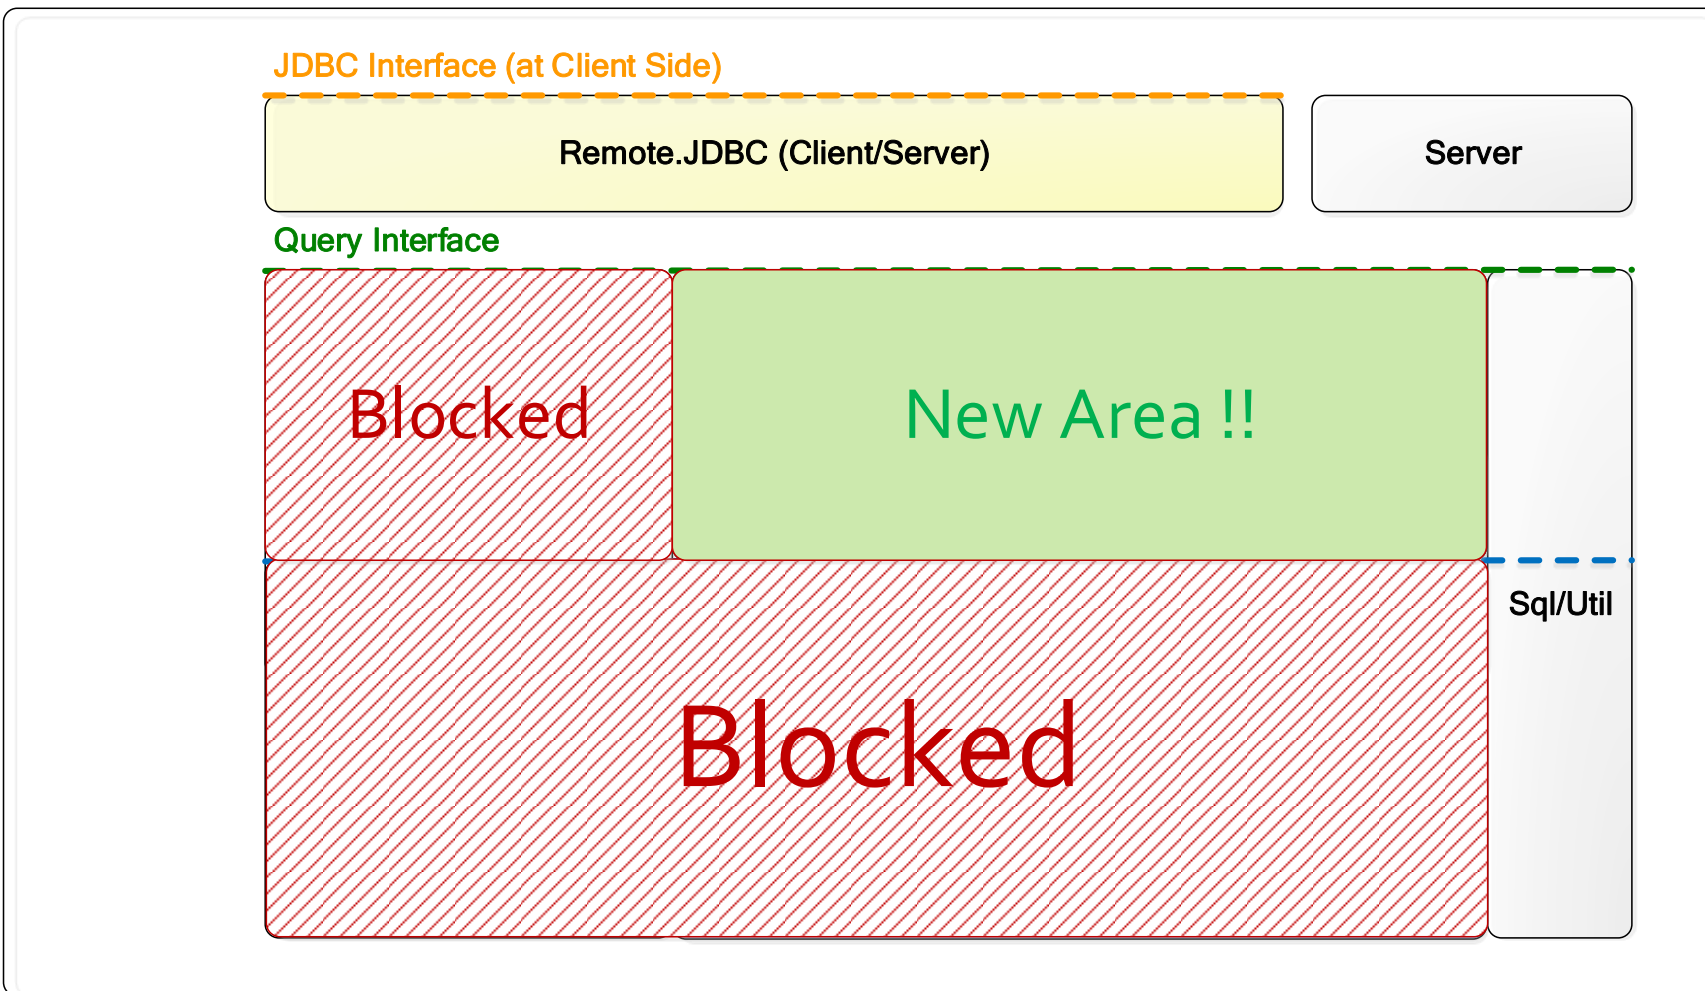
Last Time

VanillaDB



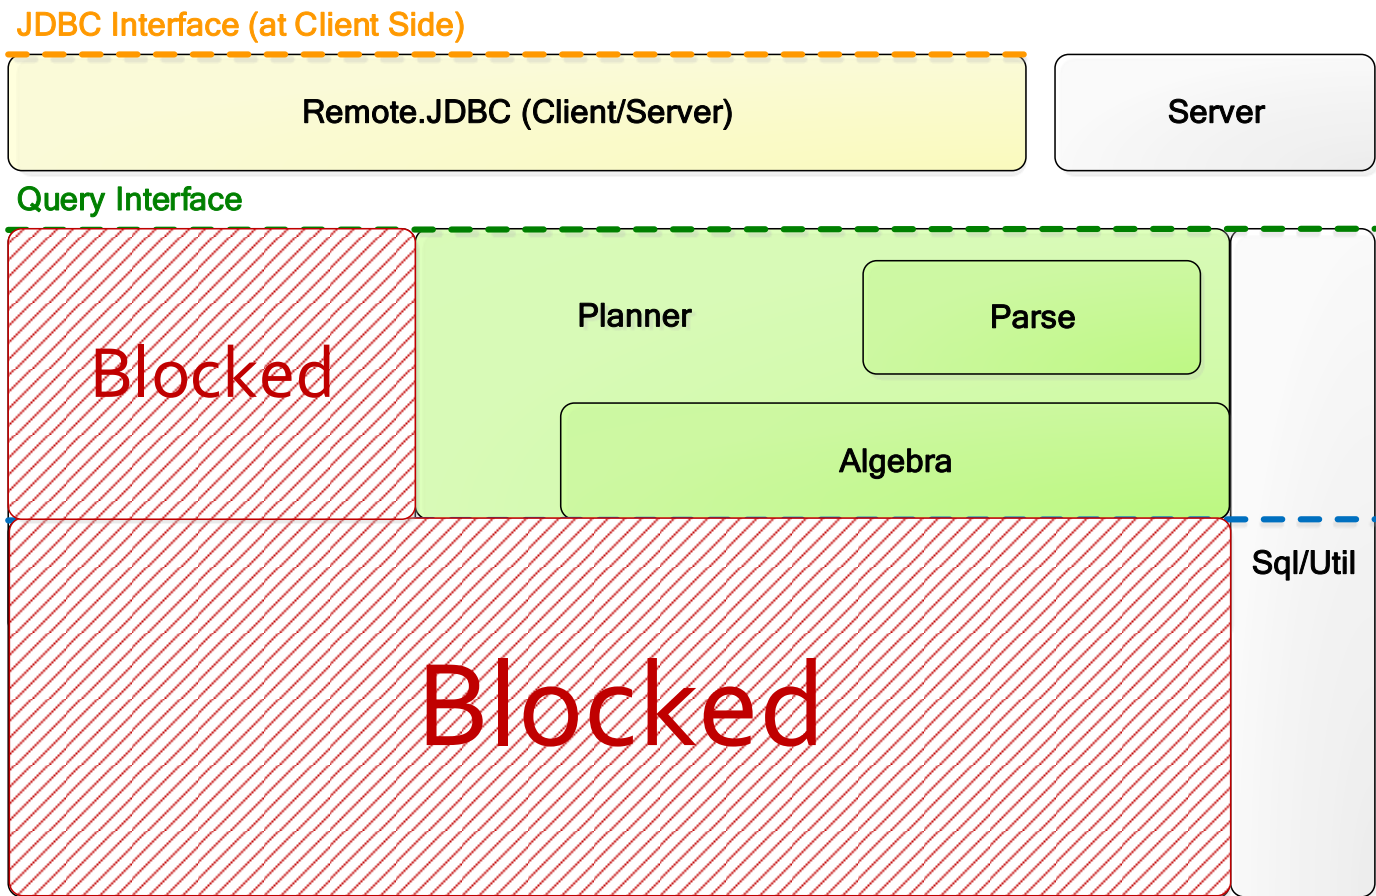
This Time

VanillaDB



This Time

VanillaDB



Outline

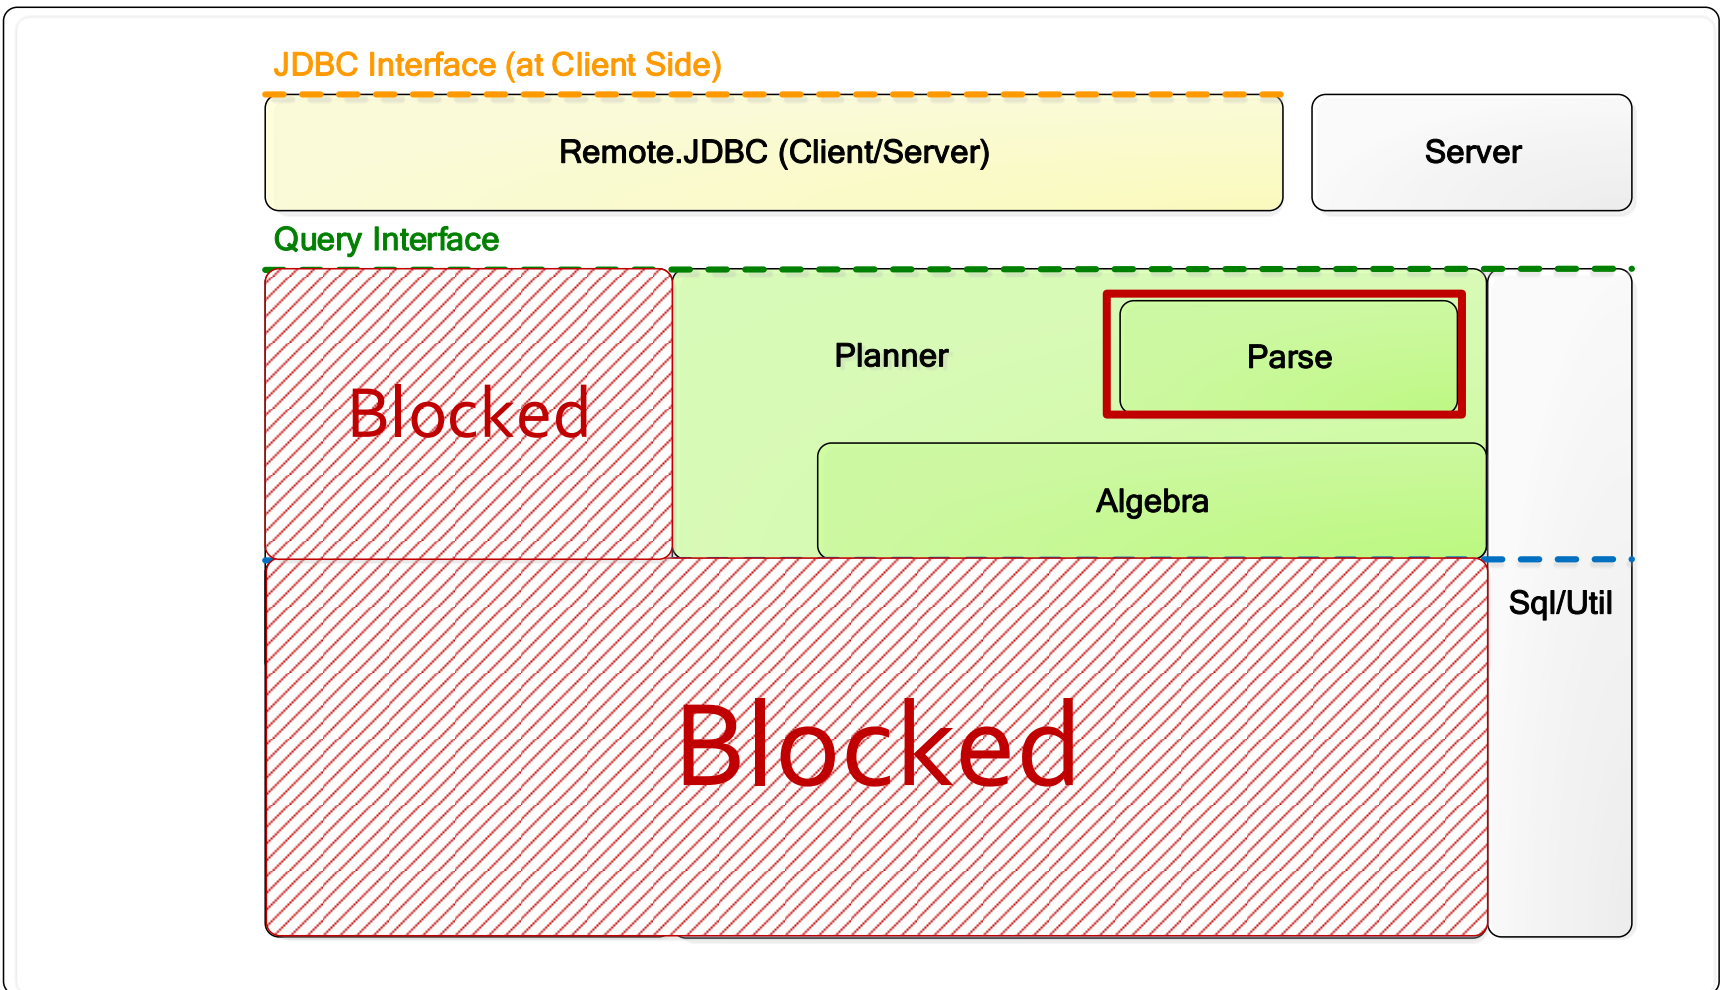
- Query package
 - Parse package
 - Algebra package
 - Planner package

Outline

- Query package
 - Parse package
 - Algebra package
 - Planner package

Where Are We ?

VanillaDB

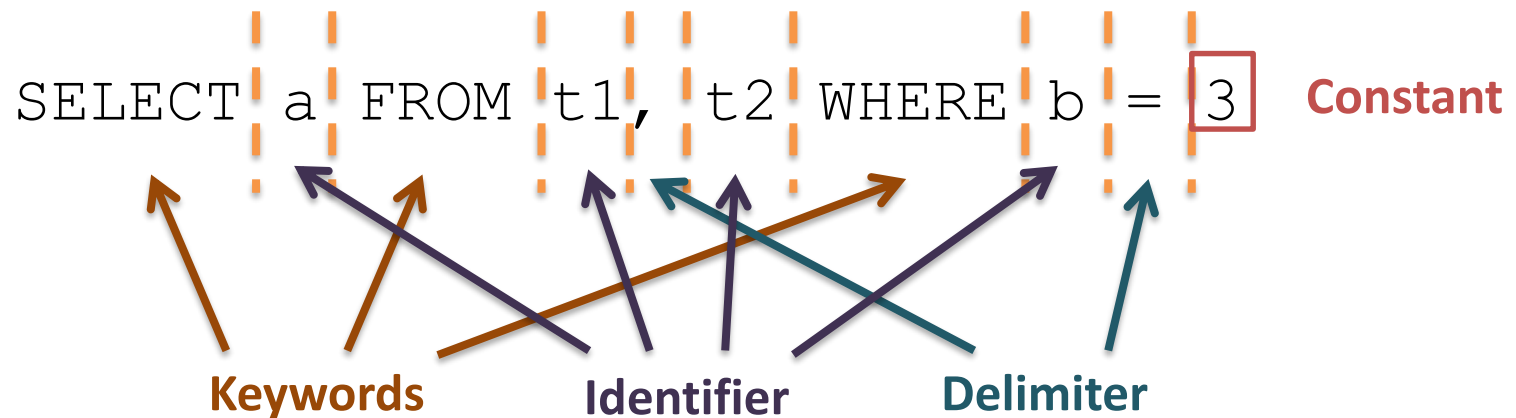


Two Steps of Parsing A SQL

- Lexer
 - Tokenizing
 - Identifying keywords, IDs, values, delimiters
- Parser
 - Checking syntax
 - Identifying the action and the parameters

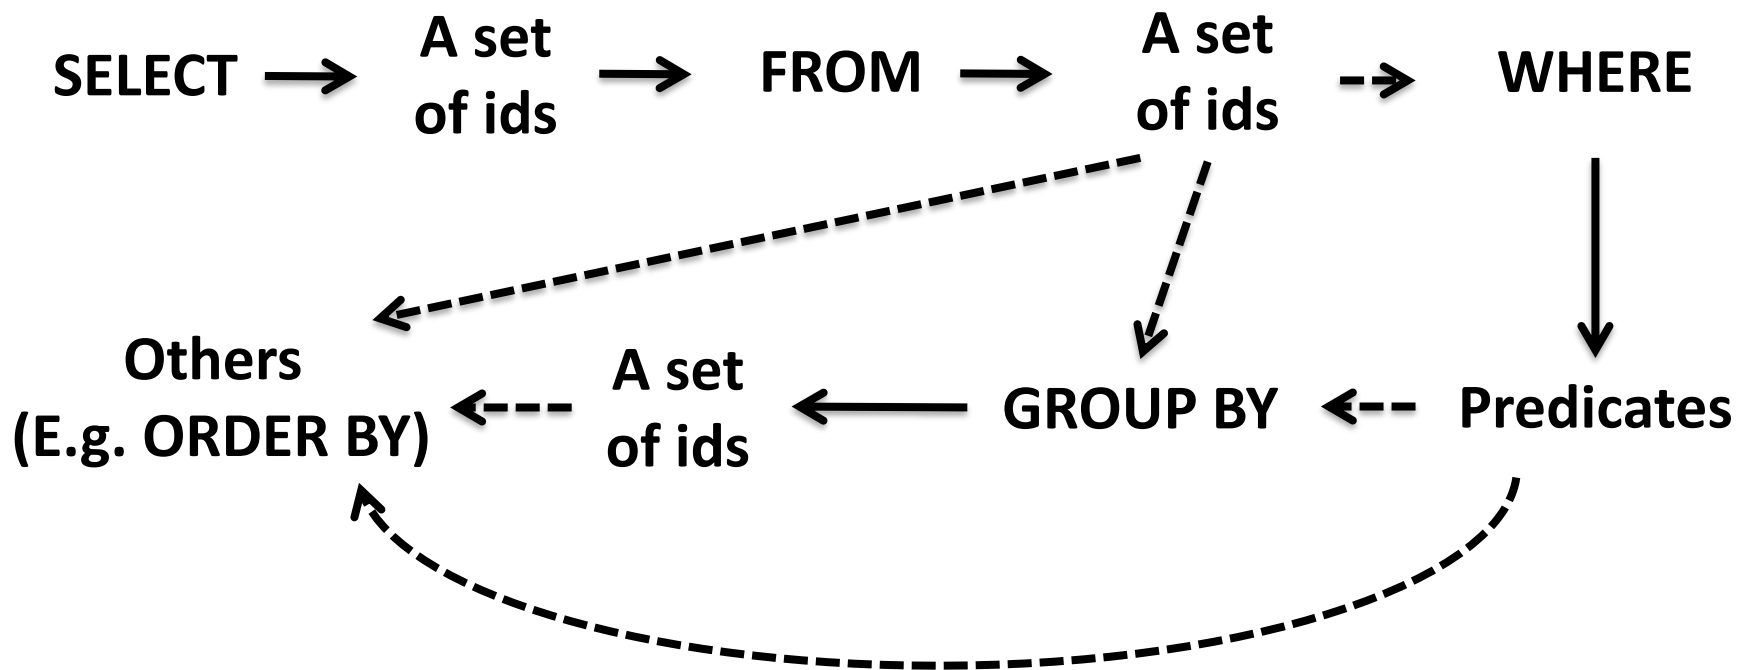
How VanillaCore Parses A SQL (1/2)

Lexer (Lexical Analyzer)



How VanillaCore Parses A SQL (2/2)

Parser

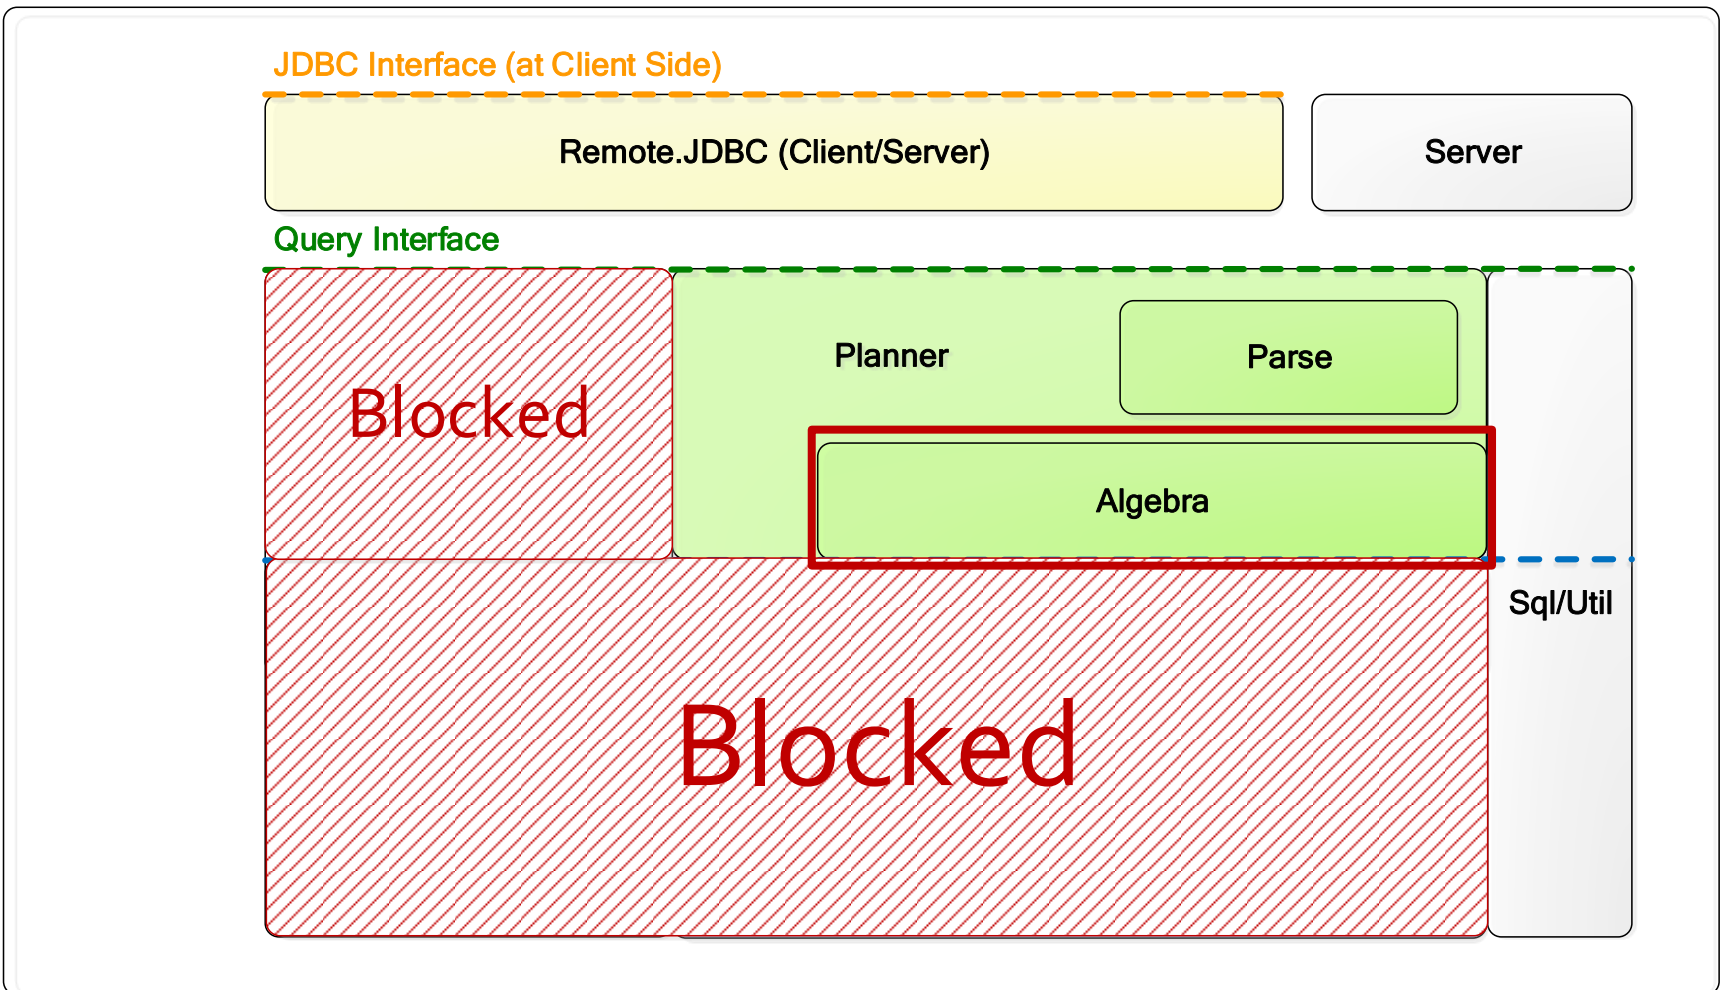


Outline

- Query package
 - Parse package
 - Algebra package
 - Planner package

Where Are We ?

VanillaDB



algebra Package

Plan
Classes

Scan
Classes

Index
Package

Materialize
Package

Multi-buffer
Package

Plan & Scan

<<interface>> Plan
+ open() : Scan + blocksAccessed() : long + schema() : Schema + histogram() : Histogram + recordsOutput() : long

<<interface>> Scan
+ beforeFirst() + next() : boolean + close() + hasField(fldname : String) : boolean

Using a Query Plan


```
VanillaDb.init("studentdb");  
Transaction tx = VanillaDb.txMgr().newTransaction(  
    Connection.TRANSACTION_SERIALIZABLE, true);
```

```
Plan pb = new TablePlan("b", tx);  
Plan pu = new TablePlan("u", tx);  
Plan pp = new ProductPlan(pb, pu);  
Predicate pred = new Predicate("...");  
Plan sp = new SelectPlan(pp, pred);
```

```
sp.blockAccessed(); // estimate #blocks accessed
```

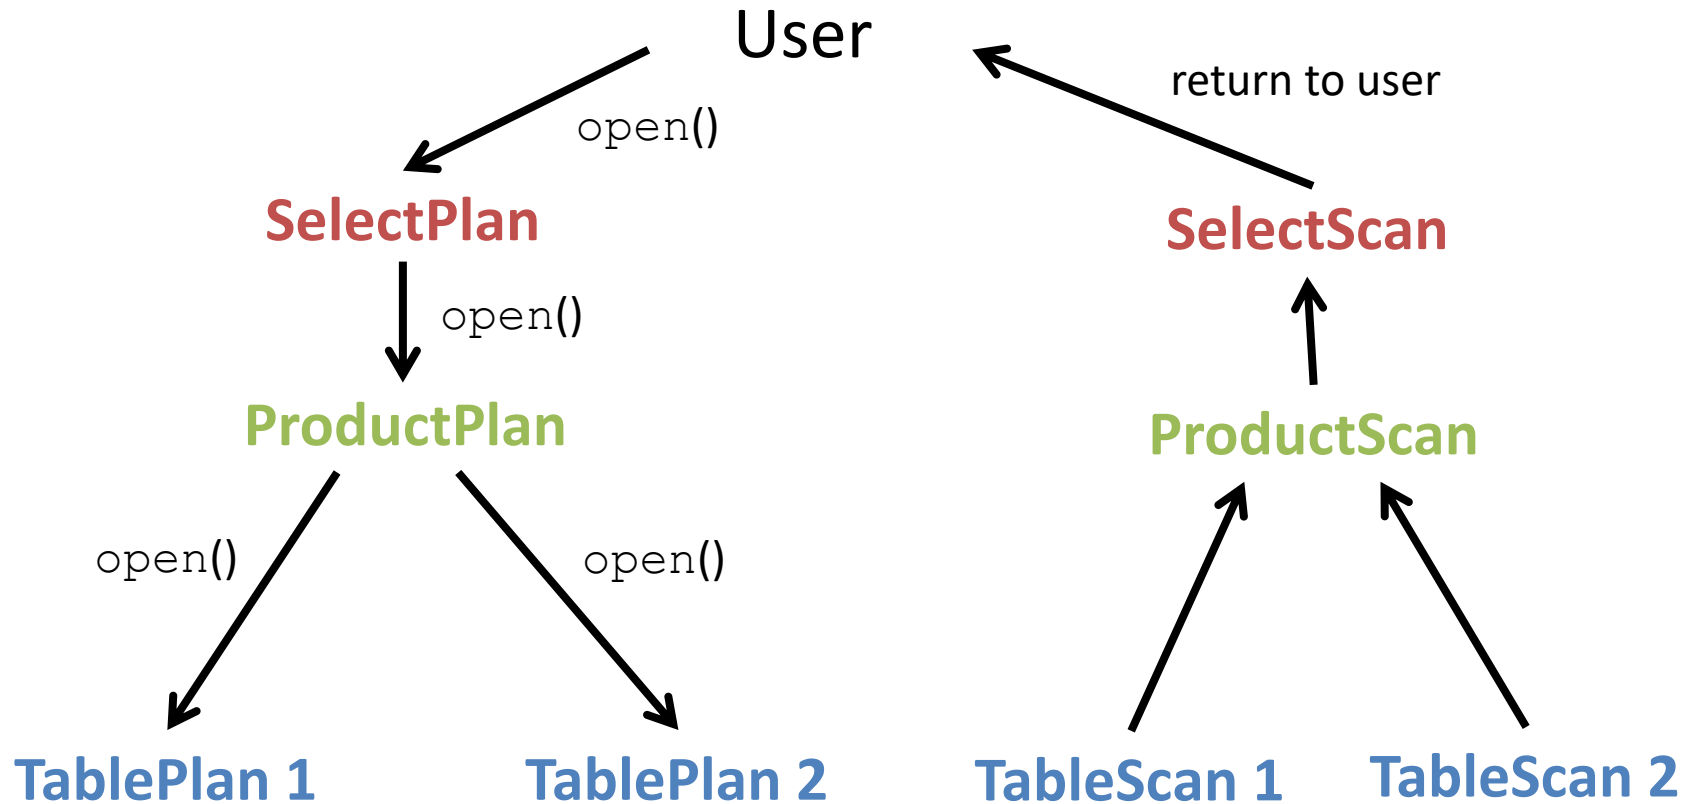
```
// open corresponding scan only if sp has low cost  
Scan s = sp.open();  
s.beforeFirst();  
while (s.next())  
    s.getVal("bid");  
s.close();
```

select(p, where...)



What Happened When We Called `open ()` ?

open ()



How Do Scans Work ?

Example

project (s)



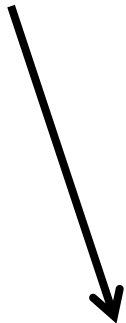
beforeFirst()

select (p)




beforeFirst()

product (b, u)



beforeFirst()



blog_id	url	created	author_id
33981	...	2009/10/31	728
33982	...	2012/11/15	729
41770	...	2012/10/20	730

u

user_id	name	balance
729	Brandon	10,235
730	Picachu	NULL

```
SELECT blog_id FROM b, u
WHERE name = "Brandon"
AND author_id = user_id;
```

Example

project (s)



beforeFirst()

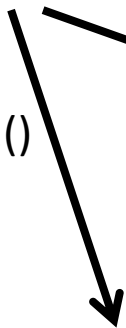
select (p)



beforeFirst()

product (b, u)

next()



beforeFirst()



b


blog_id	url	created	author_id
33981	...	2009/10/31	728
33982	...	2012/11/15	729
41770	...	2012/10/20	730

user_id	name	balance
729	Brandon	10,235
730	Picachu	NULL


```
SELECT blog_id FROM b, u
WHERE name = "Brandon"
AND author_id = user_id;
```

Example

project (s)



select (p)



```
SELECT blog_id FROM b, u
WHERE name = "Brandon"
AND author_id = user_id;
```

product (b, u)

blog_id	url	created	author_id	user_id	name	balance
33981	...	2009/10/31	728	729	Brandon	10,235

next()

b


blog_id	url	created	author_id
33981	...	2009/10/31	728
33982	...	2012/11/15	729
41770	...	2012/10/20	730

u


user_id	name	balance
729	Brandon	10,235
730	Picachu	NULL

Example

project (s)



select (p)



product (b, u)

```
SELECT blog_id FROM b, u
WHERE name = "Brandon"
AND author_id = user_id;
```

blog_id	url	created	author_id	user_id	name	balance
33981	...	2009/10/31	728	730	Picachu	NULL

next()

b



blog_id	url	created	author_id
33981	...	2009/10/31	728
33982	...	2012/11/15	729
41770	...	2012/10/20	730


u




user_id	name	balance
729	Brandon	10,235
730	Picachu	NULL

Example

project (s)




select (p)



```
SELECT blog_id FROM b, u
WHERE name = "Brandon"
AND author_id = user_id;
```

product (b, u)



b

blog_id	url	created	author_id
33981	...	2009/10/31	728
33982	...	2012/11/15	729
41770	...	2012/10/20	730

u

user_id	name	balance
729	Brandon	10,235
730	Picachu	NULL

Example


blog_id
33982

SELECT blog_id FROM b, u

blog_id	url	created	author_id	user_id	name	balance
33982	...	2012/11/15	730	730	Picachu	NULL



blog_id	url	created	author_id	user_id	name	balance
33982	...	2012/11/15	729	729	Brandon	10,235



next()

b

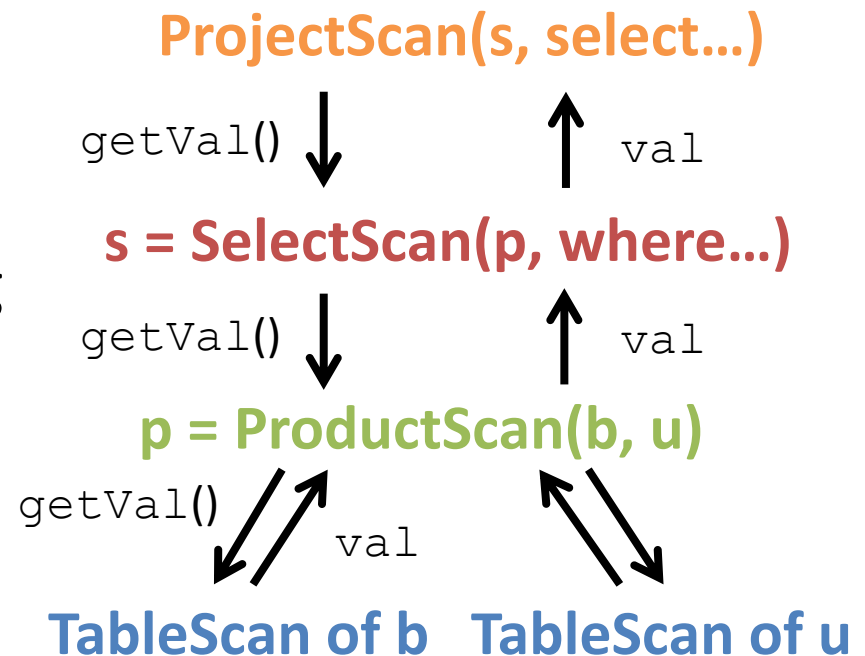
blog_id	url	created	author_id
33981	...	2009/10/31	728
33982	...	2012/11/15	729
41770	...	2012/10/20	730

u

user_id	name	balance
729	Brandon	10,235
730	Picachu	NULL

Pipelined Scanning

- The above operators implement **pipelined scanning**
 - Calling a method of a node results in recursively calling the same methods of child nodes on-the-fly
 - Records are computed one at a time as needed --- no intermediate records are saved

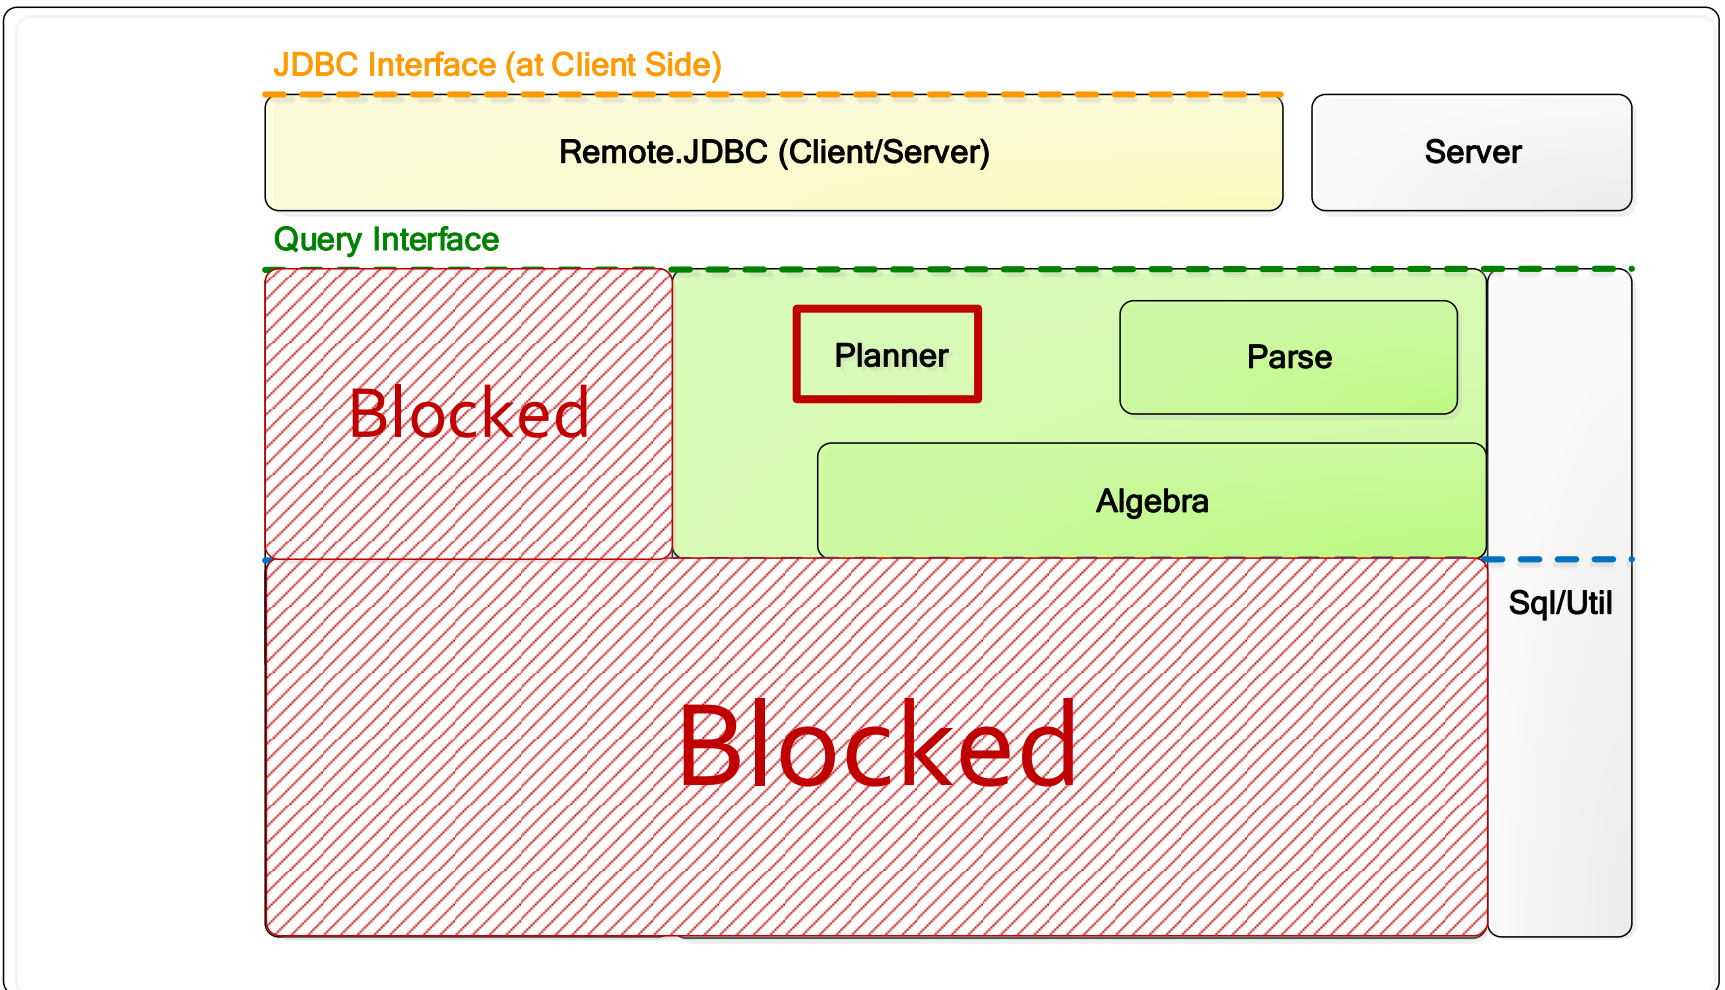


Outline

- Query package
 - Parse package
 - Algebra package
 - Planner package

Where Are We ?


VanillaDB



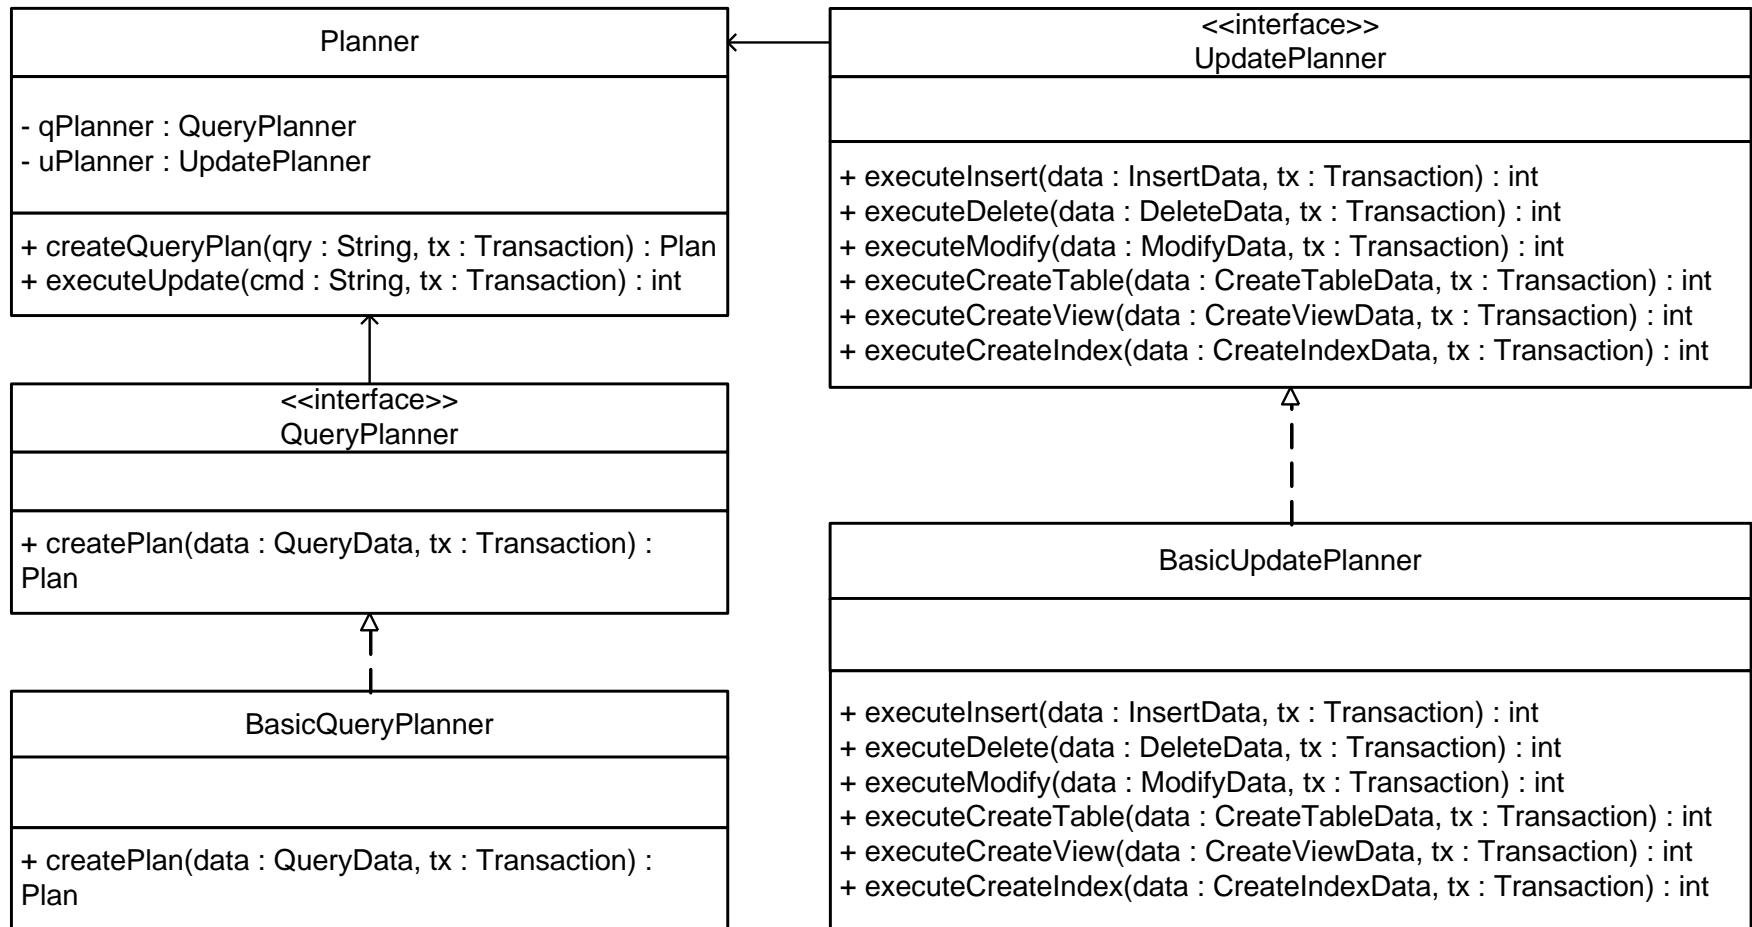
Planner

- The one puts all these together
 1. Accepts a query
 2. Creates a parser to parse the query
 3. Verifies all parameters are reasonable
 4. Generates some plan trees according to the query
 5. Choose the cheapest one

planner Package



Basic Implementation



Advanced Implementation

