

VanillaCore Walkthrough

Part 3

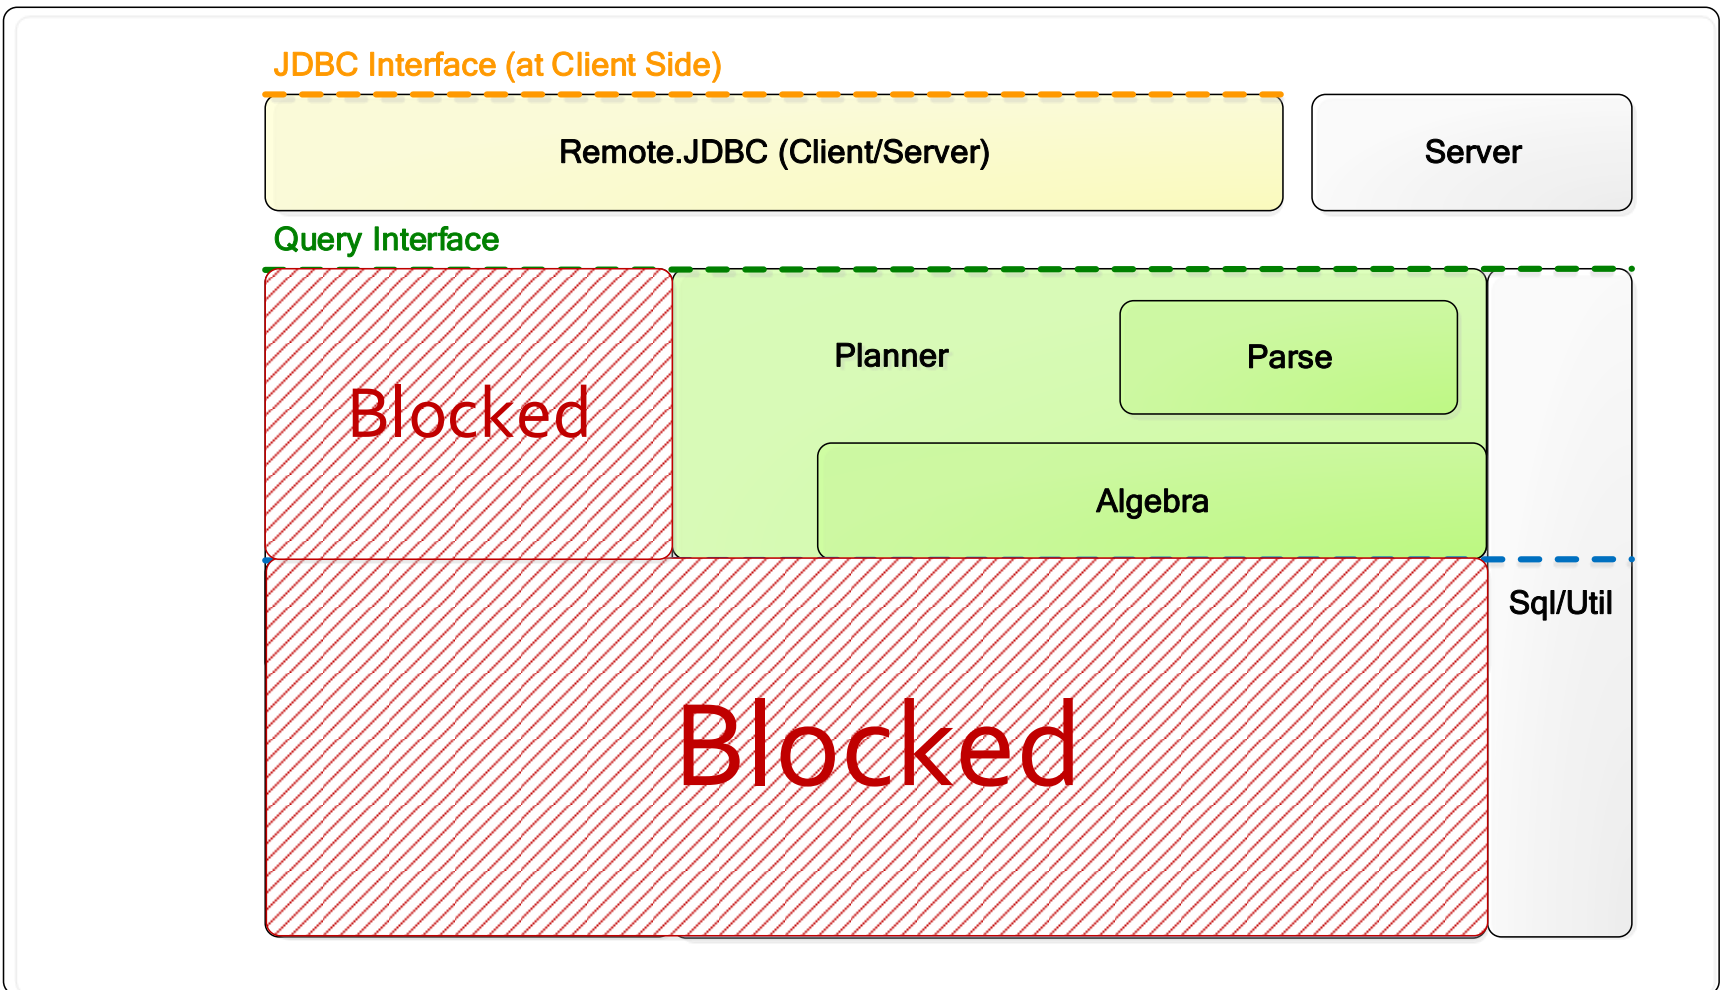
Cloud Databases

DataLab

CS, NTHU

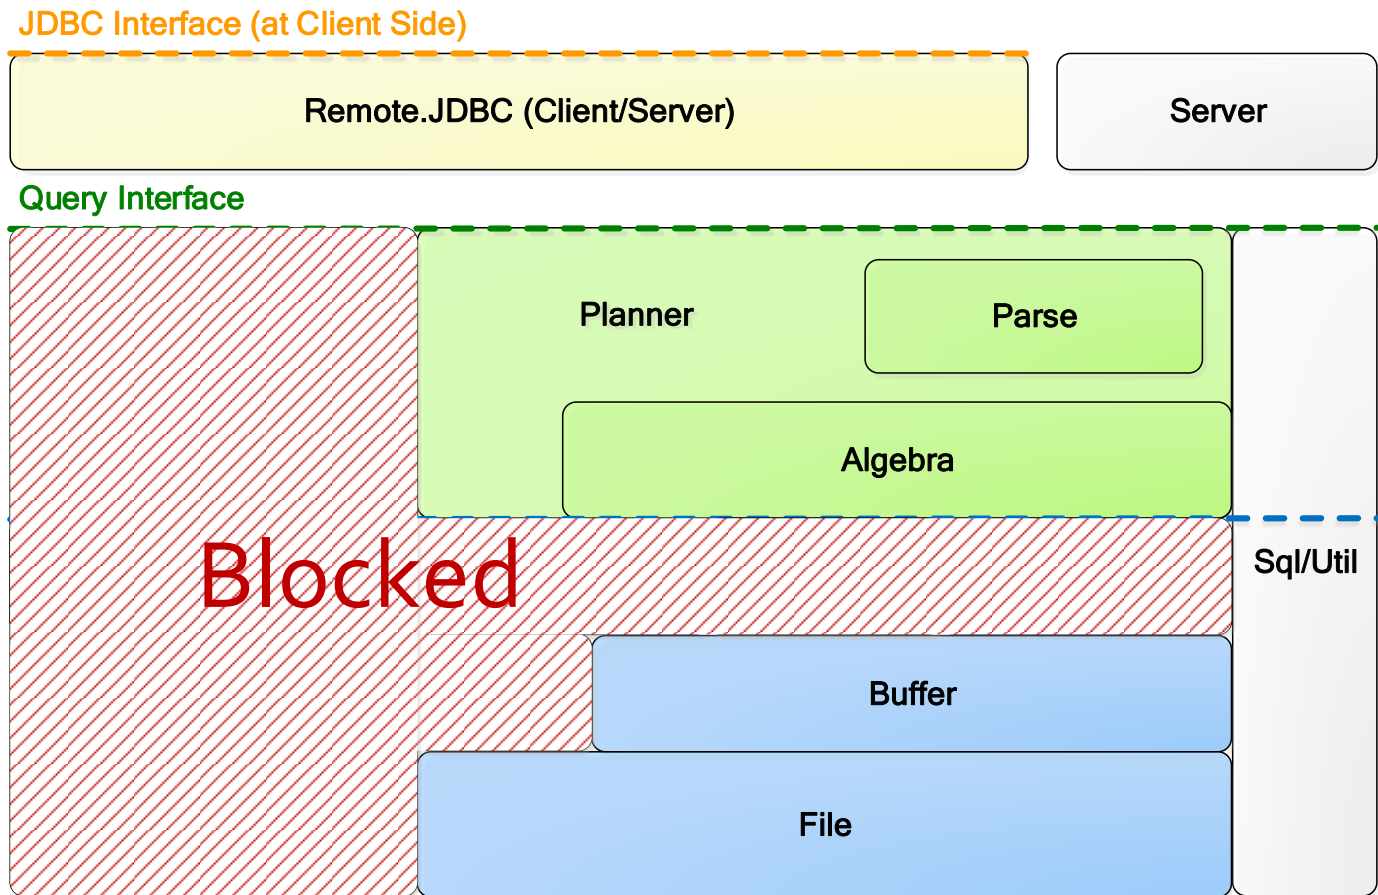
Last Time

VanillaDB



This Time

VanillaDB



Outline


- File package
- Buffer package

Outline

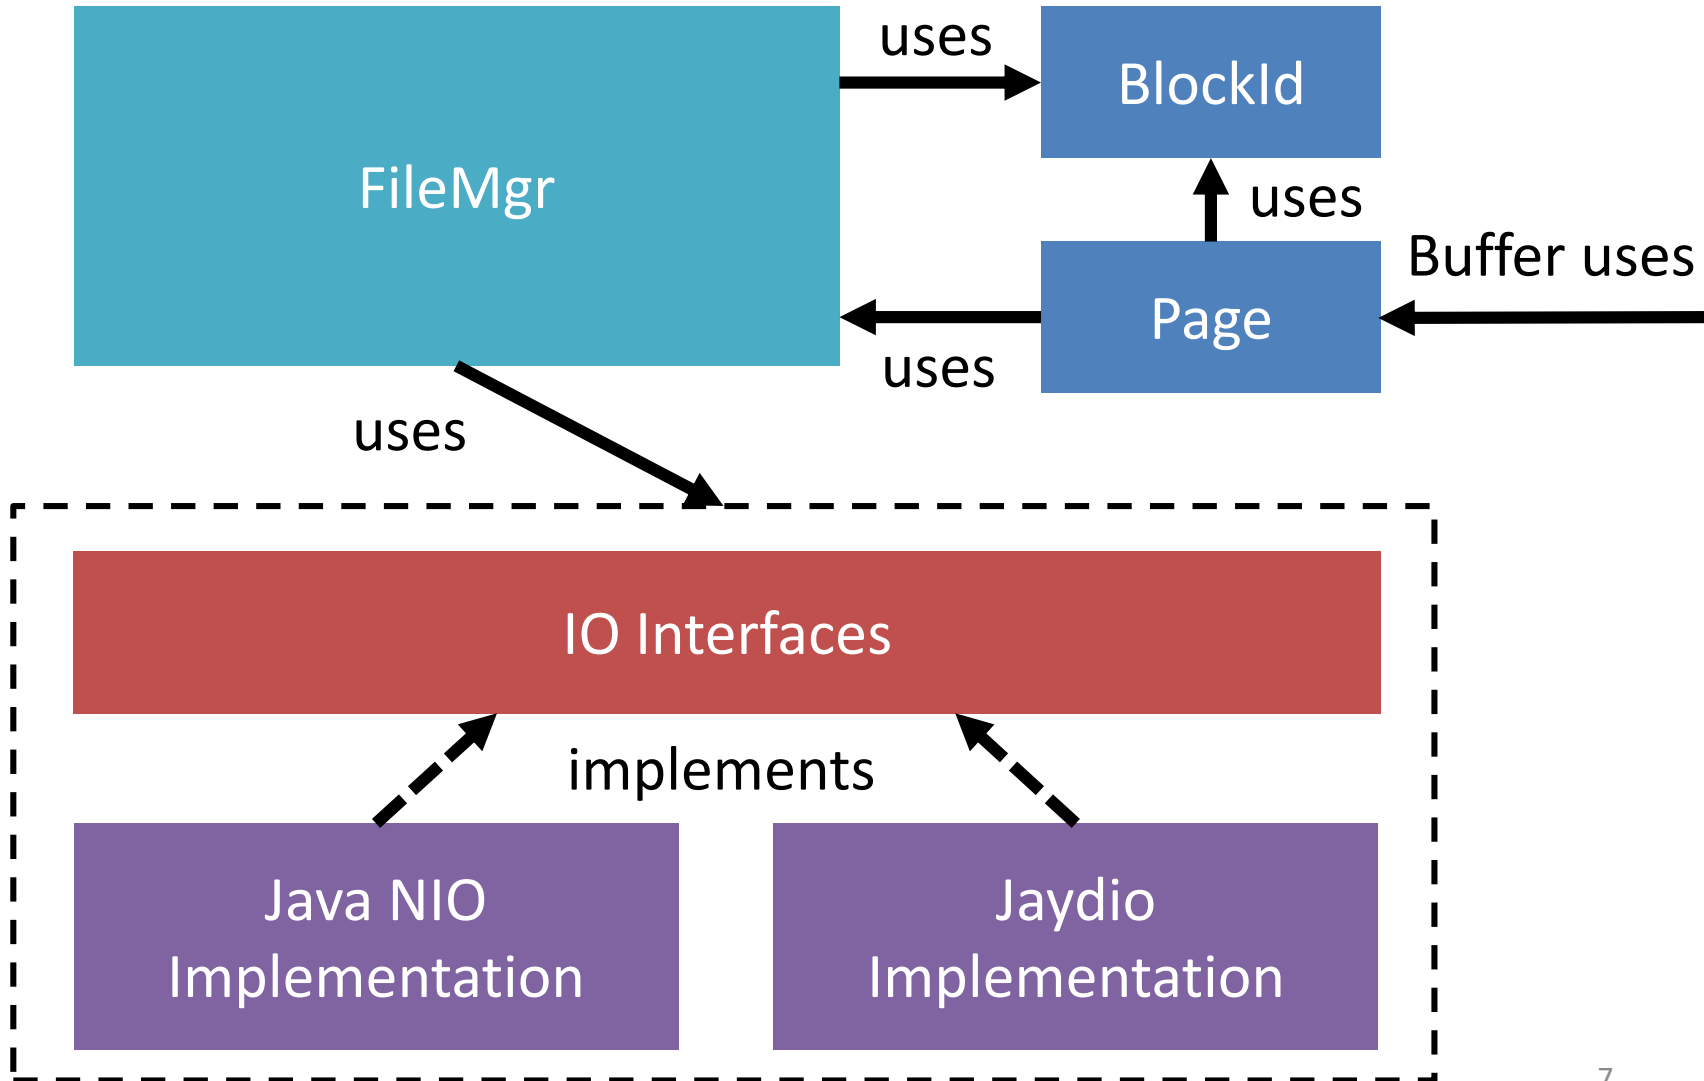
- File package
- Buffer package

Where are we?

VanillaDB



file Package



BlockId

```
public class BlockId {  
    private String fileName;  
    private long blkNum;  
  
    public BlockId(String fileName, long blkNum) {  
        this.fileName = fileName;  
        this.blkNum = blkNum;  
    }  
  
    public String fileName() {  
        return fileName;  
    }  
  
    public long number() {  
        return blkNum;  
    }  
    ...  
}
```

| BlockId |
|--|
| |
| + BlockId(filename : String, blknum : long) + fileName() : String + number() : long + equals(Object : obj) : boolean + toString() : String + hashCode() : int |

Page

| Page |
|--|
| <u><<final>> + BLOCK_SIZE : int</u> |
| <u>+ maxSize(type : Type) : int</u> <u>+ size(val : Constant) : int</u> + Page() <<synchronized>> + read(blk : BlockId) <<synchronized>> + write(blk : BlockId) <<synchronized>> + append(filename : String) : BlockId <<synchronized>> + getVal(offset : int, type : Type) : Constant <<synchronized>> + setVal(offset : int, val : Constant) + close() |

Page

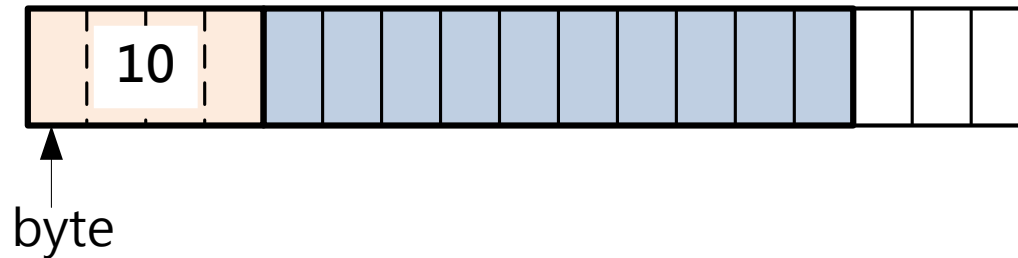
- Backed by `IoBuffer`

```
private IoBuffer contents = IoAllocator.newIoBuffer(BLOCK_SIZE);
```

- Translate constants using `Constant.asBytes()`
 - Fixed length for numeric type constants (e.g., 4 bytes for `IntegerConstant`)
 - Variable length for `VarcharConstant`
- How to reconstruct a varchar constant in getter?

Storing A Varchar

- Page stores a Varchar in two parts
 - The first is the length of those bytes
 - The second is the bytes from `asByte()`




setVal

```
public synchronized void setVal(int offset, Constant val) {
    byte[] byteval = val.asBytes();

    // Append the size of value if it is not fixed size
    if (!val.getType().isFixedSize()) {
        // check the field capacity and value size
        if (offset + ByteHelper.INT_SIZE + byteval.length > BLOCK_SIZE)
            throw new BufferOverflowException();


        byte[] sizeBytes = ByteHelper.toBytes(byteval.length);
        contents.put(offset, sizeBytes);
        offset += sizeBytes.length;
    }

    // Put bytes
    contents.put(offset, byteval);
}
```



getVal

```
public synchronized Constant getVal(int offset, Type type) {  
    int size;  
    byte[] byteVal = null;  
  
    // Check the length of bytes  
    if (type.isFixedSize()) {  
        size = type.maxSize();  
    } else {  
        byteVal = new byte[ByteHelper.INT_SIZE];  
        contents.get(offset, byteVal);  
        size = ByteHelper.toInteger(byteVal);  
        offset += ByteHelper.INT_SIZE;  
    }  
  
    // Get bytes and translate it to Constant  
    byteVal = new byte[size];  
    contents.get(offset, byteVal);  
    return Constant.newInstance(type, byteVal);  
}
```



Sizing Information

- There are static APIs providing sizing information in Page

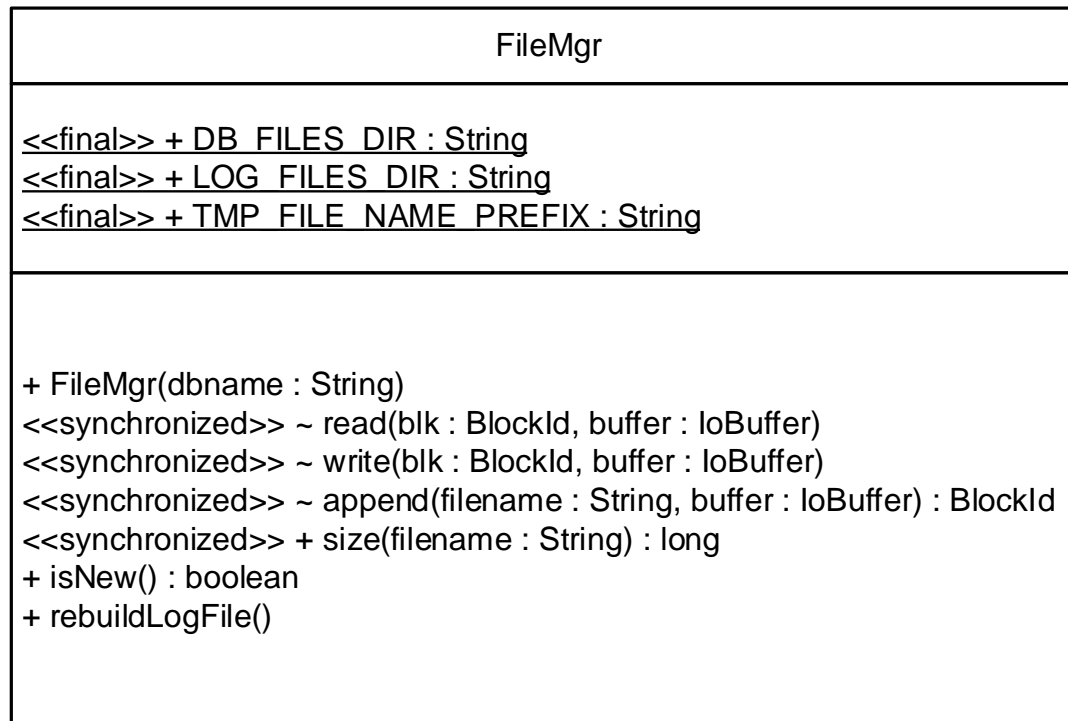
```
public static int maxSize(Type type) {  
    return type.isFixedSize() ? type.maxSize() : ByteHelper.INT_SIZE  
        + type.maxSize();  
}  
  
public static int size(Constant val) {  
    return val.getType().isFixedSize() ? val.size() : ByteHelper.INT_SIZE  
        + val.size();  
}
```

File I/Os

```
public Page() {  
}  
  
public synchronized void read(BlockId blk) {  
    fileMgr.read(blk, contents);  
}  
  
public synchronized void write(BlockId blk) {  
    fileMgr.write(blk, contents);  
}  
  
public synchronized BlockId append(String fileName) {  
    return fileMgr.append(fileName, contents);  
}
```

FileMgr

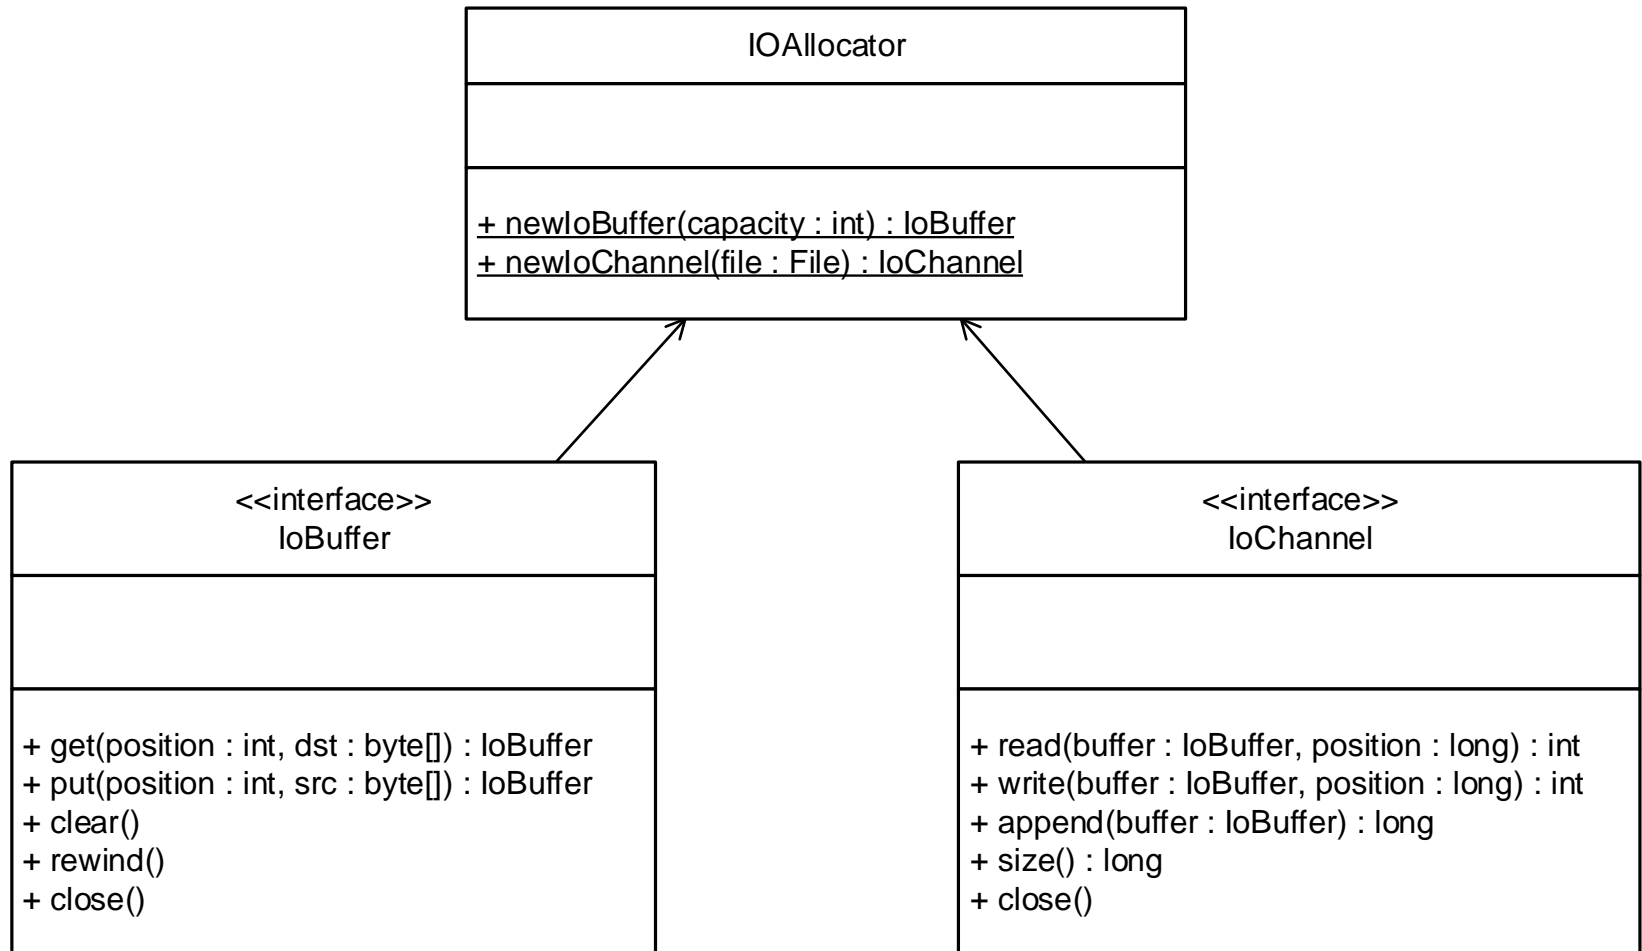
- Handles the actual I/Os
- Keeps the `IoChannel` instances of all opened files



FileMgr

- A page delegates read, write and, append to FileMgr
- Note that the file manager always reads/writes/appends a ***block-sized*** number of bytes from/to a file
 - Exactly one disk access per call

file.io



IoChannel in Java NIO

- Opens a file by creating a new `RandomAccessFile` instance and then obtain its file channel via `getChannel()`
- Files are open in “rws” mode when using Java NIO
 - The “rw” means that the file is open for reading and writing
 - The “s” means that the OS should not delay disk I/O in order to optimize disk performance; instead, every *write* operation must be written immediately to the disk

IoBuffer in Java NIO

- We don't want the memory space of `ByteBuffer` be swapped out by OS
- `ByteBuffer` has two factory methods: `allocate` and `allocateDirect`
 - `allocateDirect` tells JVM to use one of the OS's I/O buffers to hold the bytes
 - ***Not*** in Java programmable buffer, no garbage collection
 - Eliminates the redundancy of ***double buffering***

Outline

- File package
- Buffer package (TBA)