

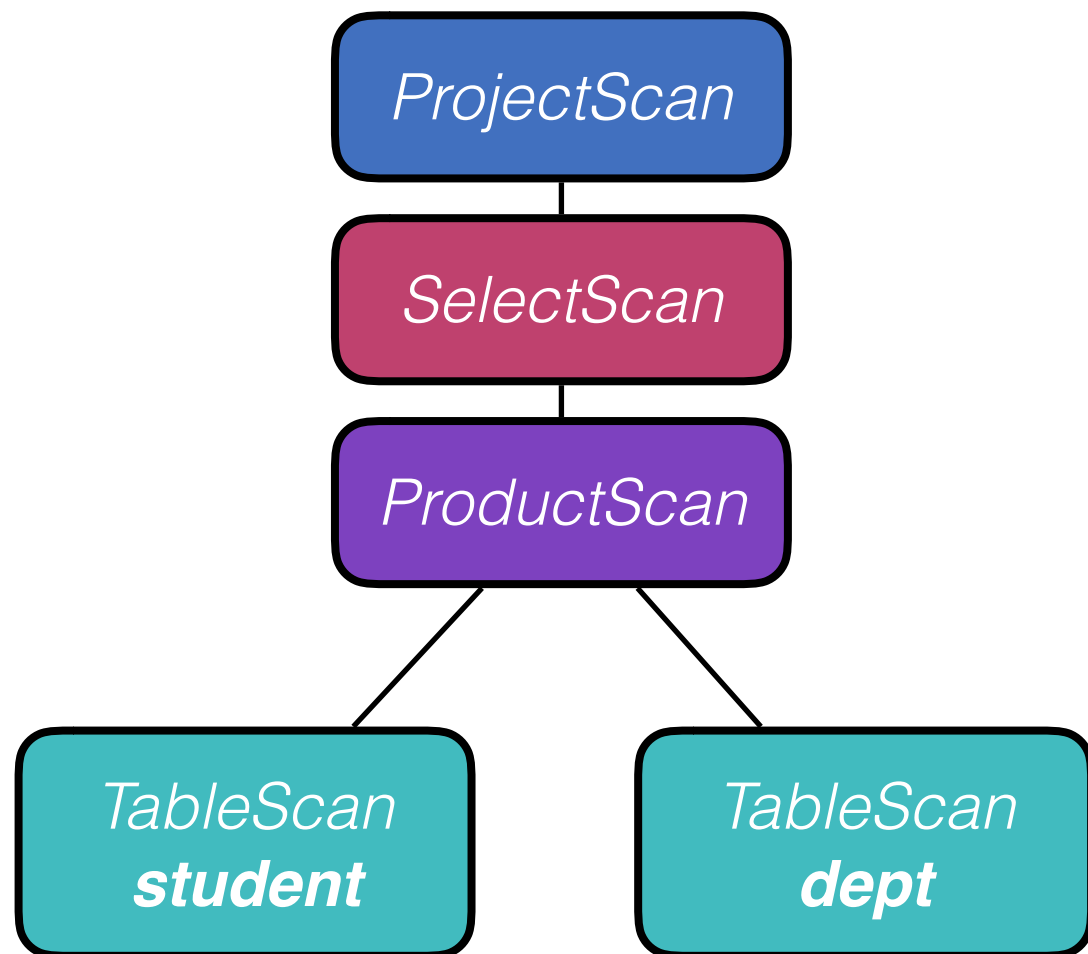
Midterm Project

ctsai@DataLAB
Cloud Database
2017 Spring

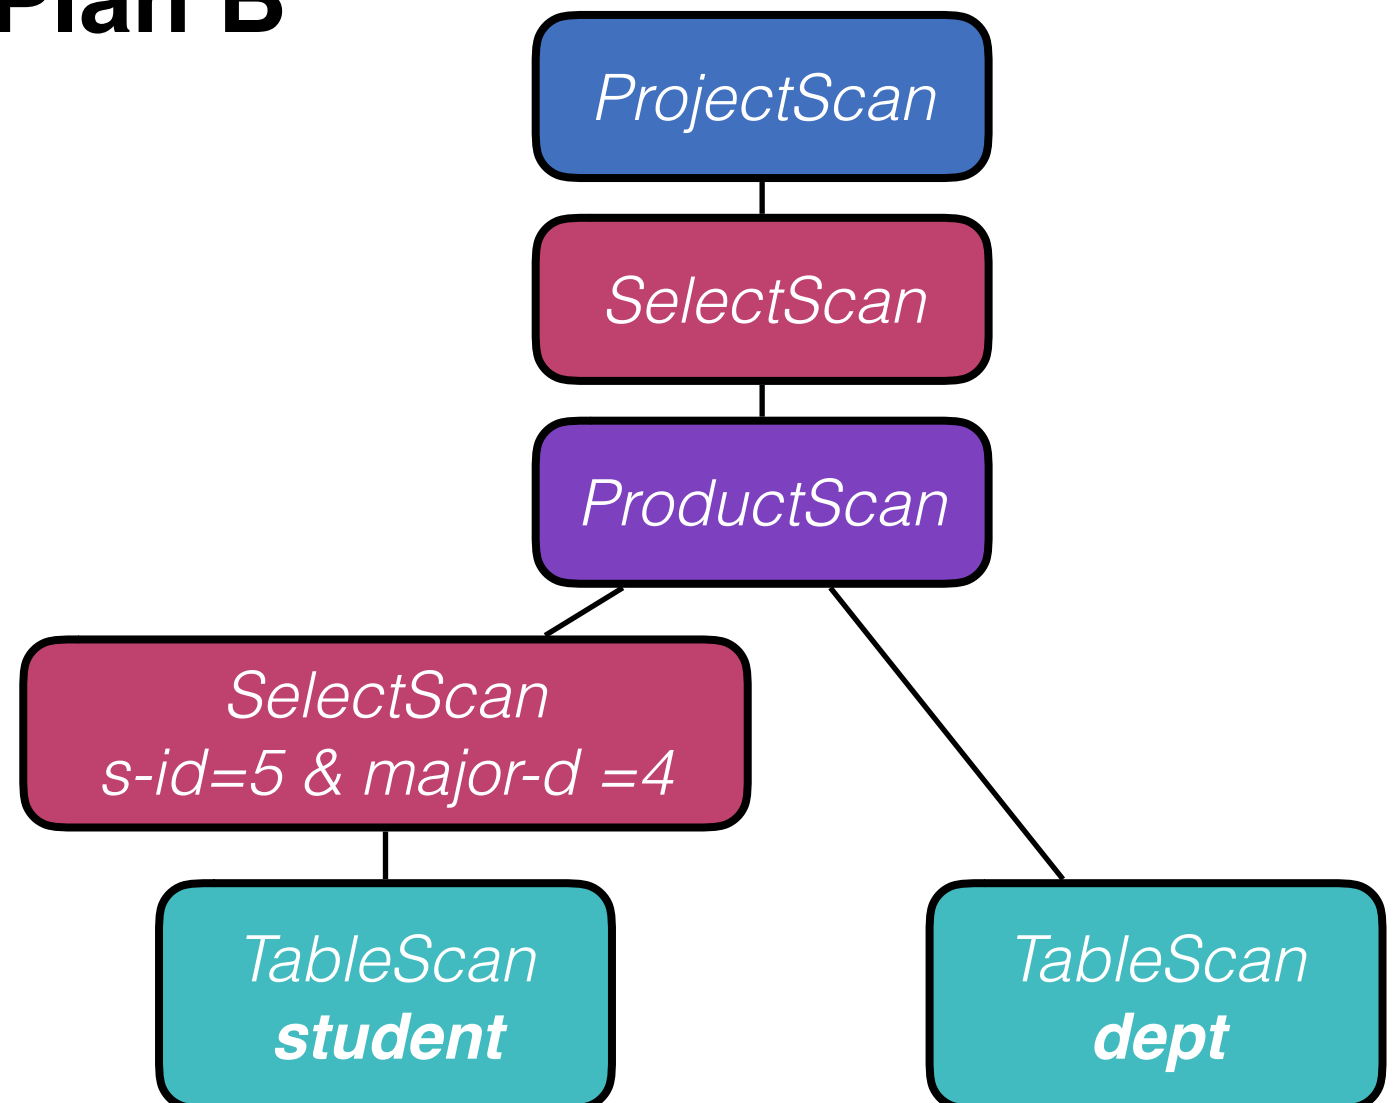
Revisit Query Processing

```
SELECT sname FROM student, dept
      WHERE major-id = d-id
            AND s-id = 5 AND major-id = 4;
```

Plan A



Plan B



Revisit Query Processing

- A good scan tree can be faster than a bad one for orders of magnitude
- Consider the product scan at middle
- Let $\text{Row}(\text{student})=10000$, $\text{Block}(\text{student})=1000$, $\text{Block}(\text{dept})=500$, and $\text{selectivity}(\text{s-id}=5 \& \text{major-id}=4)=0.01$
- Each block access requires 10ms
- Cost:
 - Plan A : $(10000 + 10000 * 500) * 10\text{ms} = 13.9 \text{ hours}$
 - Plan B : $(10000 + 10000 * 0.001 * 500) * 10\text{ms} = 8.4\text{min}$

Query Optimization

- Finding a good plan is crucial
- Cost difference between evaluation plans for a query can be enormous
 - E.g. seconds vs. days in some cases
- Midterm Project : Implement one of classic query optimizer
- What you will need ...
 - Query plan cost estimation
 - Transformation of query plans
 - Some heuristic optimization
 - Selinger optimizer <— you are going to implement it !!

How to say two Relational Expressions are equivalent ?

- Two relational algebra expressions are said to be **equivalent** if the two expressions generate the same set of tuples on every legal database instance
 - Note that the order of tuples is irrelevant
- An **equivalence rule** says that expressions of two forms are equivalent
- Some notations :


ProjectScan

Π_{L_1}

SelectScan

σ_{θ_1}

ProductScan



Equivalence Rules

1. Conjunctive selection operations can be deconstructed into a sequence of individual selections.

$$\sigma_{\theta_1 \wedge \theta_2}(E) = \sigma_{\theta_1}(\sigma_{\theta_2}(E))$$

2. Selection operations are commutative.

$$\sigma_{\theta_1}(\sigma_{\theta_2}(E)) = \sigma_{\theta_2}(\sigma_{\theta_1}(E))$$

3. Only the last in a sequence of projection operations is needed, the others can be omitted.

$$\Pi_{L_1}(\Pi_{L_2}(\dots(\Pi_{L_n}(E))\dots)) = \Pi_{L_1}(E)$$

4. Selections can be combined with Cartesian products and theta joins.

$$\sigma_{\theta}(E_1 \times E_2) = E_1 \bowtie_{\theta} E_2$$

Equivalence Rules

5. Theta-join operations (and natural joins) are commutative

$$E_1 \bowtie_{\theta} E_2 = E_2 \bowtie_{\theta} E_1$$

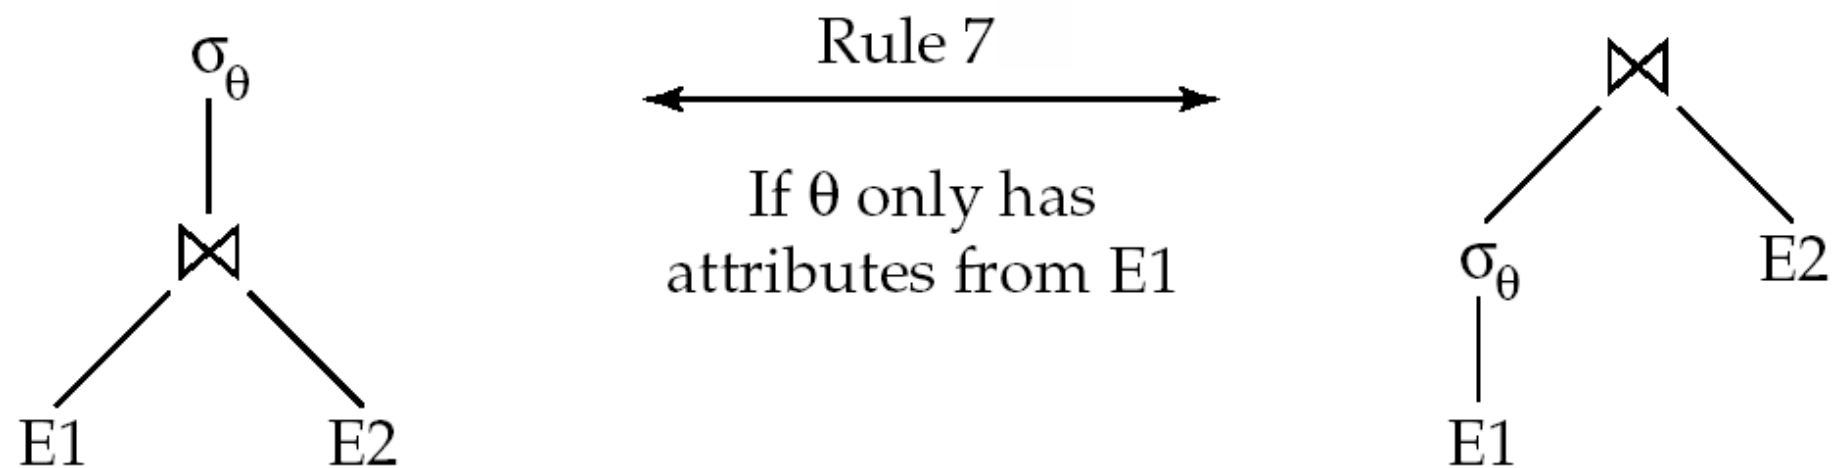


6. Natural join operations are associative:

$$(E_1 \bowtie E_2) \bowtie E_3 = E_1 \bowtie (E_2 \bowtie E_3)$$

7. The selection operation distributes over the theta join operation

$$\sigma_{\theta_0}(E_1 \bowtie_{\theta} E_2) = (\sigma_{\theta_0}(E_1)) \bowtie_{\theta} E_2$$

Equivalence Rules



Pushing Selections

- Query: Find the names of all instructors in the Music department who have taught a course in 2009, along with the titles of the courses that they taught

$$\Pi_{name, title} (\sigma_{dept_name = \text{“Music”} \wedge year = 2009} (instructor \bowtie (teaches \bowtie \Pi_{course_id, title} (course))))$$

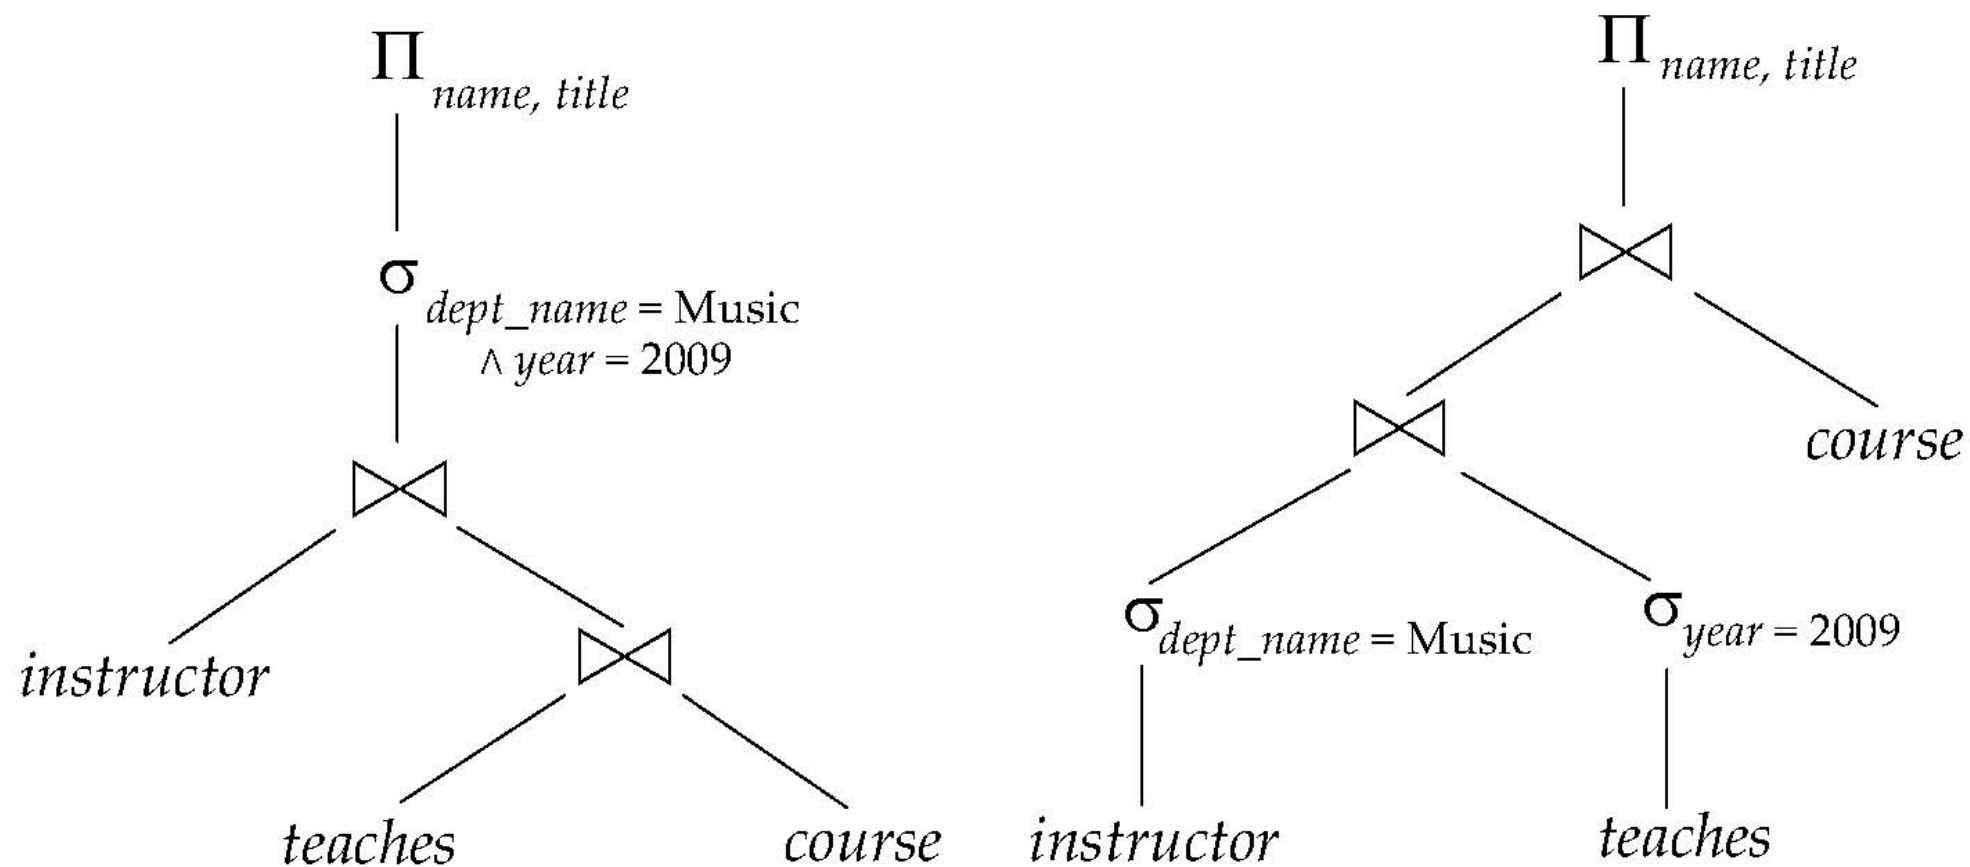
- Transformation using join associatively (Rule 6):

$$\Pi_{name, title} (\sigma_{dept_name = \text{“Music”} \wedge year = 2009} ((instructor \bowtie teaches) \bowtie \Pi_{course_id, title} (course)))$$

- Second form provides an opportunity to apply the “perform selections early” rule, resulting in the subexpression

$$\sigma_{dept_name = \text{“Music”}} (instructor) \quad \sigma_{year = 2009} (teaches)$$

Pushing Selections





(a) Initial expression tree

(b) Tree after multiple transformations

- Performing the selection as early as possible reduces the size of the relation to be joined.

Join order problem

- From previous example you know the order of join table have great effect to scan cost
- Consider finding the best join-order for $r_1 \bowtie r_2 \bowtie \dots r_n$.
 - $\{1, 2, 3\} : ((1\ 2)3), (1(2\ 3)), ((1\ 3)2), (1(3\ 2)) \dots$
 - $n! \cdot (n-1)!$, e.g. $n = 10$, the number is greater 176 billion



Selinger Optimizer

- Access Path Selection in a Relational Database Management System, SIGMOD '79 (Citation : 2412)
- *Selinger Optimizer* was used in System R of IBM and pioneered the idea of cost-based join-order optimization.
- Key implementation decision
 - Left deep tree only (((AB)C)D) (eliminate (n - 1)!))
 - Push-down selections
 - Don't consider cross products
 - Dynamic programming algorithm
 - Time : $O(n \cdot 2^n)$

Resource

- Abraham Silberschatz et al, Database System Concepts, 6 Edition, McGraw-Hill, 2010, ISBN: 0073523321
 - Lecture 13
- Database Systems, MIT 6.830 / 6.814
 - Lecture 10 & 11
- Database Internals, CSE 444
 - Query Optimization