


Overview of ARIES

ctsai@DataDB
Cloud Database
2017 Spring

Winner ? Loser ?


memory pages

page 20



Tx1 commit
Tx2 abort

page 20



Flush



- In page 20 :
 - Tx1 is a winner tx
 - Tx2 is a loser tx

Why ARIES ?

- Steal
 - Due to buffer management, dirty pages may be flushed to disk before txs commit
 - > The changes made by *loser* txs must be UNDO




Why ARIES ?

- Steal
- No Force
 - Due to performance reason, dirty page won't be flush immediately after txs commit
 - > The changes made by *winner* txs must be REDO




Three Phases of ARIES

- Analyse Phase
 - Find the earliest possibly start point of dirty page
 - Find loser txs



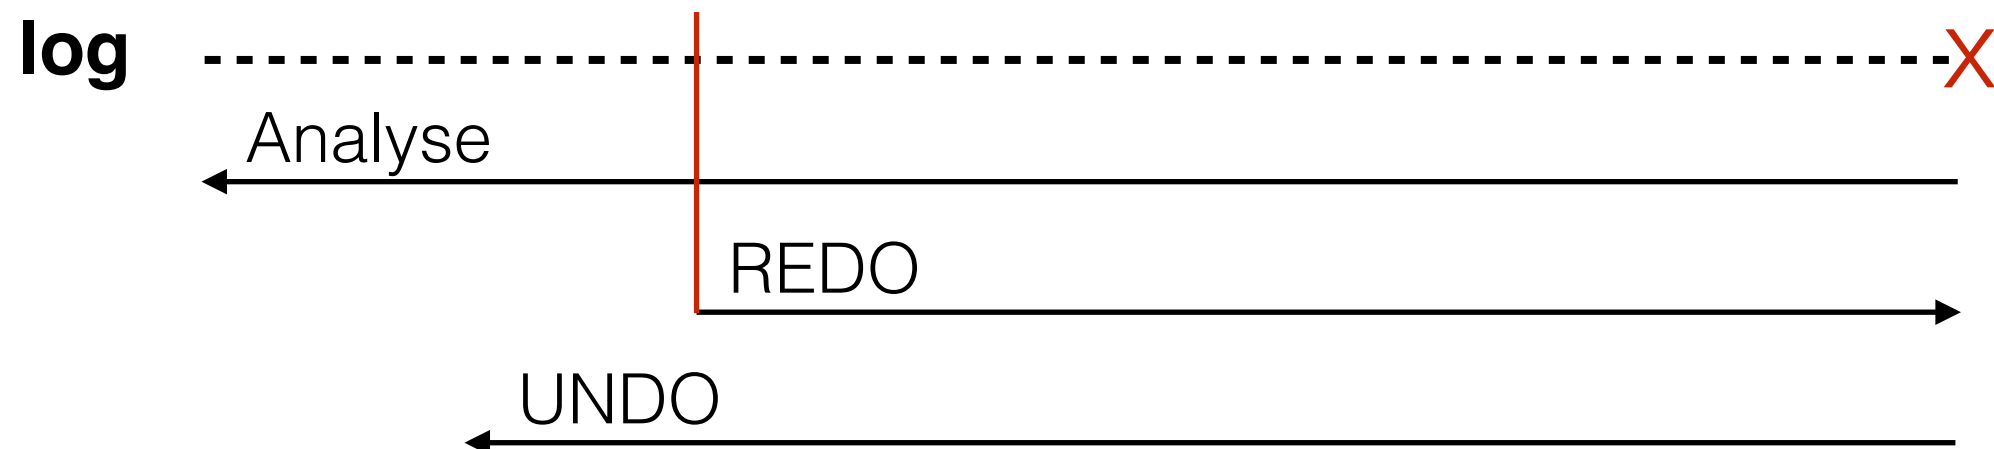
Three Phases of ARIES

- Analyse Phase
 - Find the earliest possibly start point of dirty page
 - Find loser txs
- REDO Phase
 - Repeat history (*both* winner and loser changes)
 - Recovery exact page status when the failure occurred



Three Phases of ARIES

- Analyse Phase
 - Find the earliest possibly start point of dirty page
 - Find loser txs
- REDO Phase
 - Repeat history (*both* winner and loser changes)
 - Recovery exact page status when the failure occurred
- UNDO Phase
 - Rollback *loser* txs changes

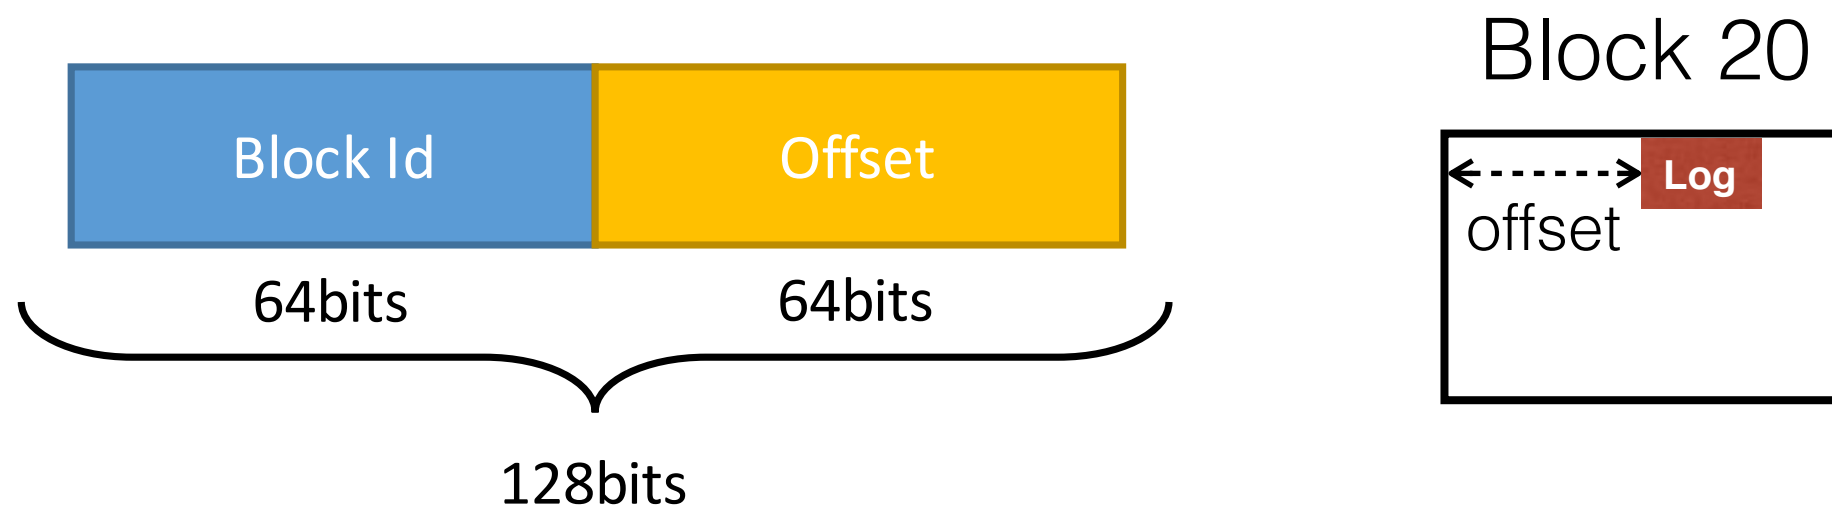


Logs in ARIES



Log Sequence Number

- Format



- Class

- LogSeqNum – Store Block , offset

Physical Log Record

- Record format :
 - Set Value Record:
<Op Code, txNum, fileName, blockNum, offset, sqlType, oldVal, newVal >
 - Index Page Insert/Delete Record :
<Op Code, txNum, fileName, blockNum, insertSlot, insertKey, insertRidBlkNum, insertRidId>
- REDO :
- UNDO :

Physical Log Record

- Record format :
 - Set Value Record:
<Op Code, txNum, fileName, blockNum, offset, sqlType, oldVal, newVal >
 - Index Page Insert/Delete Record :
<Op Code, txNum, fileName, blockNum, insertSlot, insertKey, insertRidBlkNum, insertRidId>
- REDO :
 - Apply newVal to the page
- UNDO :

Physical Log Record

- Record format :
 - Set Value Record:
<Op Code, txNum, fileName, blockNum, offset, sqlType, oldVal, newVal >
 - Index Page Insert/Delete Record :
<Op Code, txNum, fileName, blockNum, insertSlot, insertKey, insertRidBlkNum, insertRidId>
- REDO :
 - Apply newVal to the page
- UNDO :
 - Apply oldVal to the page

Physical Log Record

- Record format :
 - Set Value Record:
<Op Code, txNum, fileName, blockNum, offset, sqlType, oldVal, newVal >
 - Index Page Insert/Delete Record :
<Op Code, txNum, fileName, blockNum, insertSlot, insertKey, insertRidBlkNum, insertRidId>
- REDO :
 - Apply newVal to the page
- UNDO :
 - Apply oldVal to the page
 - Append its ***Compensation Log***

Compensation Log Record

- Feature :
 - REDO-ONLY Physical Log Record
 - UNDO procedure's REDO Log
- Record format :
 - Set Value Clr and Index Page Insert/Delete Clr:
<Op Code, UndoTxNum... , UndoNextLSN >
- REDO :
 - Apply oldVal to the page
- UNDO :
 - Do nothing

Why Clr is Redo Only ?


[0] <Start 1>

[1] <SetVal , 1 , Page 20 , 0 , 1>

[2] <SetVal , 1 , Page 20 , 1 , 2>

[3] <SetVal , 1 , Page 20 , 2 , 3>

Crash Here !



Why Clr is Redo Only ?


[0] <Start 1>

[1] <SetVal , 1 , Page 20 , 0 , 1>

[2] <SetVal , 1 , Page 20 , 1 , 2>

[3] <SetVal , 1 , Page 20 , 2 , 3>

Crash Here !



[4]<SetValClr 1, Page 20 , 3 , 2 > // Append Undo [3] Redo log

Why Clr is Redo Only ?

[0] <Start 1>

[1] <SetVal , 1 , Page 20 , 0 , 1>


[2] <SetVal , 1 , Page 20 , 1 , 2>

[3] <SetVal , 1 , Page 20 , 2 , 3>

Crash Here !

[4] <SetValClr 1, Page 20 , 3 , 2 >

Crash Again !



Why Clr is Redo Only ?

[0] <Start 1>

[1] <SetVal , 1 , Page 20 , 0 , 1>

[2] <SetVal , 1 , Page 20 , 1 , 2>

[3] <SetVal , 1 , Page 20 , 2 , 3>


Crash Here !

[4] <SetValClr 1, Page 20 , 3 , 2 >

Crash Again !

[5] <SetValClr 1, Page 20 , 2 , 3 >

// Append Undo { Undo [3] Redo log } Redo log



Why Clr is Redo Only ?

[0] <Start 1>

[1] <SetVal , 1 , Page 20 , 0 , 1>

[2] <SetVal , 1 , Page 20 , 1 , 2>

[3] <SetVal , 1 , Page 20 , 2 , 3>

Crash Here !

[4] <SetValClr 1, Page 20 , 3 , 2 >

Crash Again !

[5] <SetValClr 1, Page 20 , 2 , 3 >

Crash Again !

Redo



Analyse

Why Clr is Redo Only ?

[0] <Start 1>

[1] <SetVal , 1 , Page 20 , 0 , 1>

[2] <SetVal , 1 , Page 20 , 1 , 2>

[3] <SetVal , 1 , Page 20 , 2 , 3>

Crash Here !

[4]<SetValClr 1, Page 20 , 3 , 2 >

Crash Again !


[5]<SetValClr 1, Page 20 , 2 , 3 >

Crash Again !

[6]<SetValClr 1, Page 20 , 3 , 2 >

// Append Undo { Undo { Undo [3] Redo log } Redo log} Redo log

... WTF



Which level of REDO am
I Undoing ?



Why Clr need UndoNext ?

- Clr without UndoNext :

[0] <Start 1>

[1] <SetVal , 1 , Page 20 , 0 , 1>

[2] <SetVal , 1 , Page 20 , 1 , 2>

[3] <SetVal , 1 , Page 20 , 2 , 3>

Crash Here !

[4]<SetValClr 1, Page 20 , 3 , 2 > // Append Undo [3] Redo log

Crash Again !

[5]<SetValClr 1, Page 20 , 3 , 2 > // Append Undo [3] Redo log

[6]<SetValClr 1, Page 20 , 2 , 1 > // Append Undo [2] Redo log

[7]<SetValClr 1, Page 20 , 1 , 0 >

Why Clr need UndoNext ?

- Clr with UndoNext :

[0] <Start 1>

[1] <SetVal , 1 , Page 20 , 0 , 1>

[2] <SetVal , 1 , Page 20 , 1 , 2>

[3] <SetVal , 1 , Page 20 , 2 , 3>

Crash Here !

[4]<SetValClr 1, Page 20 , 3 , 2 , [3] > // Append Undo [3] Redo log

Crash Again !

[5]<SetValClr 1, Page 20 , 2 , 1 , [2] > // Append Undo [2] Redo log

[6]<SetValClr 1, Page 20 , 1 , 0 , [1] >

Why Clr need UndoNext ?

- UndoNext helps skip logs which have been Undone by Redo

[0] <Start 1>

[1] <SetVal , 1 , Page 20 , 0 , 1>

[2] <SetVal , 1 , Page 20 , 1 , 2>

[3] <SetVal , 1 , Page 20 , 2 , 3>


Crash Here !

[4]<SetValClr 1, Page 20 , 3 , 2 , [3] >

Crash Again !

[5]<SetValClr 1, Page 20 , 2 , 1 , [2] >

[7]<SetValClr 1, Page 20 , 1 , 0 , [1] >



Logical Log Record

- Record format :
 - Logical - Start Record
<OP Code, txNum>
 - Record File Insert/Delete End Record :
<Op Code, txNum, fileName, blockNum, slotId , logicalStartLSN>
 - Index Insert/Delete End Record :
<Op Code, txNum, tblName, fldName, searchKey, recordBlockNum, recordSlotId , logicalStartLSN>
- REDO :
 - Do nothing
- UNDO :
 - Undo *competed* logical log *logically*
 - Undo *partial* logical log *physically*
 - Append *Logical Abort* log record

Rollback a completed Logical Log Record

[0] <Start 1>

[1] <LogicalStart, 1 >

[2] <Index Page Insert , 1 , ... >

[3] <SetVal , 1 , Page 2 , 1 , 2>

[4] <SetVal , 1 , Page 20 , 2 , 3>

[5] <Record File Insert End , 1, ... , [2] >

Crash Here !

Rollback a completed Logical Log Record

[0] <Start 1>

[1] <LogicalStart, 1 >

[2] <Index Page Insert , 1 , ... >

[3] <SetVal , 1 , Page 2 , 1 , 2>

[4] <SetVal , 1 , Page 20 , 2 , 3>

[5] <Record File Insert End , 1, ... , [2] >

Logical
operations

Crash Here !

Rollback a completed Logical Log Record

[0] <Start 1>

[1] <LogicalStart, 1 >

[2] <Index Page Insert , 1 , ... >

[3] <SetVal , 1 , Page 2 , 1 , 2>

[4] <SetVal , 1 , Page 20 , 2 , 3>

[5] <Record File Insert End , 1, ... , [2] >

Physical
operations

Logical
operations

Crash Here !

Rollback a completed Logical Log Record

[0] <Start 1>

[1] <LogicalStart, 1 >

[2] <Index Page **Insert** , 1 , ... >

[3] <SetVal , 1 , Page 2 , 1 , 2>

[4] <SetVal , 1 , Page 20 , 2 , 3>

[5] <Record File Insert End , 1, ... , [2] >

Crash Here !

[6] <Start 2 >

[7] <LogicalStart, 2 >

[8] <Index Page **Delete** , 2 , ... >

[9] <SetVal , 2 , Page 2 , 2 , 1>

[10] <SetVal , 2 , Page 20 , 2 , 3>

[11] <Record File Insert End , 2, ... , [7]>

[12] <Logical Abort 1 , [1]>

Rollback a completed Logical Log Record

[0] <Start 1>

[1] <LogicalStart, 1 >

[2] <Index Page **Insert** , 1 , ... >

[3] <SetVal , 1 , Page 2 , 1 , 2>

[4] <SetVal , 1 , Page 20 , 2 , 3>

[5] <Record File Insert End , 1, ... , [2] >

Crash Here !

[6] <Start 2 >

[7] <LogicalStart, 2 >

[8] <Index Page **Delete** , 2 , ... >

[9] <SetVal , 2 , Page 2 , 2 , 1>

[10] <SetVal , 2 , Page 20 , 2 , 3>

[11] <Record File Insert End , 2, ... , [7]>

[12] <Logical Abort 1 , [1]>

Physical
operations

Logical
operations