


SQL Queries

Shan-Hung Wu & DataLab
CS, NTHU



PostgreSQL

- [Download and install](#)
- For Mac users, try [PostgreSQL.app](#)

Today's Example Code

- You can find the example we use today from here:
 - <https://shwu10.cs.nthu.edu.tw/courses/databases/2019-spring/class-materials/blob/master/01-postgresql-example.md>

Using PostgreSQL

```
$ createdb <db>
$ psql <db> [user]
> \h or \?
> SELECT now(); -- SQL commands
```

- Default schema: `public`
 - `\dn` for listing all schemas
- Multiple lines until `;`
- `--` for comments
- ***Case insensitive***
 - Use `""` to distinguish lower and upper cases
 - E.g., `SELECT "authorId" FROM posts;`

Structured Query Language (SQL)

- Data Definition Language (DDL) on schema
 - CREATE TABLE ...
 - ALTER TABLE ...
 - DROP TABLE ...
- Data Manipulation Language (DML) on records
 - INSERT INTO ... VALUES ...
 - SELECT ... FROM ... WHERE ...
 - UPDATE ... SET ... WHERE ...
 - DELETE FROM ... WHERE ...

Schema

users

<u>id</u>	name	karma
729	Bob	35
730	John	0

friend

<u>uid1</u>	<u>uid2</u>	since
729	730	14928063
729	882	14827432

foreign keys

posts

<u>id</u>	text	authorId	ts
33981	'Hello DB!'	729	1493897351
33982	'Show me code'	729	1493854323

Creating Tables/Relations

```
CREATE TABLE posts (  
  id          serial PRIMARY KEY NOT NULL,  
  text        text NOT NULL,  
  "authorId" integer NOT NULL  
              REFERENCES users ON DELETE CASCADE,  
  ts          bigint NOT NULL  
              DEFAULT (extract(epoch from now())) ,  
  ...  
);
```

- Column types:
 - Integer, bigint, real, double, etc.
 - varchar(10), text, etc.
- Non-null constraint
- Default values

Creating Tables/Relations

```
CREATE TABLE posts (  
  id          serial PRIMARY KEY NOT NULL,  
  text        text NOT NULL,  
  "authorId" integer NOT NULL  
              REFERENCES users ON DELETE CASCADE,  
  ts          bigint NOT NULL  
              DEFAULT (extract(epoch from now())) ,  
  ...  
);
```

- Primary key:
 - Unique (no duplicate values among rows)
 - Usually of type “serial” (auto-filled integer)
 - Index automatically created

Creating Tables/Relations

```
CREATE TABLE posts (  
  id          serial PRIMARY KEY NOT NULL,  
  text        text NOT NULL,  
  "authorId" integer NOT NULL  
              REFERENCES users ON DELETE CASCADE,  
  ts          bigint NOT NULL  
              DEFAULT (extract(epoch from now())) ,  
  ...  
);
```

- Foreign key: post.authorId must be a valid user.id
- When deleting a user (row):
 - NO ACTION (default): user not deleted, error raised
 - CASCADE: user **and all referencing posts** deleted

Inserting Rows

```
INSERT INTO posts(text, "authorId", ...)
VALUES ('Today is a good day!', 5, ...);
```

- String values should be *single* quoted
- Inserting dummy rows:

```
INSERT INTO posts(text, "authorId")
SELECT
    'Dummy word ' || i || '.',
    round(random() * 10) + 1
FROM generate_series(1, 20) AS s(i);
```

Queries

```
SELECT *  
FROM posts  
WHERE ts > 147988213 AND text ILIKE '%good%'  
ORDER BY ts DESC, id ASC  
LIMIT 2;
```

- To see how a query is processed:

```
EXPLAIN ANALYZE -- show plan tree  
SELECT *  
FROM posts  
WHERE ts > 147988213 AND text ILIKE '%good%'  
ORDER BY ts DESC, id ASC  
LIMIT 2;
```

(Batch) Updating Rows

```
UPDATE posts SET ts = ts + 3600 WHERE "authorId" = 10;
```

- ***All*** rows satisfying the WHERE clause will be updated
- $ts + 3600$ is an ***expression***
 - Can be evaluated to a single value

Handling “Big” Data

```
INSERT INTO posts(text, "authorId")
SELECT
    'Dummy word ' || i || '.',
    rount(random() * 10) + 1
FROM generate_series(1, 1000000) AS s(i);
```

- Some queries will be slow:

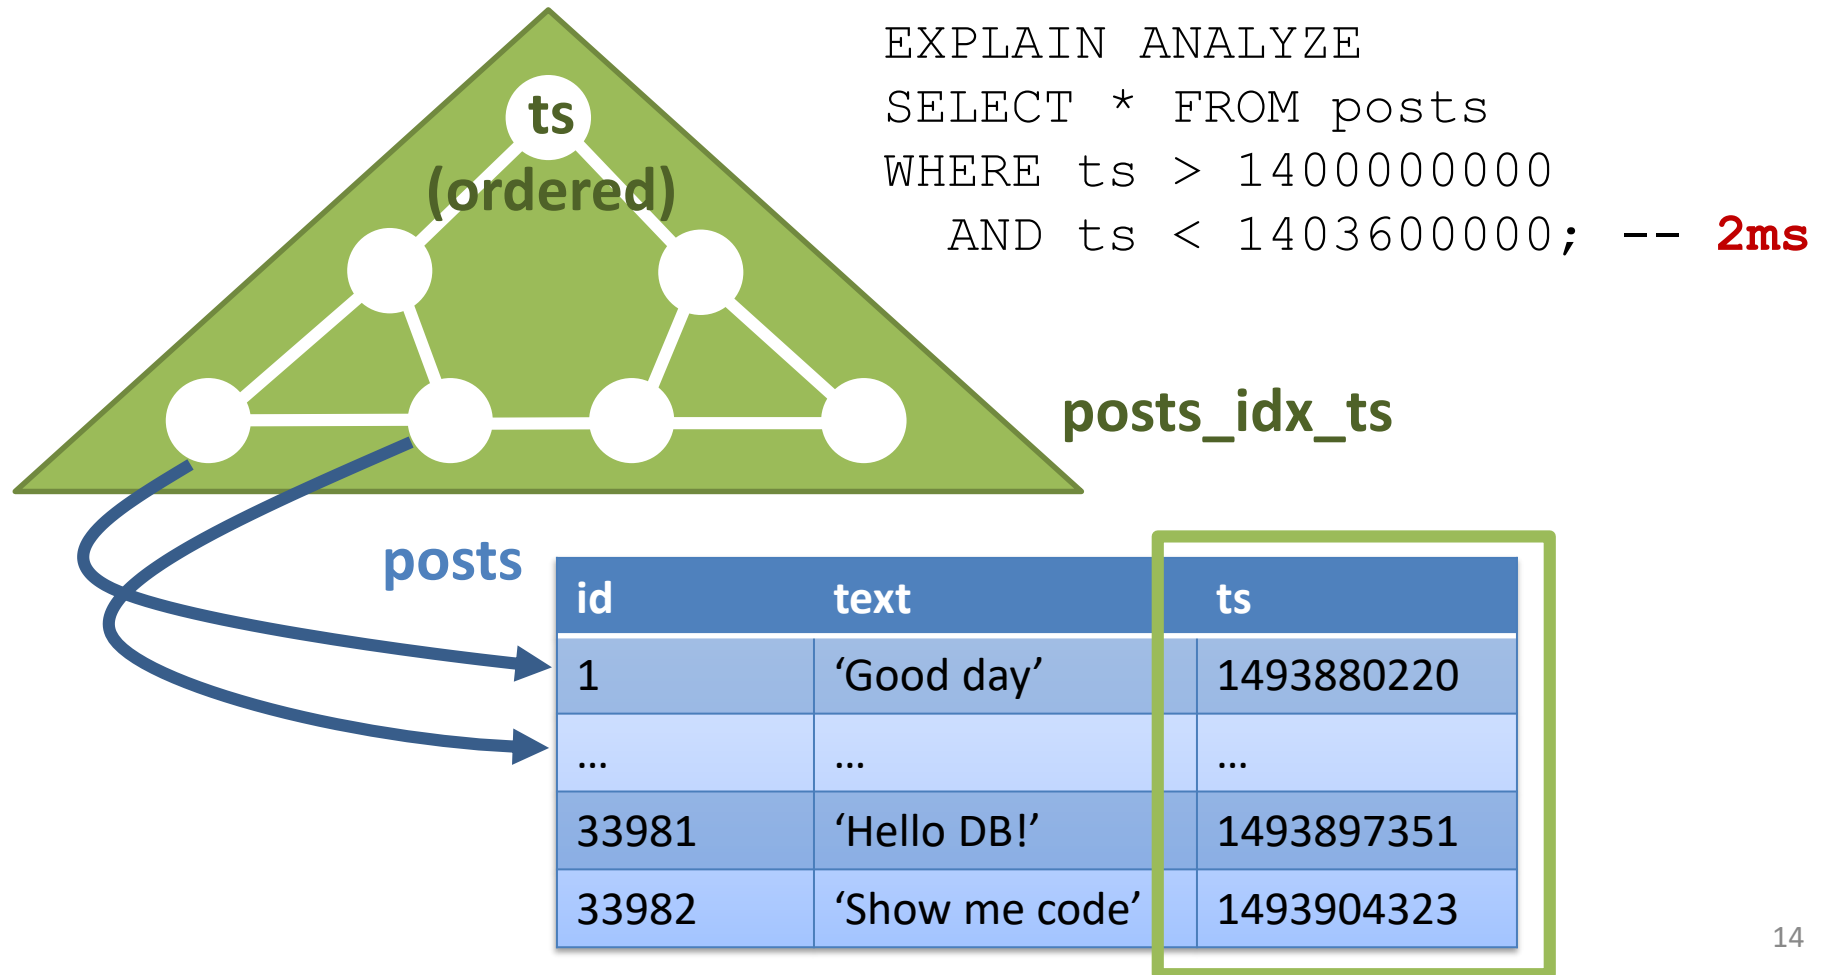
```
EXPLAIN ANALYZE SELECT * FROM posts
WHERE id > 500000 AND id < 501000; -- 1ms
```

```
EXPLAIN ANALYZE SELECT * FROM posts
WHERE ts > 14000000000 AND ts < 14036000000; -- 230ms
```

Using Index

```
CREATE INDEX posts_idx_ts
ON posts
USING btree(ts);
\di -- list indices
```

```
EXPLAIN ANALYZE
SELECT * FROM posts
WHERE ts > 1400000000
AND ts < 1403600000; -- 2ms
```



Index for ILIKE?

```
CREATE INDEX posts_idx_text ON posts  
USING btree(text);
```

```
EXPLAIN ANALYZE SELECT * FROM posts  
WHERE text ILIKE '% word 500000%'; -- 1.5s
```

- B-tree indices are **not** helpful for text searches
- Use GIN (generalized inverted index) instead:

```
CREATE EXTENSION pg_trgm;  
\dx -- list extensions  
CREATE INDEX posts_idx_text_trgm ON posts  
USING gin(text gin_trgm_ops);
```

```
EXPLAIN ANALYZE SELECT * FROM posts  
WHERE text ILIKE '%word 500000%'; -- 50ms
```