


# Using a DBMS

Shan-Hung Wu & DataLab  
CS, NTHU

# Assigned Reading

- [Java concurrency](#)
- "Database Management Systems," 3ed, by Ramakrishnan



# Coverage

- Java concurrency
- Chaps 2 and 3 on how to store your data into a DBMS
  - ER model and relational model
- Chaps 4 and 5 on queries
  - SQL language (DDL and DML)
  - Relational algebra
- Chap 19\* on how to store your data *well*
  - Easy maintenance
  - Answering most queries fast

# Coverage

- Java concurrency
- Chaps 2 and 3 on how to store your data into a DBMS
  - ER model and relational model
- Chaps 4 and 5 on queries
  - SQL language (DDL and DML)
  - Relational algebra
- Chap 19\* on how to store your data *well*
  - Easy maintenance
  - Answering most queries fast

# Starting a New Thread

```
public class HelloRunnable implements Runnable {  
    @Override  
    public void run() {  
        System.out.println("Hello from a thread!");  
    }  
  
    public static void main(String args[]) {  
        (new Thread(new HelloRunnable())).start();  
    }  
}
```

**or**

```
public class HelloThread extends Thread {  
    @Override  
    public void run() {  
        System.out.println("Hello from a thread!");  
    }  
  
    public static void main(String args[]) {  
        (new HelloThread()).start();  
    }  
}
```

# What Happened?

```
public class HelloRunnable implements Runnable {  
    @Override  
    public void run() {  
        System.out.println("Hello from a thread!");  
    }  
  
    public static void main(String args[]) {  
        (new Thread(new HelloRunnable())).start();  
    }  
}
```

- A new stack is allocated for `run()`, in addition to that of `main()`
- Your CPU spends time on executing `run()` ***in parallel with*** `main()`

# Multiple Stacks, Single Heap

- The heap in memory scheme?
  - Stores objects
  - *Shared by all threads*
- Can two threads access the same object? Yes
- How? Passing the same object to their constructors

```
public static void main(String args[]) {  
    Counter counter = ...;  
    (new Thread(new HelloRunnableA(counter))).start(); // thread A  
    (new Thread(new HelloRunnableB(counter))).start(); // thread B  
}
```

# Concurrent Access

- Given the same object counter
- Suppose both threads execute in `run()`:

```
...  
int c = counter.get();  
c++; // c--;  
counter.set(c);
```

```
class Counter {  
    private int c = 0;  
    public void set(int c) {  
        this.c = c;  
    }  
    public int get () {  
        return c;  
    }  
}
```

- Thread A's result will be lost if
  1. Thread A: Get c
  2. Thread B: Get c
  3. Thread A: Increment retrieved value; result is 1
  4. Thread B: Decrement retrieved value; result is -1
  5. Thread A: Set result in c; c is now 1.
  6. Thread B: Set result in c; c is now -1.



# Synchronization

```
public class SynchronizedCounter {  
    private int c = 0;  
    public synchronized void set(int c) {  
        this.c = c;  
    }  
    public synchronized int get() {  
        return c;  
    }  
}
```

- Same as...

- Only one thread can enter sync. block of an obj. at a time
- Problem solved?

```
public class SynchronizedCounter {  
    private int c = 0;  
    public void set(int c) {  
        synchronized(this) { this.c = c; }  
    }  
    public int get() {  
        synchronized(this) { return c; }  
    }  
}
```

# Still Wrong!

- Two threads in `run()`:

```
... // counter is a SynchronizedCounter instance
int c = counter.get();
c++; // c--;
counter.set(c);
```

- Thread A's result will still be lost if
  1. Thread A: Get c
  2. Thread B: Get c
  3. Thread A: Increment retrieved value; result is 1
  4. Thread B: Decrement retrieved value; result is -1
  5. Thread A: Set result in c; c is now 1.
  6. Thread B: Set result in c; c is now -1.

# Synchronization at Right Place

- Solution1: callers lock `counter` during the entire increment/decrement period:

```
synchronized(counter) {  
    int c = counter.get();  
    c++; // or c--;  
    counter.set(c);  
}
```

- Solution2: callee provides atomic methods

```
public class SynchronizedCounter {  
    private int c = 0;  
    public void synchronized increment() {  
        c++;  
    }  
    public int get() {  
        return c;  
    }  
}
```

# Blocking and Waiting States

- Threads are **blocked** outside a critical section if someone is in
- Thread *A* in a critical section of `o` can stop and enter the **waiting** state by calling `o.wait()`
  - Gives up the lock, so some other blocking thread *B* can enter the critical section
  - If *B* calls `o.notifyAll()`, *A* competes for the lock again and resume

# Wrap `wait()` in a Loop

- It's a good practice to wrap `wait()` in a loop to prevent bugs
- Queue length: 10

## Threads A, B:

```
// enqueue
synchronized(queue) {
    while (queue.size() == 10) {
        queue.wait();
    }
    queue.add(...);
    queue.notifyAll();
}
```

## Threads C, D:

```
// dequeue
synchronized(queue) {
    while (queue.size() == 0) {
        queue.wait();
    }
    ... = queue.remove();
    queue.notifyAll();
}
```

# Coverage

- Java concurrency
- Chaps 2 and 3 on how to store your data into a DBMS
  - ER model and relational model
- Chaps 4 and 5 on queries
  - SQL language (DDL and DML)
  - Relational algebra
- Chap 19\* on how to store your data *well*
  - Easy maintenance
  - Answering most queries fast

# DBMS $\neq$ Database

- A database is a collection of your data stored in a computer
- A DBMS (DataBase Management System) is a software that manages databases

# Storing Data

- Let's say, you have data in memory to store
- How does data in memory (heap) look like?
  - Objects
  - References to objects
- Objects formatted by classes you defined
- Could we store these objects and references directly?



# Data Model

- Definition: A ***data model*** is a framework for describing the structure of databases in a DBMS
- Common data models: ***ER model*** and ***relational model***
- A DBMS supporting the relational model is called the relational DBMS


# Why ER Model?

- Allows thinking your data in OOP way
- **Entity**
  - An object (or instance of a class)
  - With attributes
- **Entity group**
  - A class
  - Must define the ID attribute for each entity
- **Relationship** between entities
  - References (“has-a” relationship)
  - Could be 1-1, 1-many, and many-many

# Why Relational Model?

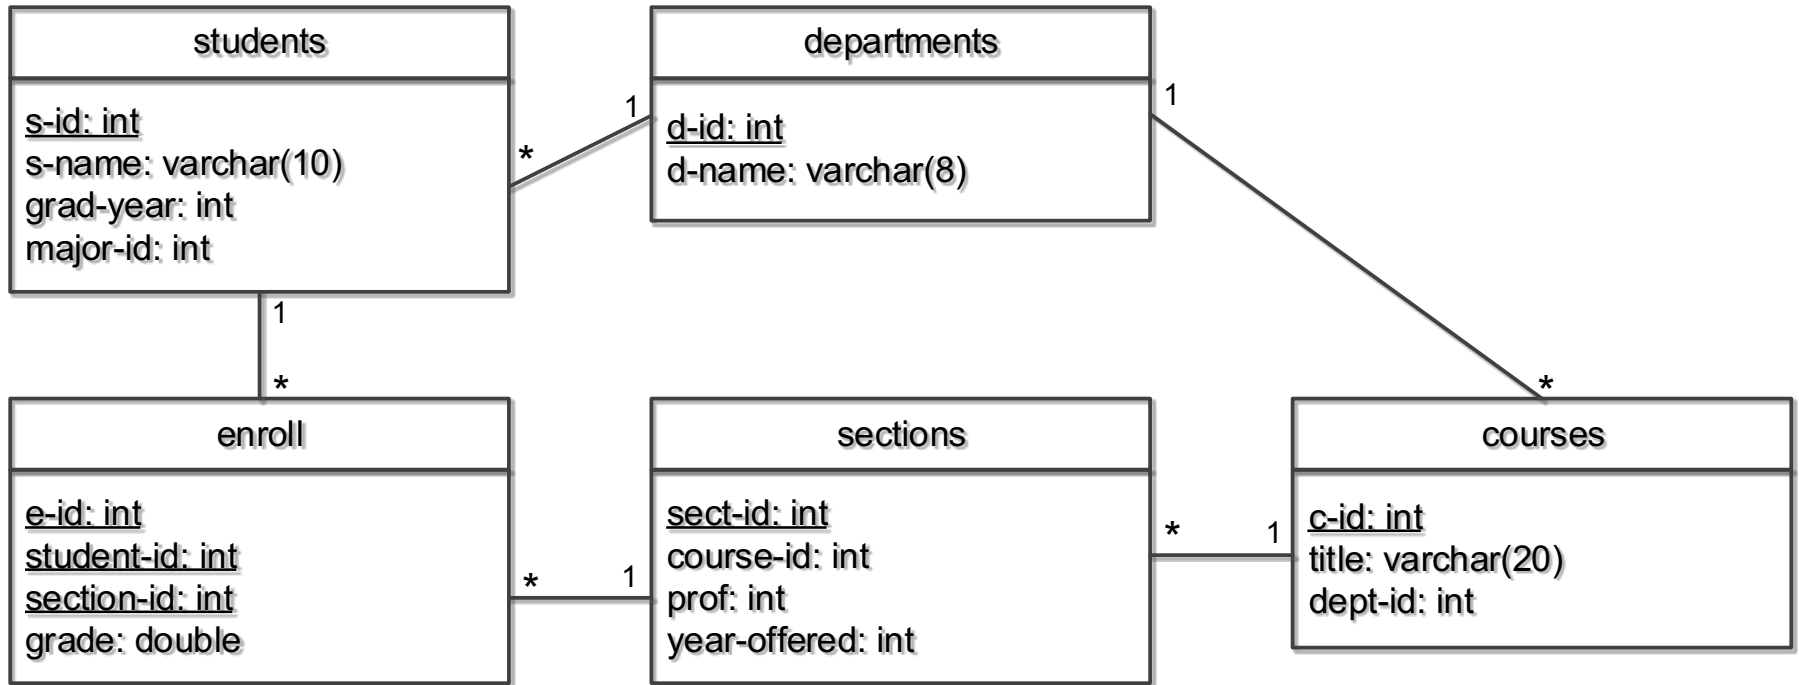
- To realize an ER model using a collocation of tables/***relations***
  - Simplifies data management and query processing
- Still logic (not how your data stored physically)

# Example: Student DB



- Relation (table)
  - Realization of 1) an entity group via table; or 2) a relationship
  - **Fields/attributes** as columns
  - **Records/tuples** as rows

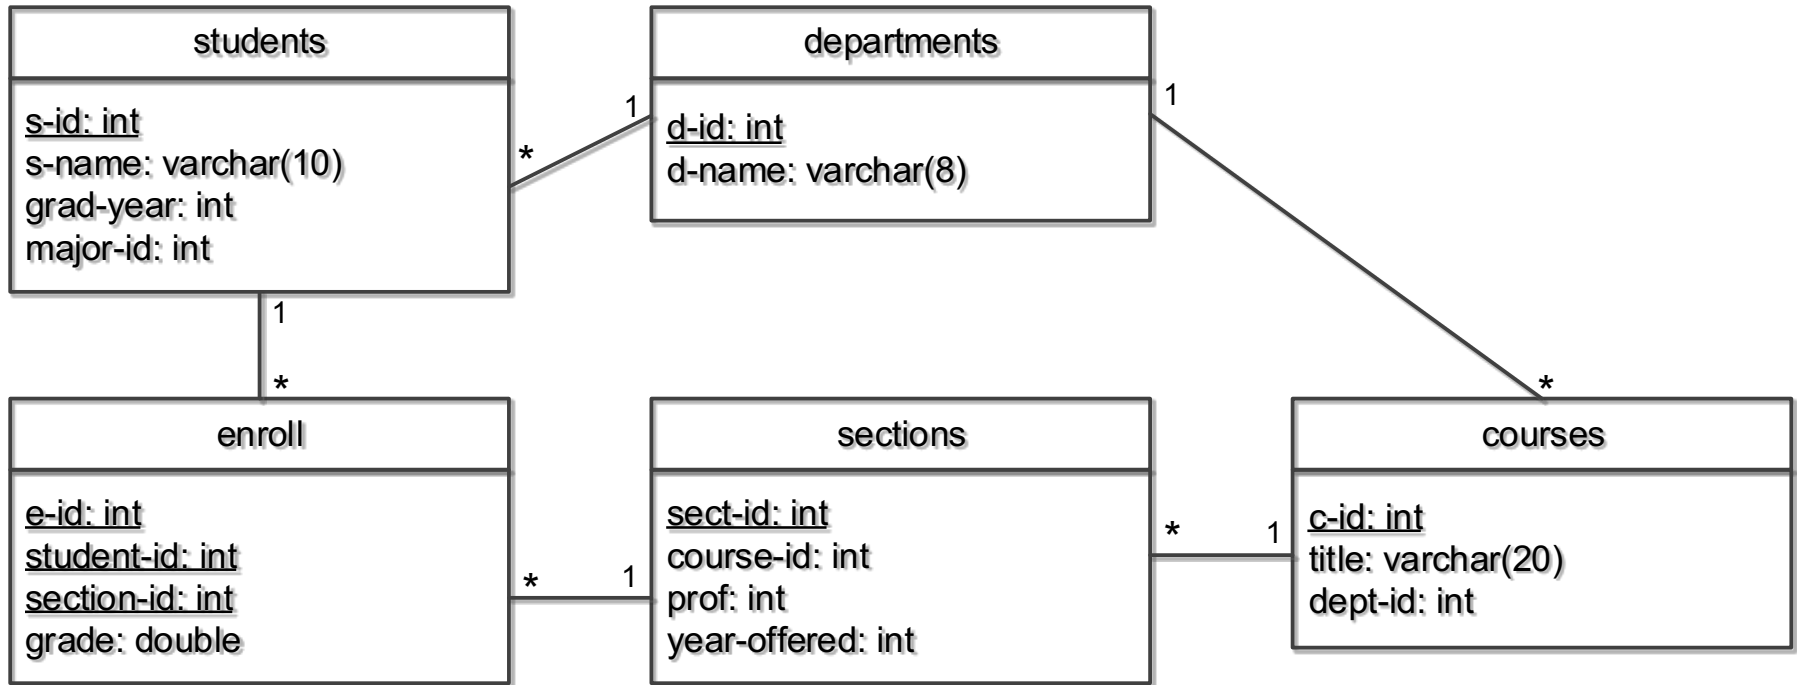
# Example: A student DB



- **Primary Key**

- Realization of ID via a group of fields

# Example: A student DB



- **Foreign key**
  - Realization of relationship
  - A record can point to the primary key of the other record
  - Only 1-1 and 1-many
  - Intermediate relation is needed for many-many

# Schema

- Definition: A *schema* is the structure of a particular database
- The schema of a relation/table is its fields and field types

# Coverage

- Java concurrency
- Chaps 2 and 3 on how to store your data into a DBMS
  - ER model and relational model
- Chaps 4 and 5 on queries
  - SQL language (DDL and DML)
  - Relational algebra
- Chap 19 on how to store your data *well*
  - Easy maintenance
  - Answering most queries fast



# Queries

- Data Definition Language (DDL) on schema
  - CREATE TABLE ...
  - ALTER TABLE ...
  - DROP TABLE ...
- Data Manipulation Language (DML) on records
  - INSERT INTO ... VALUES ...
  - SELECT ... FROM ... WHERE ...
  - UPDATE ... SET ... WHERE ...
  - DELETE FROM ... WHERE ...

# Data Model and Queries (1/3)

Q1: among a set of blog pages, find those pages written by Steven Sinofsky after 2011

Step1: structure your data by following the *relational data model*

- Identify *records* (e.g., web pages, authors, etc.) with the same *fields* in your data and place them into respective *tables*

## blog\_pages

blog_id	url	created	author_id
33981	ms.com/...	2012/10/31	729
33982	apache.org/...	2012/11/15	4412

← record

field



## users

user_id	name	balance
729	Steven Sinofsky	10,235
730	Picachu	NULL

# Data Model and Queries (2/3)

Q1: among a set of blog pages, find those pages written by Steven Sinofsky after 2011

```
CREATE TABLE blog_pages (  
    blog_id INT NOT NULL AUTO_INCREMENT,  
    url VARCHAR(60),  
    created DATETIME,  
    author_id INT);
```

```
INSERT INTO blog_pages (url, created, author_id)  
VALUES ('ms.com/...', 2012/09/18, 729);
```

## blog\_pages

blog_id	url	created	author_id
33981	ms.com/...	2012/10/31	729
33982	apache.org/...	2012/11/15	4412

# Data Model and Queries (3/3)

Q1: among a set of blog pages, find those pages written by Steven Sinofsky after 2011

Step2: issue queries

```
SELECT b.blog_id
FROM blog_pages b, users u
WHERE b.author_id=u.user_id
      AND u.name='Steven Sinofsky'
      AND b.created >= 2011/1/1;
```

# How Is a Query Answered?

```
SELECT b.blog_id
```

```
FROM blog_pages b, users u
```

```
WHERE b.author_id=u.user_id  
      AND u.name='Steven Sinofsky'  
      AND b.created >= 2011/1/1;
```

**product(b, u)**

blog_id	url	created	author_id	user_id	name	balance
33981	...	2009/10/31	729	729	Steven Sinofsky	10,235
33981	...	2009/10/31	729	730	Picachu	NULL
33982	...	2012/11/15	4412	729	Steven Sinofsky	10,235
33982	...	2012/11/15	4412	730	Picachu	NULL
41770	...	2012/10/20	729	729	Steven Sinofsky	10,235
41770	...	2012/10/20	729	730	Picachu	NULL

**b**

blog_id	url	created	author_id
33981	...	2009/10/31	729
33982	...	2012/11/15	4412
41770	...	2012/10/20	729

**u**

user_id	name	balance
729	Steven Sinofsky	10,235
730	Picachu	NULL

# How Is a Query Answered?

```
SELECT b.blog_id
```

```
FROM blog_pages b, users u
```

```
WHERE b.author_id=u.user_id  
      AND u.name='Steven Sinofsky'  
      AND b.created >= 2011/1/1;
```

**select(p, where...)**

blog_id	url	created	author_id	user_id	name	balance
41770	...	2012/10/20	729	729	Steven Sinofsky	10,235



**p = product(b, u)**

blog_id	url	created	author_id	user_id	name	balance
33981	...	2009/10/31	729	729	Steven Sinofsky	10,235
33981	...	2009/10/31	729	730	Picachu	NULL
33982	...	2012/11/15	4412	729	Steven Sinofsky	10,235
33982	...	2012/11/15	4412	730	Picachu	NULL
41770	...	2012/10/20	729	729	Steven Sinofsky	10,235
41770	...	2012/10/20	729	730	Picachu	NULL

# How Is a Query Answered?

```
SELECT b.blog_id
```

```
FROM blog_pages b, users u
```

```
WHERE b.author_id=u.user_id  
      AND u.name='Steven Sinofsky'  
      AND b.created >= 2011/1/1;
```

**project(s, select...)**

blog_id
41770


**s = select(p, where...)**



blog_id	url	created	author_id	user_id	name	balance
41770	...	2012/10/20	729	729	Steven Sinofsky	10,235

# Query Algebra


- Operators
  - Product, select, project, join, group-by, etc.
- Operands
  - Tables, output of other operators, predicates, etc.





# Query Plan

- A tree that answers a query



- ***Not unique!***
- A DBMS automatically seeks for the best query plan

# Coverage

- Java concurrency
- Chaps 2 and 3 on how to store your data into a DBMS
  - ER model and relational model
- Chaps 4 and 5 on queries
  - SQL language (DDL and DML)
  - Relational algebra
- Chap 19\* on how to store your data ***well***
  - Easy maintenance
  - Answering most queries fast

# How Good are Your Data?

- Let's say, if you want to track the topics of a blog page
- Is this a good table?

## blog\_pages

blog_id	url	created	author_id	topic	topic_admin
33981	ms.com/...	2012/10/31	729	programming	5638
33981	ms.com/...	2012/10/31	729	databases	5649
33982	apache.org/...	2012/11/15	4412	programming	5638
33982	apache.org/...	2012/11/15	4412	os	7423

# Insertion Anomaly

## blog\_pages

blog_id	url	created	author_id	topic	topic_admin
33981	ms.com/...	2012/10/31	729	programming	5638
33981	ms.com/...	2012/10/31	729	databases	5649
33982	apache.org/...	2012/11/15	4412	programming	5638
33982	apache.org/...	2012/11/15	4412	os	7423

33983	apache.org/...	2013/02/15	7412		
-------	----------------	------------	------	--	--




- A blog cannot be inserted without knowing all fields of topics (except setting them to null)

# Update Anomaly

## blog\_pages

blog_id	url	created	author_id	topic	topic_admin
33981	ms.com/...	2012/10/31	729	<i>win prog.</i>	5638
33981	ms.com/...	2012/10/31	729	databases	5649
33982	apache.org/...	2012/11/15	4412	programming	5638
33982	apache.org/...	2012/11/15	4412	os	7423




- If you forget to update all duplicated cells, you get inconsistent data

# Deletion Anomaly

## blog\_pages

blog_id	url	created	author_id	topic	topic_admin
33981	ms.com/...	2012/10/31	729	<i>programming</i>	<i>5638</i>
33981	ms.com/...	2012/10/31	729	databases	5649
33982	apache.org/...	2012/11/15	4412	programming	5638
33982	apache.org/...	2012/11/15	4412	os	7423



- Deleting topics force you to delete the blog fields too

# Normalization

- Avoids these anomaly through schema normalization
  - 3<sup>rd</sup> normal form
  - BCNF normal form
- Idea: break your one, big table into multiple small, modular tables
  - Reuse tables
  - Avoid bias towards any particular pattern of querying