

# Indexing

Shan-Hung Wu

CS, NTHU

# Outline

- Overview
  - API in VanillaCore
- Hash-Based Indexes
- B-Tree Indexes
- Query Processing
- Transaction Management

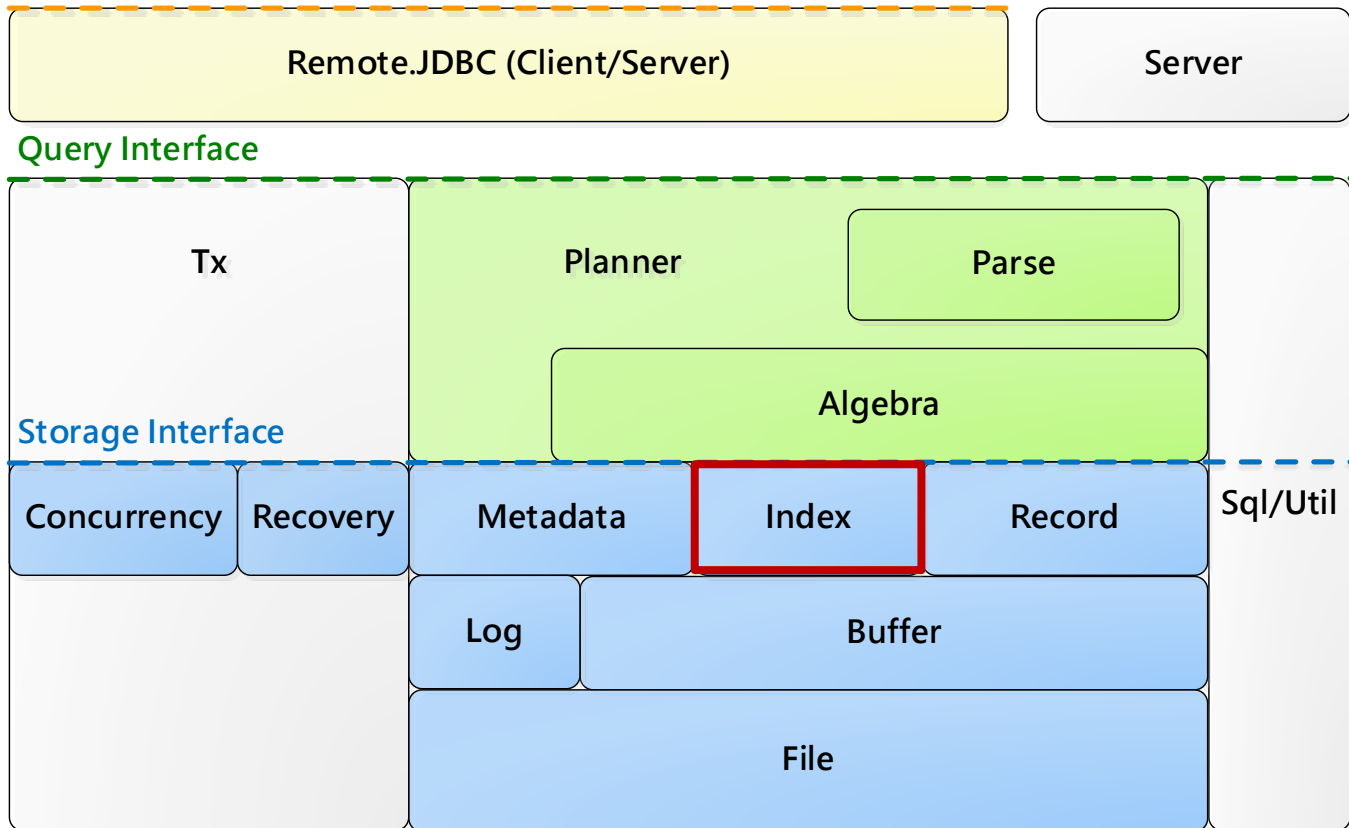
# Outline

- Overview
  - API in VanillaCore
- Hash-Based Indexes
- B-Tree Indexes
- Query Processing
- Transaction Management

# Where are we?

VanillaCore

JDBC Interface (at Client Side)



# Why Index?


- Query:
  - `SELECT * FROM students WHERE dept = 10`
- Record file for students:

1 joe 10 2015	2 kay 20 2013	4 rob 20 2011	5 tom 10 2013	6 bob 20 2016	9 jim 20 2011
r1	r2	r3	r4	r5	r6

- Selectivity is usually low
- Full table scan results in poor performance

# What is an Index?


- Query:
  - `SELECT * FROM students WHERE dept = 10`
- **Index**: a data structure (file) defined on *fields* that speeds up data accessing
  - Input: field values or ranges
  - Output: rids




# Terminology (1/2)

- Every index has an associated **search key**
  - I.e., one or more fields

Search key: dept



- **Primary index** vs. **secondary index**
  - If search key contains primary key or not
- Index entry/record:
  - <data value, data rid>



# Terminology (2/2)

- An index is designed to speed up *equality* or *range selections* on the search key
  - ... WHERE dept = 10
  - ... WHERE dept > 30 AND dept < 100



# Outline

- Overview
  - API in VanillaCore
- Hash-Based Indexes
- B-Tree Indexes
- Query Processing
- Transaction Management

# SQL Statements for Index Creation

- The SQL:1999 standard does not include any statement for creating or dropping indeice
- Creating index:
  - `CREATE INDEX <name> ON <table>(<fields>) USING <method>`
  - E.g., `CREATE INDEX idxdept ON students(dept) USING btree`
- In VanillaCore, an index only supports *one* indexed field

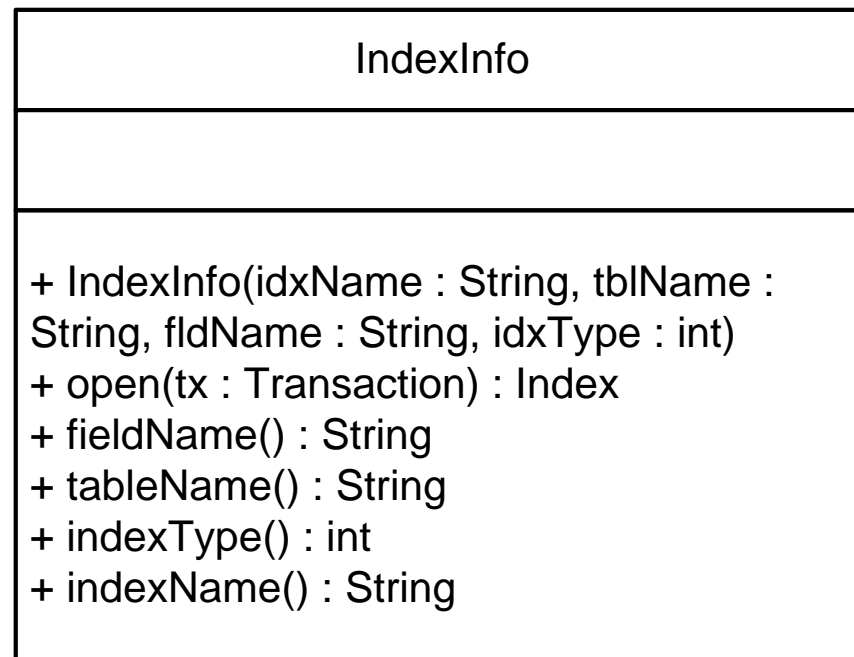
# The Index Class in VanillaCore

- An abstract class in `storage.index`
  - `beforeFirst()` resets iterator and search value
  - `next()` moves to the next rid *matching search value*

<div>&lt;&lt;abstract&gt;&gt; Index</div>
<div><u>&lt;&lt;final&gt;&gt; + IDX_HASH : int</u> <u>&lt;&lt;final&gt;&gt; + IDX_BTREE : int</u></div>
<div><u>+ searchCost(idxType : int, fldType : Type, totRecs : long, matchRecs : long) : long</u> <u>+ newInstance(ii : IndexInfo, fldType : Type, tx : Transaction) : Index</u>  &lt;&lt;abstract&gt;&gt; + beforeFirst(searchkey : ConstantRange) &lt;&lt;abstract&gt;&gt; + next() : boolean &lt;&lt;abstract&gt;&gt; + getDataRecordId() : RecordId &lt;&lt;abstract&gt;&gt; + insert(key : Constant, dataRecordId : RecordId) &lt;&lt;abstract&gt;&gt; + delete(key : Constant, dataRecordId : RecordId) &lt;&lt;abstract&gt;&gt; + close() &lt;&lt;abstract&gt;&gt; + preLoadToMemory()</div>

# IndexInfo

- Factory class for `Index` via `open ( )`
- Stores information about an index
- Similar to `TableInfo`



# Using an Index

- `SELECT sname FROM students WHERE dept=10`

```
Transaction tx = VanillaDb.txMgr().newTransaction(
    Connection.TRANSACTION_SERIALIZABLE, false);

// Open a scan on the data table
Plan studentPlan = new TablePlan("students", tx);
TableScan studentScan = (TableScan) studentPlan.open();

// Open index on the field dept of students table
Map<String, IndexInfo> idxmap =
    VanillaDb.catalogMgr().getIndexInfo("students", tx);
Index deptIndex = idxmap.get("dept").open(tx);

// Retrieve all index records having dataval of 10
deptIndex.beforeFirst(ConstantRange
    .newInstance(new IntegerConstant(10)));
while (deptIndex.next()) {
    // Use the rid to move to a student record
    RecordId rid = deptIndex.getDataRecordId();
    studentScan.moveToRecordId(rid);
    System.out.println(studentScan.getVal("sname"));
}

deptIndex.close();
studentScan.close();
tx.commit();
```

# Updating Indexes

- INSERT INTO students (sid,sname,dept,gradyear)  
VALUES (7,'sam',10,2014)

```
Transaction tx = VanillaDb.txMgr().newTransaction(
    Connection.TRANSACTION_SERIALIZABLE, false);
TableScan studentScan = (TableScan) new TablePlan("students", tx).open();

// Create a map containing all indexes of students table
Map<String, IndexInfo> idxMap = VanillaDb.catalogMgr().getIndexInfo(
    "students", tx);
Map<String, Index> indexes = new HashMap<String, Index>();
for (String fld : idxmap.keySet())
    indexes.put(fld, idxMap.get(fld).open(tx));

// Insert a new record into students table
studentScan.insert();
studentScan.setVal("sid", new IntegerConstant(7));
studentScan.setVal("sname", new VarcharConstant("sam"));
studentScan.setVal("dept", new IntegerConstant(10));
studentScan.setVal("grad", new IntegerConstant(2014));

// Insert a record into each of the indexes
RecordId rid = studentScan.getRecordId();
for (String fld : indexes.keySet()) {
    Constant val = studentScan.getVal(fld);
    Index idx = indexes.get(fld);
    idx.insert(val, rid);
}

for (Index idx : indexes.values())
    idx.close();
studentScan.close();
tx.commit();
```

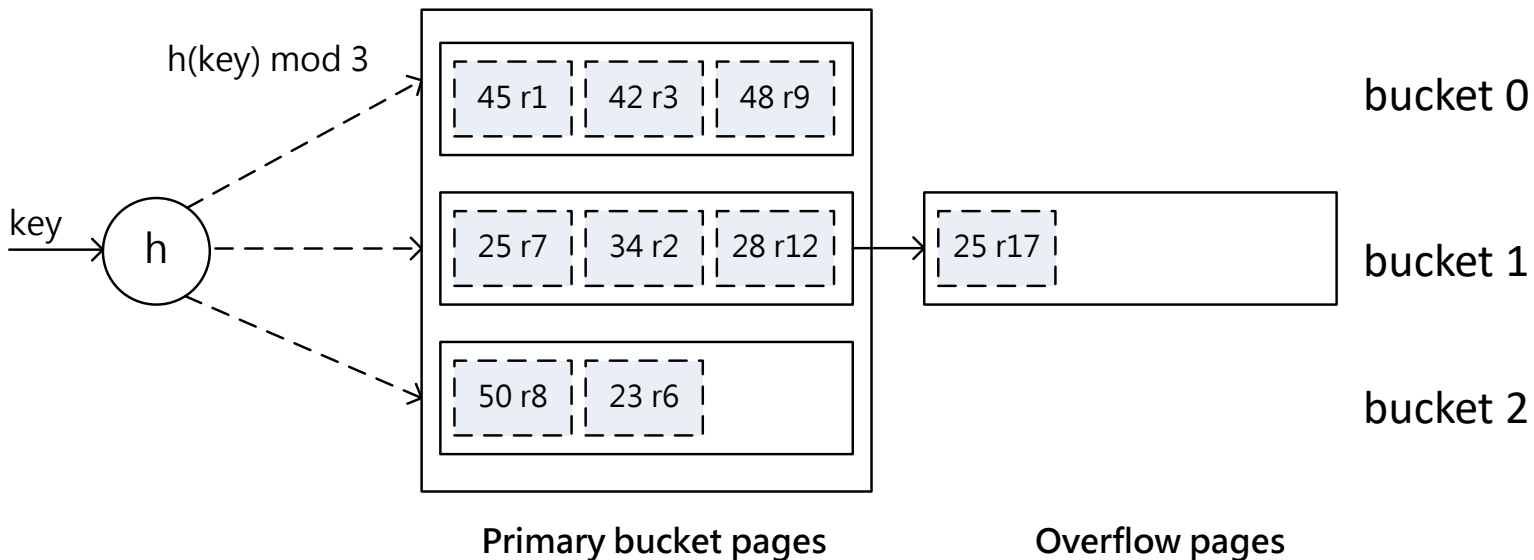
- *Faster reads* at the  
cost of *slower writes*

# Outline

- Overview
  - API in VanillaCore
- Hash-Based Indexes
- B-Tree Indexes
- Query Processing
- Transaction Management

# Hash-Based Indexes


- Designed for equality selections
- Uses a **hashing function**
  - Search values  $\rightarrow$  **bucket** numbers
- Bucket
  - Primary page plus zero or more overflow pages
- Based on static or dynamic hashing techniques






# Static Hashing

- The number of bucket  $N$  is fixed
- Overflow pages if needed
- $h(k) \bmod N = \text{bucket to which data entry with key } k \text{ belongs}$
- Records having the same hash value are stored in the same bucket



# Search Cost of Static Hashing

- How to compute the #block-access?
- Assume index has B blocks and has N buckets
- Then each bucket is about  $B/N$  blocks long



#rec = 13  
rpb = 3  
 $B = \lceil 13/3 \rceil = 5$   
 $N = 3$   
#blockAccess = 2

# Hash Index in VanillaCore

- Related Package
  - `storage.index.hash.HashIndex`

HashIndex
<u>&lt;&lt;final&gt;&gt; + NUM_BUCKETS : int</u>
<u>+ searchCost(ifldType : Type, totRecs : long, matchRecs : long) : long</u>  + HashIndex(ii : IndexInfo, fldtype : Type, tx : Transaction) + beforeFirst(searchRange : ConstantRange) + next() : boolean + getDataRecordId() : RecordId + insert(key : Constant, dataRecordId : RecordId) + delete(key : Constant, dataRecordId : RecordId) + close() + preLoadToMemory()

# HashIndex

- Stores each bucket in a record file
  - Name: {index-name}{bucket-num}
- `beforeFirst ( )`
  1. Hashes the search value, and
  2. Opens the corresponding record file
- The index record [key, blknum, id]

key	block	id
45	235	20

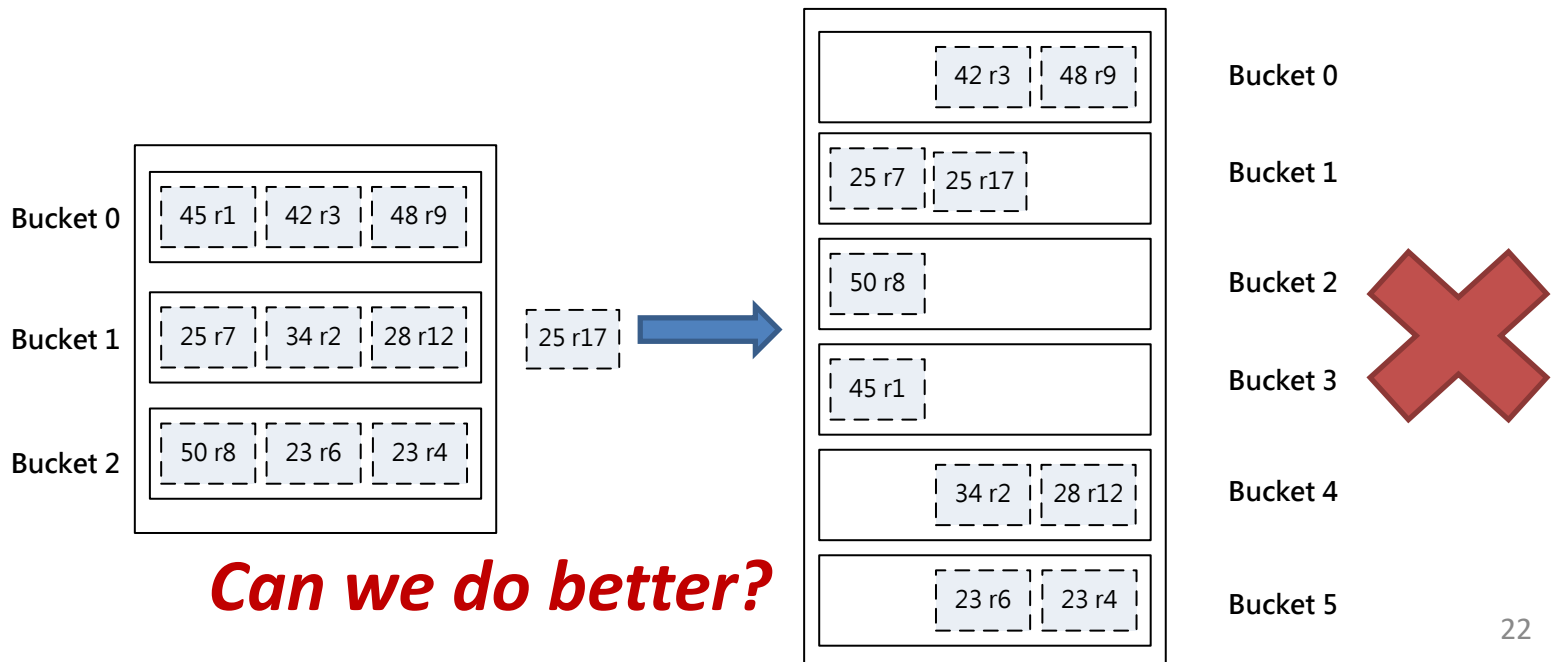
RecordId

# Limitations of Static Hashing (1/2)

- Search cost:  $B/N$
- Increase efficiency  $\rightarrow$  increase  $N$  (#buckets)
  - Best when 1 block per bucket
- However, a large #buckets leads to wasted space
  - Empty pages waiting the index to grow into it

# Limitations of Static Hashing (2/2)

- Hard to decide N
- Why not double #buckets when a bucket is full?
  - Redistributing records is costly

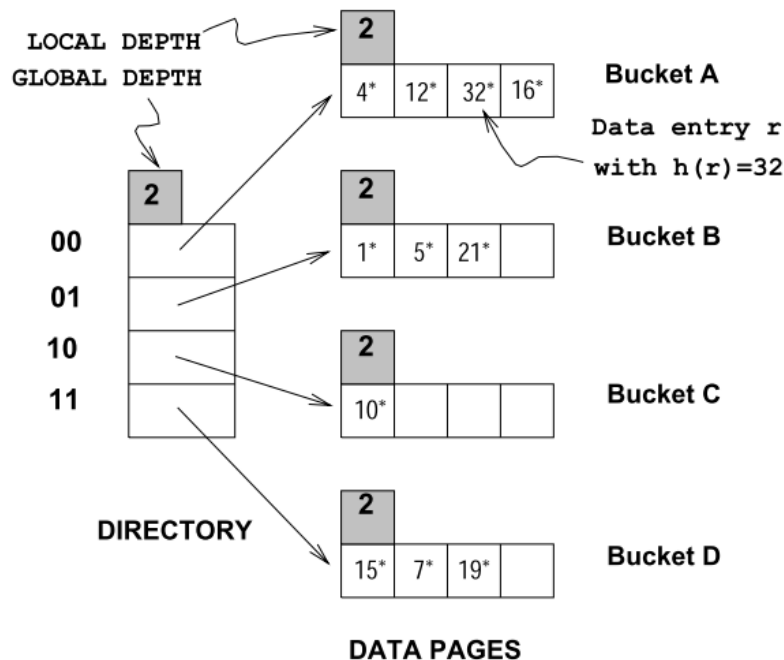


# Extendable Hash Indexes

- Use *directory*: pointers to buckets
- Double #buckets by doubling the directory
- Splitting just the bucket that overflowed

# Extendable Hash Indexes

- Directory is array of size 4
- To find bucket for  $r$ , take last 'global depth' #bits of  $h(r)$



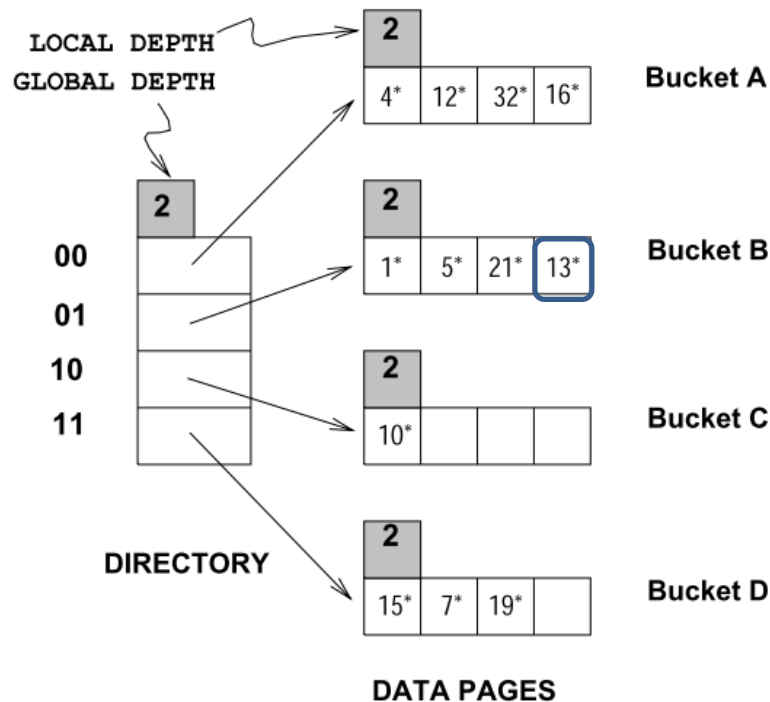
**Global depth** of directory:  
Max #bits needed to tell  
which bucket an entry belongs to

**Local depth** of a bucket:  
#bits used to determine if an  
entry belongs to this bucket



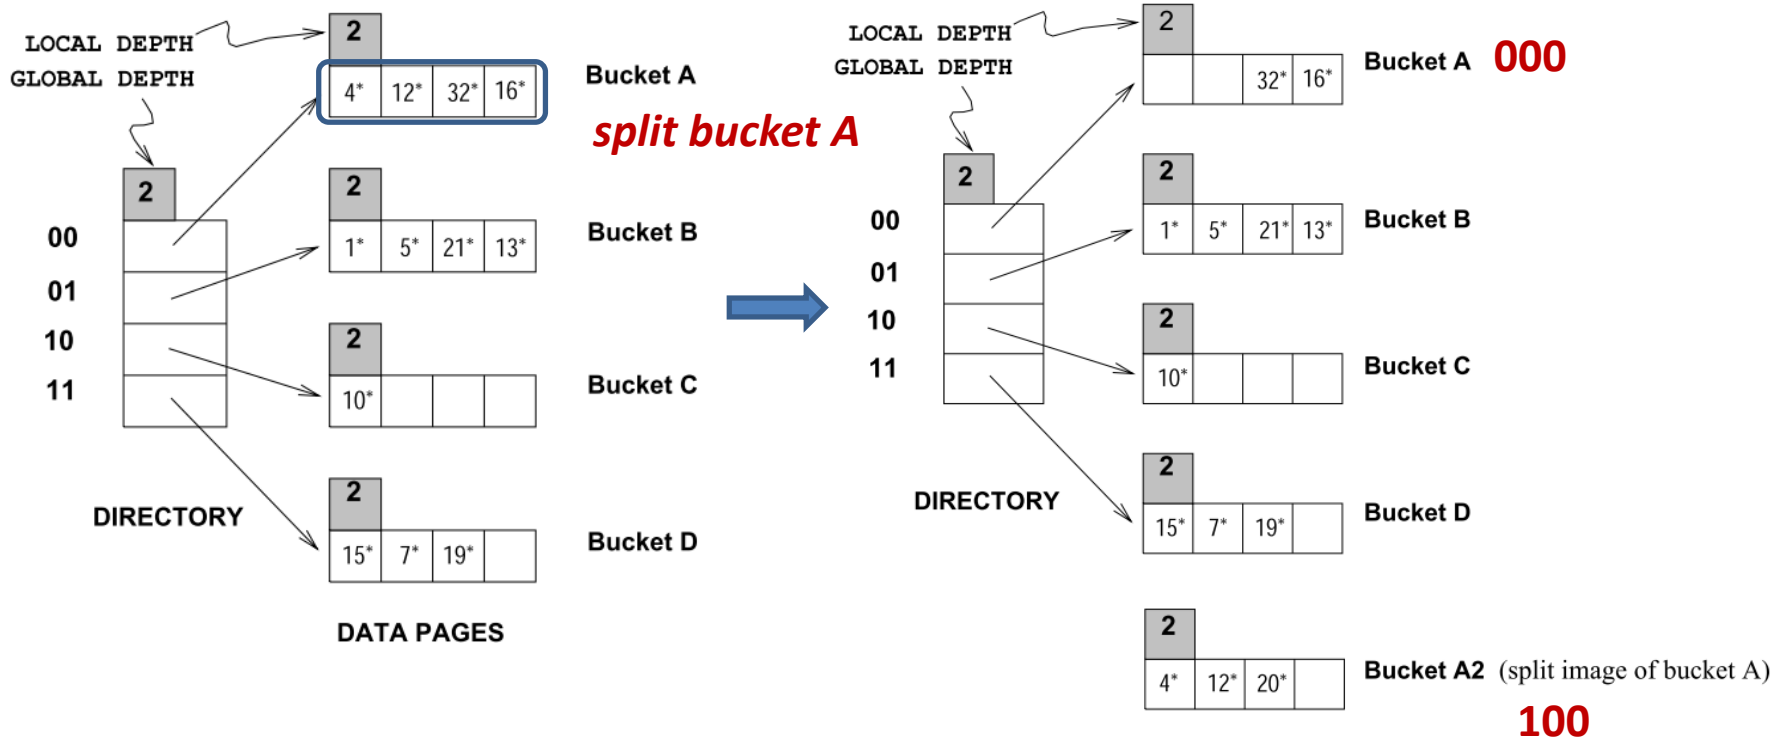
# Example (1/4)

- After inserting entry  $r$  with  $h(r)=13$ 
  - Binary number: 11**01**



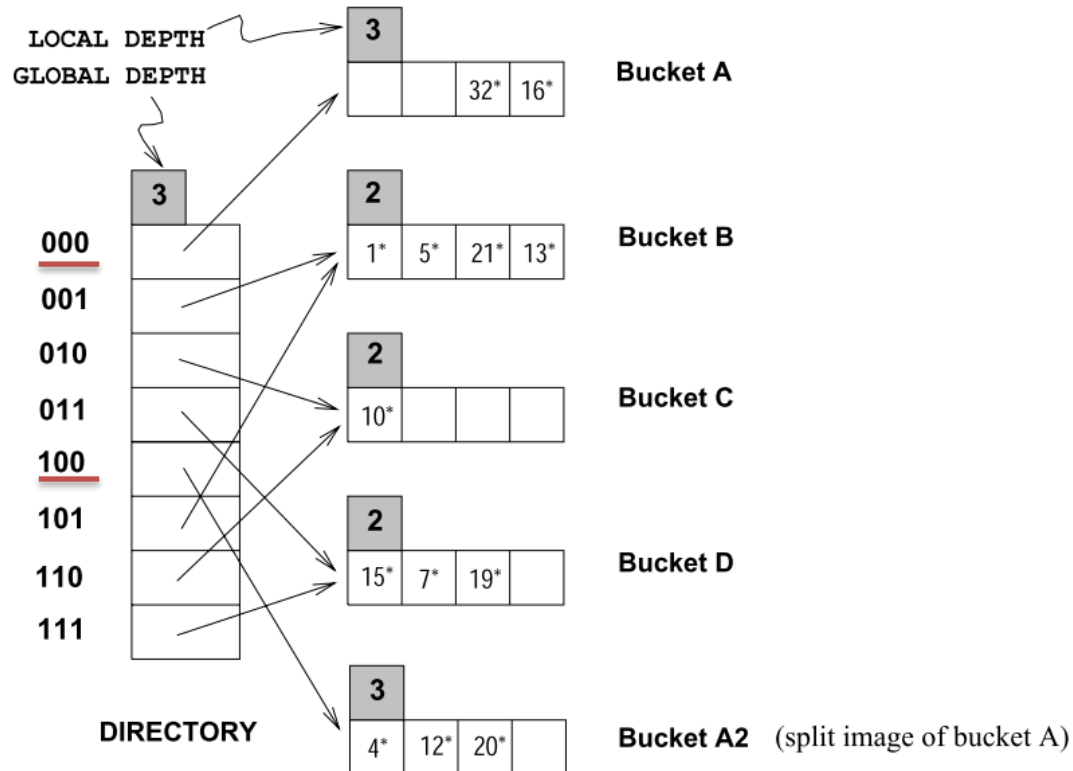
## Example (2/4)

- While inserting entry  $r$  with  $h(r)=20$ 
  - Binary number: 10100



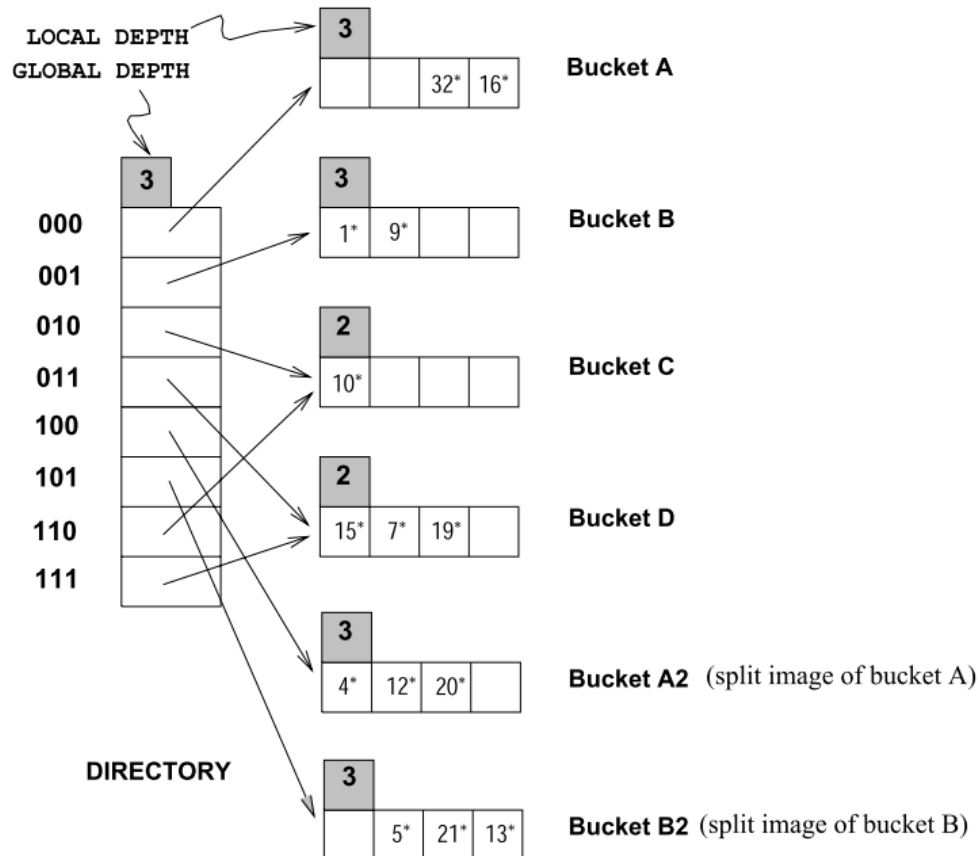
# Example (3/4)

- After inserting entry  $r$  with  $h(r)=20$
- Update the global depth
  - Some buckets will have local depth less than global depth



# Example (4/4)

- After inserting entry  $r$  with  $h(r)=9$



# Remarks

- At most 1 page split for each insert
- Cheap doubling
  - When local depth of bucket = global depth
  - Only 3 page access (1 directory page, 2 data pages)
- No overflow page?
  - Still has, but only when there are a lot of records with same key value

# Outline

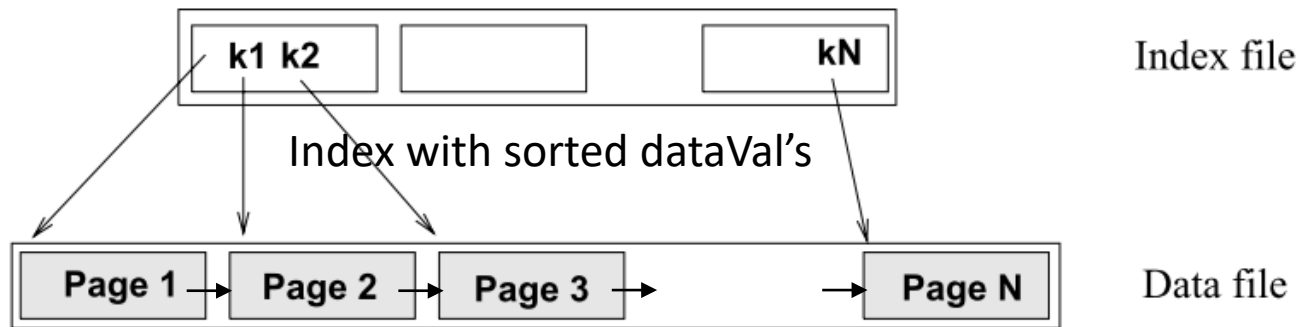
- Overview
  - API in VanillaCore
- Hash-Based Indexes
- **B-Tree Indexes**
- Query Processing
- Transaction Management

# Is Hash-Based Index Good Enough?

- Hash-based indexes are good for equality selections
- However, cannot support *range searches*
  - E.g., . . . WHERE dept>100
- We now consider an index structured as a *search tree*
  - Speeds up search by *sorting* values
  - Supports *both* range and equality searches

# Power of Sorting

- Create an “index” file
  - where dataVal's are sorted
- Query: “Find all students with dept > 100”
  - Do **binary search** to find first such student, then scan the index till end to find others

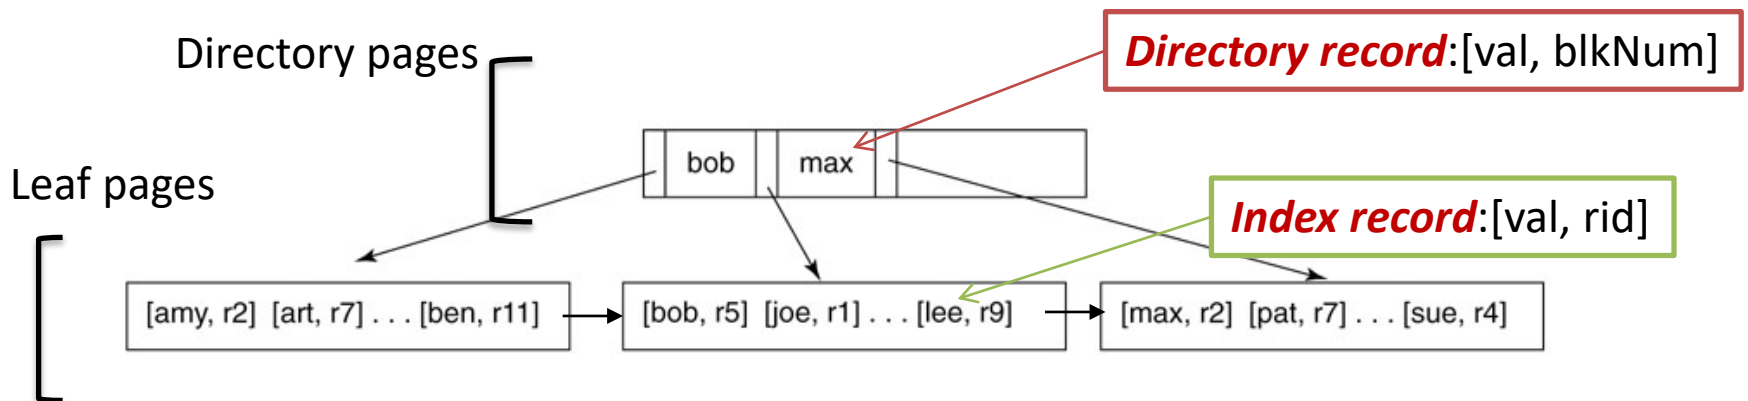


- However, slow update:  $O(\#data\text{-}records)$



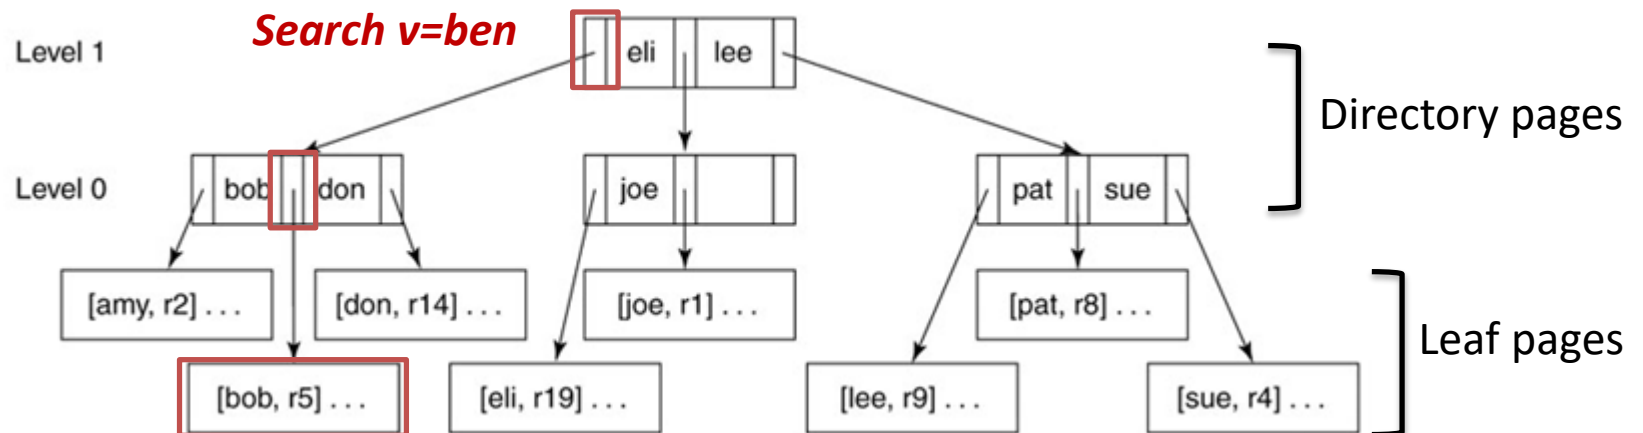
# B-Tree Index

- The most widely used index
- Index records are sorted on dataVal in each page
- M-way balanced search tree:
  - $O(\log_M(\#data-records))$  for equality search & update
  - $O(\#data-records)$  for range search



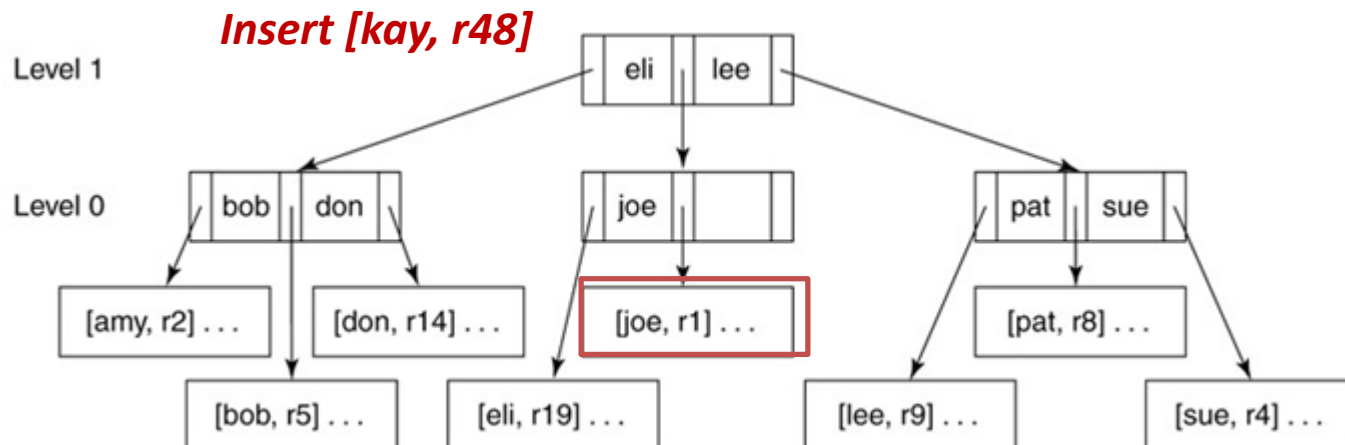
# Searching

- “Finding all index records having a specified dataVal v”
  1. Search begins at root
  2. Fetches child block pointed to by parent until leaf
- Search cost:  $O(\text{tree height})$ , usually  $< 5$



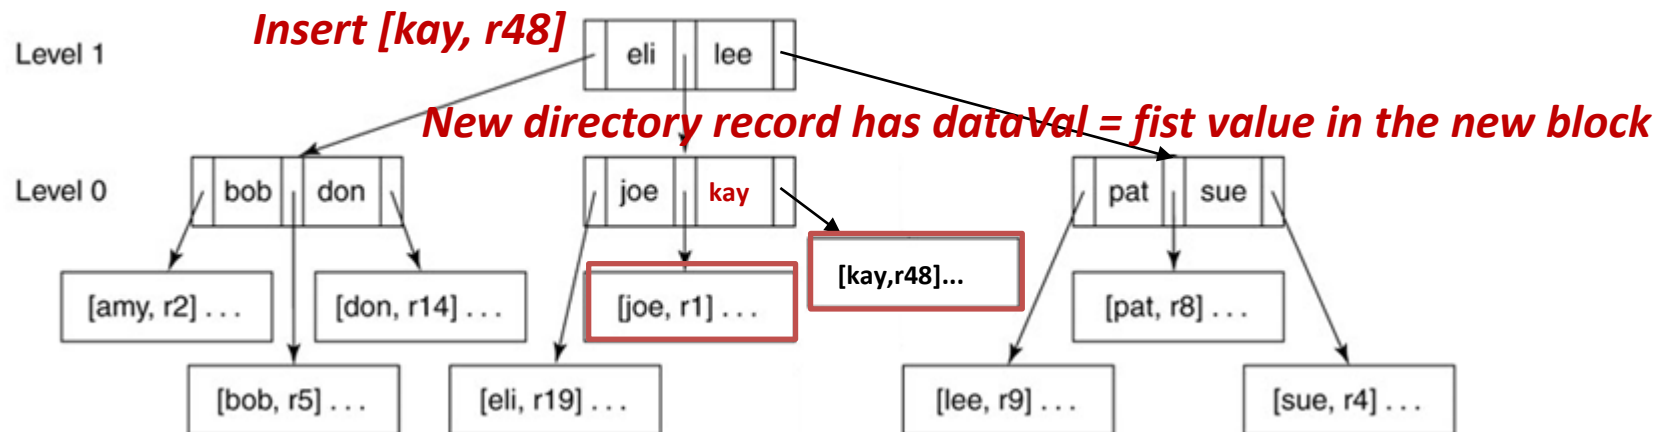
# Insertion

1. Search the index with the inserted dataVal
  2. Insert the new index record into the target leaf block
- What if the block has no more room?
    - Remember extendable hashing? *Spilt it!*



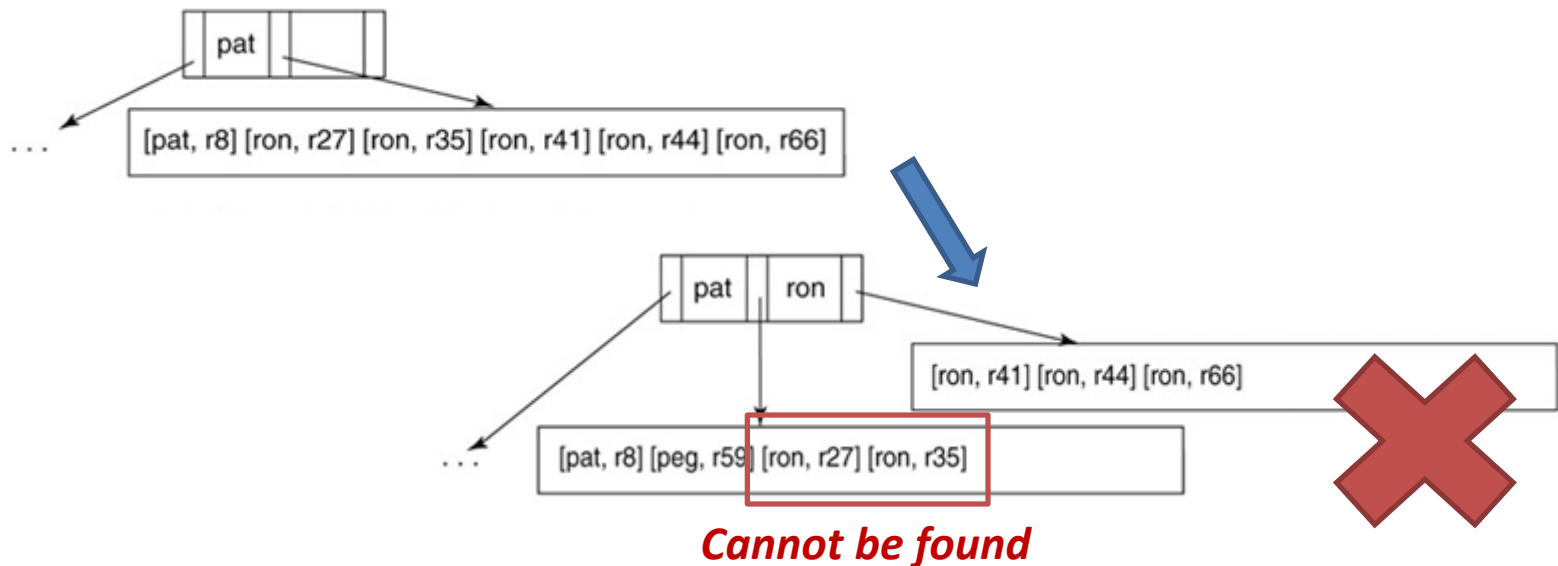
# Splitting

1. Allocate a new block in the index file
  2. Move the high-valued half of the index record into this new block
  3. Create a directory record for the new block
  4. Insert the new directory record into the same level-0 directory block
  5. Recursively split directory block if necessary
- Update cost:  $O(\text{tree height})$



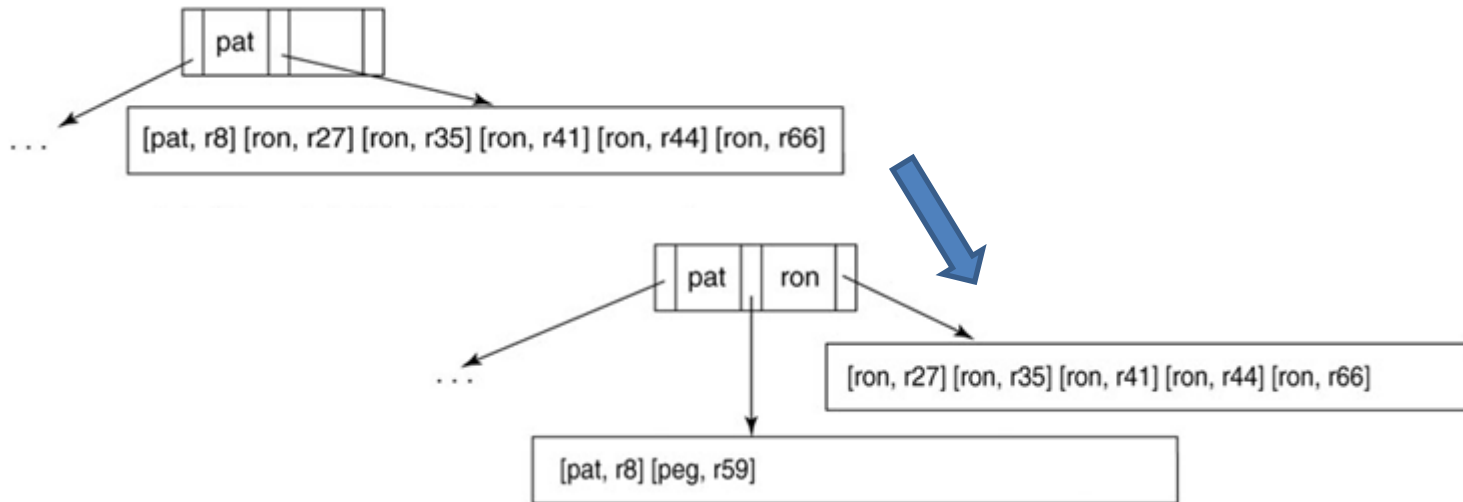
# Duplicate DataVals (1/2)

- When splitting a leaf block, we must place all records with same dataVal in same block



# Duplicate DataVals (2/2)

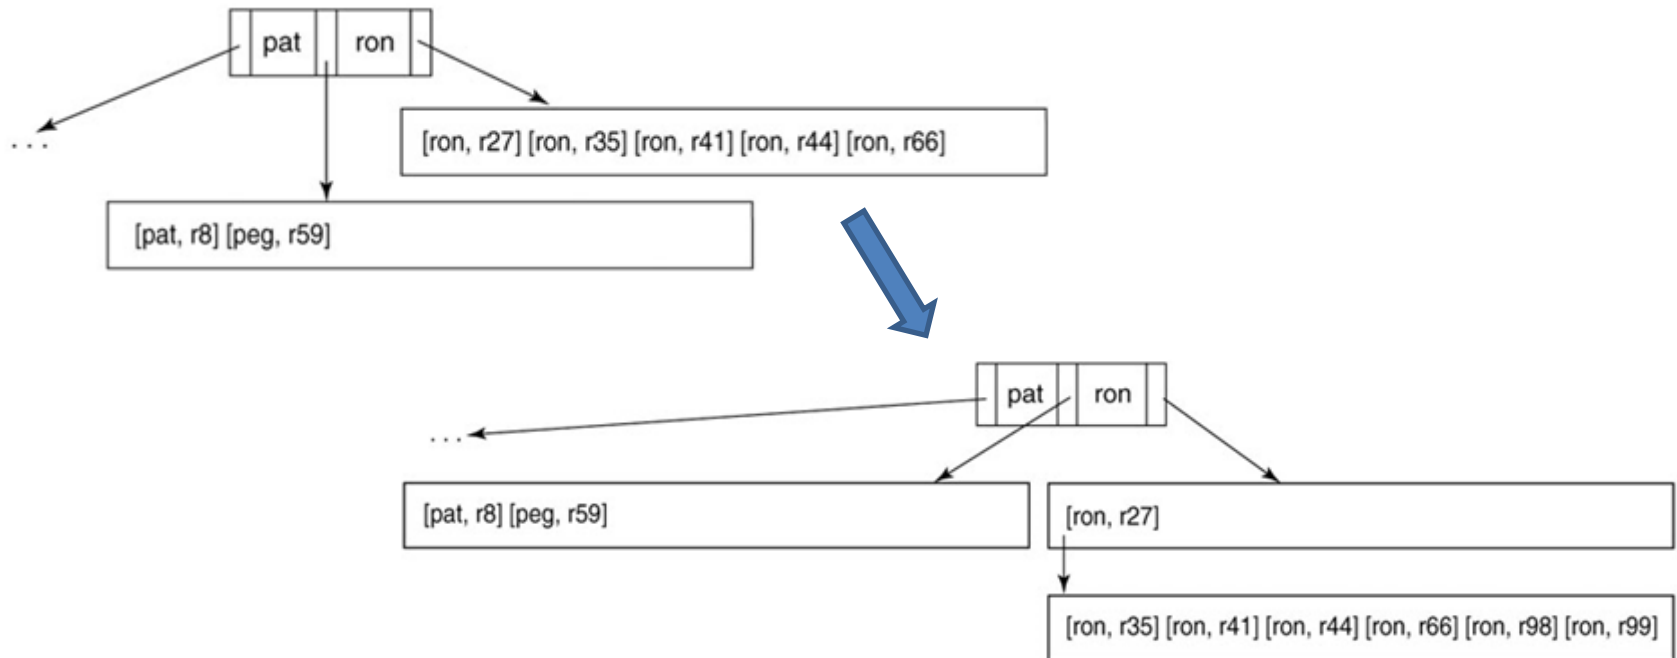
- E.g., insert [ron, r27]



- What if there are too many records with same dataVal?


# Overflow Blocks (1/2)

- Keep records of the same dataVal
- Chained by primary blocks



# Deletion


1. Search the index with the target dataVal
2. Delete the index record in a leaf block
3. Move the next records one-slot ahead
4. Merge blocks if #records is less than a threshold
5. Recursive delete on parents






# B-tree Index in VanillaCore

- Related package
  - `storage.index.btree`
- B-tree page
  - Directory pages



- Leaf pages



- Supports search, insert, but **not** delete

# Outline

- Overview
  - API in VanillaCore
- Hash-Based Indexes
- B-Tree Indexes
- **Query Processing**
- Transaction Management

# Related Relational Algebra

- Related package: `query.algebra.index`
- `IndexSelectPlan`
- `IndexJoinPlan`

# Update Planner

- Related package: `query.planner.index`
- `IndexUpdatePlanner`

# Outline

- Overview
  - API in VanillaCore
- Hash-Based Indexes
- B-Tree Indexes
- Query Processing
- **Transaction Management**

# Index Locking

- Why, given that we have S2PL already?
  - Can we just lock data objects (after index search)?
- No! You need to lock indices
- To ensure the consistency of the index structures
- To prevent phantom due to modification

# Maintaining Structure Consistency

- How?
- Naïve: simply s-/x-lock on an index
- But an index is one of the most frequently accessed meta-structures in a DBMS
- Can you improve the performance?
- Idea: early lock release

# Specialized Locking Protocols for Hash Index

- Search:
  - Repeat until rid found:
    - S-lock the current block of bucket
    - Release the S-lock of the previous block
    - Perform index lookup in current block
  - S-lock on data object
  - ***Release the S-lock of index block***
  - Perform data access insert/delete
  - Hold the data locks following S2PL



# Specialized Locking Protocols for Hash Index

- Search:
  - S-lock on the bucket file
  - Perform index lookup
  - S-lock on data object
  - ***Release the S-lock of bucket***
  - Perform data access
  - Hold the data lock following S2PL

# Specialized Locking Protocols

- Data access with For each static hash index:
  - S-/X-lock on the bucket file
  - Perform index lookup/insert/delete
  - ***Release the index locks***
  - S-/X-lock on data object
  - Perform data access insert/delete
  - Hold the data locks following S2PL

# Specialized Locking Protocols

- Data access with a B-tree index:
  - *Crab-locking* along the B-tree
  - Perform index lookup/insert/delete
  - *Release the leaf locks*
  - S-/X-lock on data object
  - Perform data access insert/delete
  - Hold the data locks following S2PL
- Deadlock free

# How about Phantom due to Updates?

- Idea: hold the lock of B-tree leave until tx end
- Limitation: only prevents phantoms due to single-table updates
- Be careful about deadlock!
  - This protocol is no longer deadlock free
  - A better deadlock handling is required

# Recovery

- Since locks are released early, logical logging and recovery is required

**You Have Assignment!**

# Assignment: Preventing Update Phantoms

- Modify index locking protocol to prevent phantoms due to updates
- Hint: revisit lock mode and data access path
  - No update phantom in SERIALIZED isolation mode
  - Other isolation modes need to be compatible with SERIALIZED mode

# Assignment: Preventing Update Phantoms

- Report
  - New table for lock mode & data access path
  - How you modify relate components (e.g., ConcurrencyMgr, Planner, etc.)
    - API changes and/or new classes
  - Show your results before and after your modification given a sample stored procedure
  - Compare the throughputs before and after your modification using the given benchmark & loader
- Due 2015/06/03 (Wed.) 23:59:59