


Query Optimization

Shan-Hung Wu and DataLab
CS, NTHU

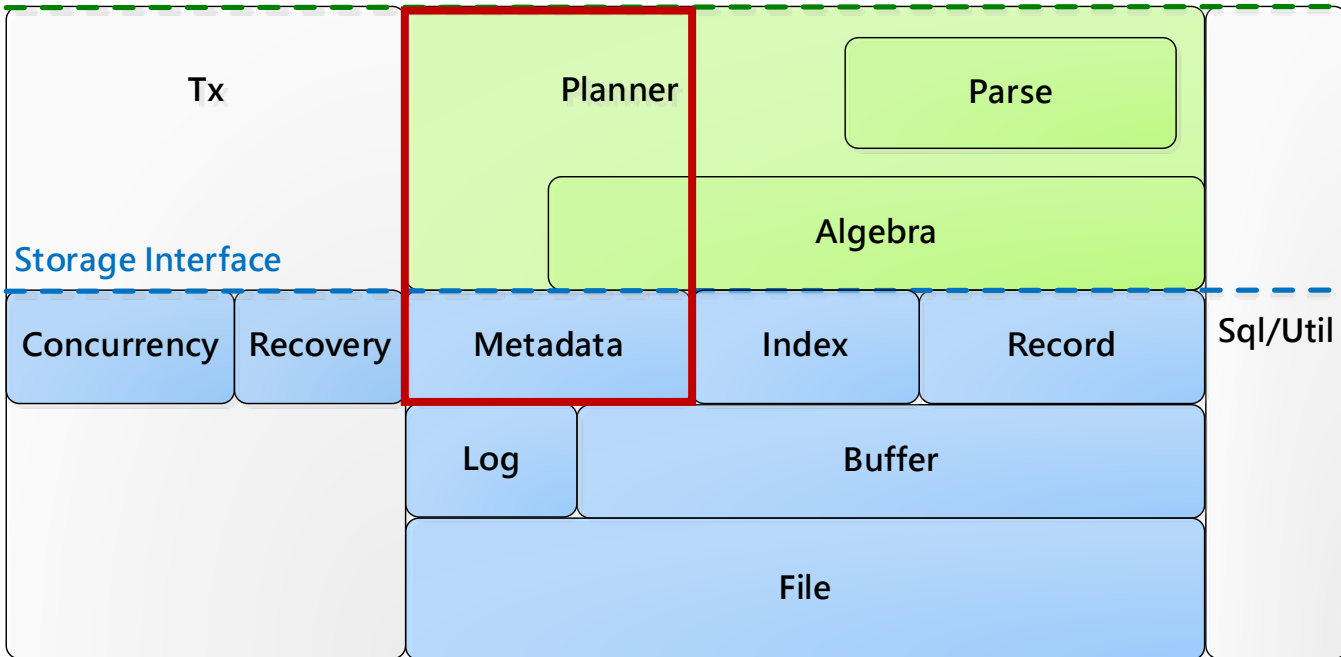
Where Are We?

VanillaCore

JDBC Interface (at Client Side)



Query Interface



Outline

- Overview
- Cost Estimation
 - Cardinality Estimation
 - Histogram-based Estimation
 - Types of Histograms
- Heuristic Query Optimizer
 - Basic Planner
 - Pushing Select Down
 - Join Ordering
 - Heuristic Query Planner in VanillaCore
- Selinger-Style Query Optimizer

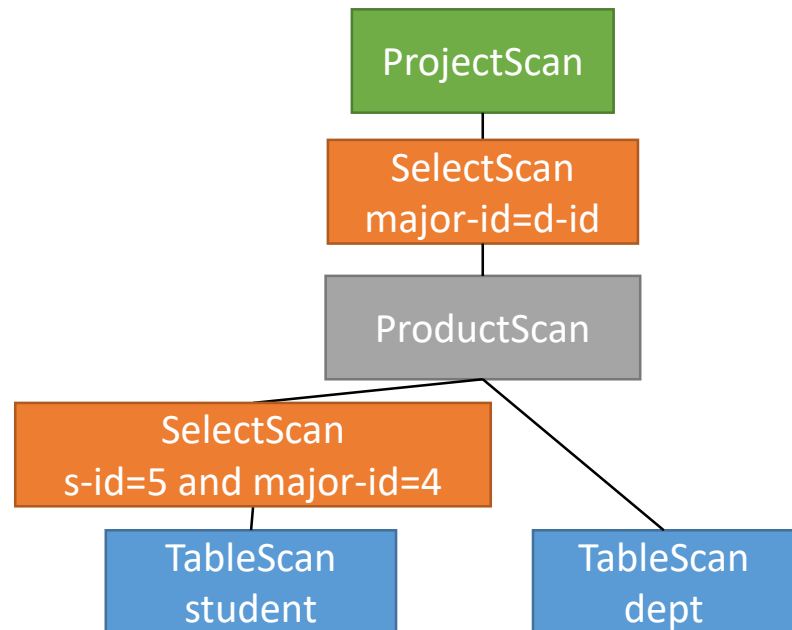
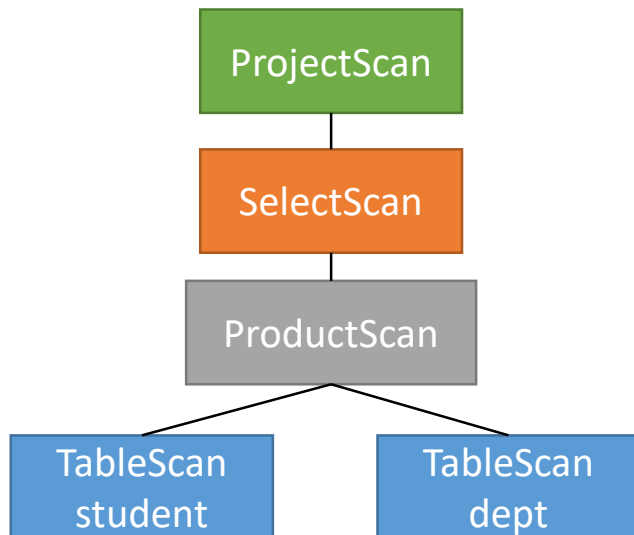
Outline

- Overview
- Cost Estimation
 - Cardinality Estimation
 - Histogram-based Estimation
 - Types of Histograms
- Heuristic Query Optimizer
 - Basic Planner
 - Pushing Select Down
 - Join Ordering
 - Heuristic Query Planner in VanillaCore
- Selinger-Style Query Optimizer

SQL and Relational Algebra

- A SQL command can be expressed as multiple trees in relational algebra

```
SELECT sname FROM student, dept
WHERE major-id = d-id AND s-id = 5 AND major-id = 4;
```



Query Optimization

- A good scan tree can be faster than a bad one for orders of magnitude
- Query optimizer:
 1. Generate candidate plan trees
 2. Estimate cost of each corresponding scan tree (not discussed yet)
 3. Pick and open the “best” one to execute query
- Goal (ideally): find the one with least cost
- Goal (in practice): avoid bad trees

Outline

- Overview
- Cost Estimation
 - Cardinality Estimation
 - Histogram-based Estimation
 - Types of Histograms
- Heuristic Query Optimizer
 - Basic Planner
 - Push Select Down
 - Join Order Problem
 - Heuristic Planner in VanillaCore
- Selinger-Style Query Optimizer

Metric for Cost

- Cost of a query?
 - To user: query delay
 - Low delay also implies better system throughput
-
- Typically, I/O delay dominates query delay

Cost Estimation

- For each plan/table p , we estimate $B(p)$
 - #blocks accessed by the corresponding scan
- Usually, estimating $B(p)$ requires more knowledge:
 - $R(p)$: #records output
 - *Search cost* (#blocks) of index, if used
 - $V(p,f)$: #distinct values for field f in p

Estimating $B(p)$

p	$B(p)$
TablePlan	Actual #blocks cached by StatMgr (via periodic table scanning)
ProjectPlan(c)	$B(c)$
SelectPlan(c)	$B(c)$
IndexSelectPlan(t)	$\text{IndexSearchCost}(R(t), R(p)) + R(p)$
ProductPlan(c1, c2)	$B(c1) + (R(c1) * B(c2))$
IndexJoinPlan(c1, t2)	$B(c1) + (R(c1) * \text{IndexSearchCost}(R(t2), 1)) + R(p)$

- $B(c)$ is evaluated recursively down to the table level

For Any p , We Need to Estimate $R(p)$ and Index Search Cost


- Index Search Cost:
 - `HashIndex.searchCost()`
 - `BTreeIndex.searchCost()`
- Estimating $R(p)$ is called *cardinality estimation*

Outline

- Overview
- Cost Estimation
 - Cardinality Estimation
 - Histogram-based Estimation
 - Types of Histograms
- Heuristic Query Optimizer
 - Basic Planner
 - Pushing Select Down
 - Join Ordering
 - Heuristic Query Planner in VanillaCore
- Selinger-Style Query Optimizer

Naïve Approach

- Uniform assumption
 - All values in field appear with the same probability



- Few statistics are enough:

R(c)	#records in child plan c
V(c, f)	#distinct values in field f in c
Max(c, f)	Max value in field f in c
Min(c, f)	Min value in field f in c

$$p = \text{Select}(c, f=x)$$

- $R(p)$?

$R(c)$	#records in child plan c
$V(c, f)$	#distinct values in field f in c
$\text{Max}(c, f)$	Max value in field f in c
$\text{Min}(c, f)$	Min value in field f in c

- $\text{Selectivity}(f=x): \frac{1}{V(c, f)}$
- $R(p): \text{Selectivity}(f=x) * R(c)$



$$p = \text{Select}(c, f > x)$$


- $R(p)$?

$R(c)$	#records in child plan c
$V(c, f)$	#distinct values in field f in c
$\text{Max}(c, f)$	Max value in field f in c
$\text{Min}(c, f)$	Min value in field f in c

- $\text{Selectivity}(f > x): \frac{\text{Max}(c, f) - x}{\text{Max}(c, f) - \text{Min}(c, f)}$
- $R(p): \text{Selectivity}(f > x) * R(c)$

• $R(p)$?

- $\text{Selectivity}(f > x): \frac{\text{Max}(c, f) - x}{\text{Max}(c, f) - \text{Min}(c, f)}$
- $R(p): \text{Selectivity}(f > x) * R(c)$




Outline

- Overview
- Cost Estimation
 - Cardinality Estimation
 - **Histogram-based Estimation**
 - Types of Histograms
- Heuristic Query Optimizer
 - Basic Planner
 - Pushing Select Down
 - Join Ordering
 - Heuristic Query Planner in VanillaCore
- Selinger-Style Query Optimizer


Naïve Estimation is Inaccurate

- In the real world, values in a field are seldom uniform distributed
- $p = \text{Select}(c, f=14)$
- Estimated $R(p) = \frac{1}{15} * R(c) = 3$
- Actually, $R(p) = 9$



Histogram


- Approximates value distribution in every field
- Partitions field values into a set of *buckets*



- More #buckets, more accurate approximation
 - Tradeoff between accurate and storage cost

Buckets

- Each bucket b collects statistics of a value range
 - Assumes ***uniform distribution of records and values*** in b



- $R(p, f, b)$: #records
- $V(p, f, b)$: #distinct values
- $\text{Range}(p, f, b)$: value range

Cardinality Estimation

- Not matter what p is, we have

$$R(p) = \sum_{b \in p.hist.buckets(f)} R(p, f, b)$$

for any f

- Problem: how to construct the histogram?

Range Selection (1/2)


- $p = \text{Select}(c, f \text{ in Range})$
- For each bucket b **in** f :
 - $\text{Selectivity} = \frac{|Range(c, f, b) \cap Range|}{|Range(c, f, b)|}$
 - $R(p, f, b) = R(c, f, b) * \text{selectivity}$
 - $V(p, f, b) = V(c, f, b) * \text{selectivity}$
 - $Range(p, f, b) = Range(c, f, b) \cap Range$
- Assumptions:
 - #Records in a bucket are uniformly distributed
 - Values in a bucket are uniformly distributed

Given $\forall f, b$:

$R(c, f, b)$

$V(c, f, b)$

$Range(c, f, b)$



Range Selection (2/2)

- $p = \text{Select}(c, f \text{ in Range})$
- For each bucket b in $f' \neq f$:
 - $\text{Reduction} = \frac{\sum_b R(p, f, b)}{R(c)}$
 - $R(p, f', b) = R(c, f', b) * \text{Reduction}$
 - $V(p, f', b) = \min(V(c, f', b), R(p, f', b))$
 - $\text{Range}(p, f', b) = \text{Range}(c, f', b)$

Given $\forall f, b$:

$R(c, f, b)$

$V(c, f, b)$

$\text{Range}(c, f, b)$

- Assumptions:
 - Values in different fields are independent with each other

Product

- $p = \text{Product}(c1, c2)$
- For each (b, f) in $c1$:
 - $R(p, f, b) = R(c1, f, b) * R(c2)$
 - $V(p, f, b) = V(c1, f, b)$
 - $\text{Range}(p, f, b) = \text{Range}(c1, f, b)$
- For each (b, f) in $c2$:
 - $R(p, f, b) = R(c2, f, b) * R(c1)$
 - $V(p, f, b) = V(c2, f, b)$
 - $\text{Range}(p, f, b) = \text{Range}(c2, f, b)$

Given $\forall f, b$:

$R(c1, f, b)$

$V(c1, f, b)$


$\text{Range}(c1, f, b)$

$R(c2, f, b)$


$V(c2, f, b)$

$\text{Range}(c2, f, b)$

Join Selection (1/2)



Match rate with recs in b2



- $p = \text{Select}(c, f=g)$
- For each bucket **b1** in f and **b2** in g:
 - If $\text{Range}(c, f, b1) \cap \text{Range}(c, f, b2) = \emptyset$ discard and b1 and b2
 - $\min V = \min(V(c, f, b1), V(c, f, b2))$
 - $R1 = R(c, f, b1) * \frac{\min V}{V(c, f, b1)} * \frac{1}{V(c, g, b2)} * \frac{R(c, g, b2)}{R(c)}$
 - $R2 = R(c, g, b2) * \frac{\min V}{V(c, g, b2)} * \frac{1}{V(c, f, b1)} * \frac{R(c, f, b1)}{R(c)}$
 - $R(p, f, b1) = R(p, g, b2) = \min(R1, R2)$
 - $V(p, f, b1) = V(p, g, b2) = \min(V(c, f, b1), V(c, g, b2))$
 - $\text{Range}(p, f, b1) = \text{Range}(p, g, b2) = \text{Range}(c, f, b1) \cap \text{Range}(c, f, b2)$
- Assumptions:
 - Values in bucket are uniformly distributed
 - All values in the range having smaller number of values appear in the range having larger number of values
 - Values in different fields are independent with each other

Match rate with recs not in b2

Join Selection (2/2)

- $p = \text{Select}(c, f=g)$
- For each bucket b in $f' \notin \{f, g\}$:
 - $\text{Reduction} = \frac{\sum_b R(p, f, b)}{R(c)}$
 - $R(p, f', b) = R(c, f', b) * \text{Reduction}$
 - $V(p, f', b) = \min(V(c, f', b), R(p, f', b))$
 - $\text{Range}(p, f', b) = \text{Range}(c, f', b)$
- Assumptions:
 - Values in different fields are independent with each other

Cost Estimation in VanillaCore


- `B(p): p.blocksAccessed()`
- Histogram-based cardinality estimation:
 - `R(p): p.histogram().recordsOutput()`
 - `V(p,f): p.histogram().distinctVaues(f)`
- Each plan builds its own histogram in constructor
- Important utility methods to trace:
 - `SelectPlan.constantRangeHistogram()`
 - `ProductPlan.productHistogram()`
 - `SelectPlan.joinFieldHistogram()`
 - `AbstractJointPlan.joinHistogram()`

Outline

- Overview
- Cost Estimation
 - Cardinality Estimation
 - Histogram-based Estimation
 - Types of Histograms
- Heuristic Query Optimizer
 - Basic Planner
 - Pushing Select Down
 - Join Ordering
 - Heuristic Query Planner in VanillaCore
- Selinger-Style Query Optimizer


Table Histogram at Lowest-Level

- Data structure that approximates value distribution
- Partitions field values into a set of **buckets**
- Each bucket b collects statistics of a value range
 - Assumes **uniform distribution of records and values** in b
- Given a fixed #buckets, how to decide bucket ranges?



Equi-Width Histogram


- Partition strategy: all buckets have the same range
- $|\text{Range}(b)| = \frac{\text{Max}(p,f) - \text{Min}(p,f) + 1}{\#Buckets}$



- Problem: some buckets may be wasted

Equi-Depth Histogram



- Partition strategy: all buckets have the same #recs
- $\text{Depth} = \frac{R(p)}{\#Buckets}$





- Problem: records/values in a bucket may **not** be uniformly distributed

Max-Diff Histogram

- Partition strategy: split buckets at values with max. diff in #rec ($\text{MaxDiff}(F)$) or area ($\text{MaxDiff}(A)$):
 1. #recs: uniform #records in each bucket



2. Area: uniform #records *and values* in each bucket



Histogram in VanillaCore

- Table histograms are statistics metadata
 - `org.vanilladb.core.storage.metadata.statistics`
- Accessed (by TablePlan) via `StatMgr.getTableStatInfo()`

Histogram
<pre>+ Histogram() + Histogram(fldnames : Set<String>) ~ Histogram(dists : Map<String, Collection<Bucket>>) + Histogram(hist : Histogram) + fields() : Set<String> + buckets(fldname : String) : Collection<Bucket> + addField(fldname : String) + addBucket(fldname : String, bkt : Bucket) + setBuckets(fldname : String, bkts : Collection<Bucket>) + recordsOutput() : double + distinctValues(fldname : String) : double + toString() : String + toString(int) : String</pre>

Bucket
<pre>+ Bucket(valrange : ConstantRange, freq : double, distvals : double) + Bucket(valrange : ConstantRange, freq : double, distvals : double, pcts : Percentiles) + valueRange() : ConstantRange + frequency() : double + distinctValues() : double + distinctValues(range : ConstantRange) : double + valuePercentiles() : Percentiles + toString() : String + toString(int) : String</pre>

Building Histogram (1/2)

- When system starts up:
- StatMgr:
 - Scans table and calls `SampledHistogramBuilder.sample()`
 - When done, calls `SampledHistogramBuilder.newMaxDiffHistogram()`
- Histogram types:
 - `MaxDiff(A)` : when field value is numeric
 - `MaxDiff(F)` : otherwise

Building Histogram (2/2)

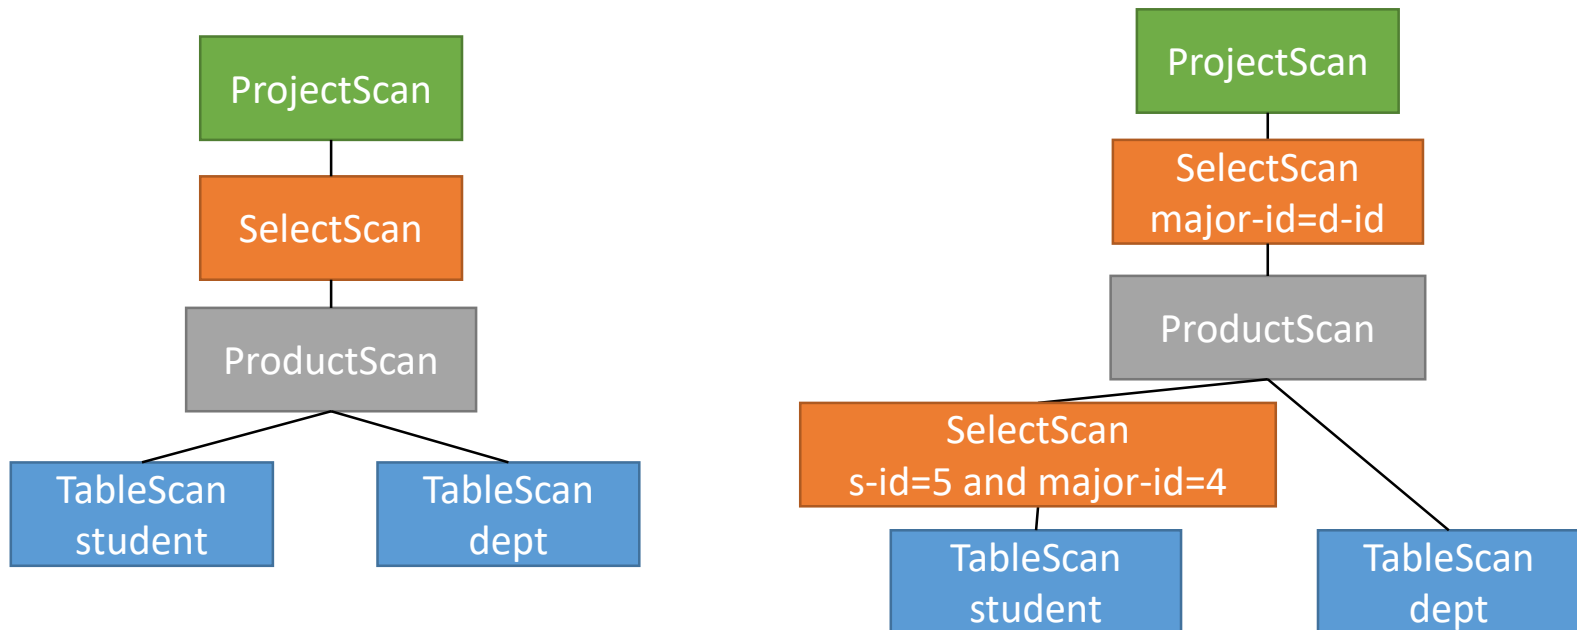
- At runtime:
- StatMgr tacks #recs updated for each table
 - QueryPlanner calls StatMgr.countRecordUpdates() after executing modify/insert/delete queries
- Rebuilds histogram *in background* when StatMgr.getTableStatInfo() is called
 - If #recs updated > threshold (e.g., 100)
- StatisticsRefreshTask:
 - Scans table and calls SampledHistogramBuilder.sample()
 - When done, calls SampledHistogramBuilder.newMaxDiffHistogram()

Outline

- Overview
- Cost Estimation
 - Cardinality Estimation
 - Histogram-based Estimation
 - Types of Histograms
- **Heuristic Query Optimizer**
 - Basic Planner
 - Pushing Select Down
 - Join Ordering
 - Heuristic Query Planner in VanillaCore
- Selinger-Style Query Optimizer

Query Optimization

- Query optimizer:
 1. Generate candidate plan trees
 2. Estimate cost of each corresponding scan tree
 3. Pick and open the “best” one to execute query

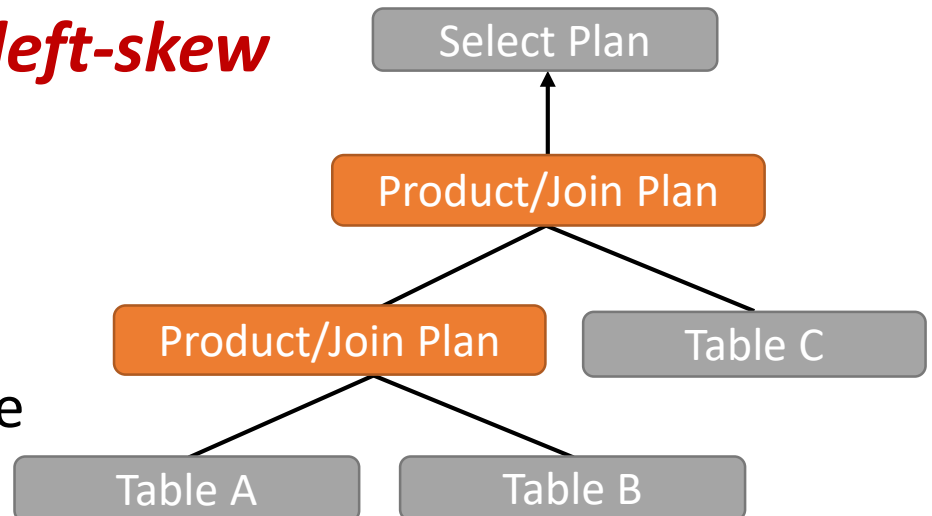


In Reality...

- Generating all candidate plan trees are too costly
 - #trees with n products/joins = Catalan number:

$$\frac{1}{n+1} \binom{2n}{n}$$

- Compromise: consider *left-skew* candidate trees only
- Query planner's goal
 - Avoiding bad trees
 - Not finding the best tree



Why Left-Skew Trees Only?

- Tend to be better than plans of other shapes
- Because many join algorithms scan right child c2 multiple times
- Normally, we don't want c2 to be a complex subtree


BasicQueryPlanner

```
public Plan createPlan(QueryData data, Transaction tx) {
    // Step 1: Create a plan for each mentioned table or view
    List<Plan> plans = new ArrayList<Plan>();
    for (String tblname : data.tables()) {
        String viewdef = VanillaDb.catalogMgr().getViewDef(tblname, tx);
        if (viewdef != null)
            plans.add(VanillaDb.newPlanner().createQueryPlan(viewdef, tx));
        else
            plans.add(new TablePlan(tblname, tx));
    }
    // Step 2: Create the product of all table plans
    Plan p = plans.remove(0);
    for (Plan nextplan : plans)
        p = new ProductPlan(p, nextplan);
    // Step 3: Add a selection plan for the predicate
    p = new SelectPlan(p, data.pred());
    // Step 4: Add a group-by plan if specified
    if (data.groupFields() != null) {
        p = new GroupByPlan(p, data.groupFields(), data.aggregationFn(), tx);
    }
    // Step 5: Project onto the specified fields
    p = new ProjectPlan(p, data.projectFields());
    // Step 6: Add a sort plan if specified
    if (data.sortFields() != null)
        p = new SortPlan(p, data.sortFields(), data.sortDirections(), tx);
    // Step 7: Add a explain plan if the query is explain statement
    if (data.isExplain())
        p = new ExplainPlan(p);
    return p;
}
```

- Product/join order follows what's written in SQL

Cost & Bettlenecks

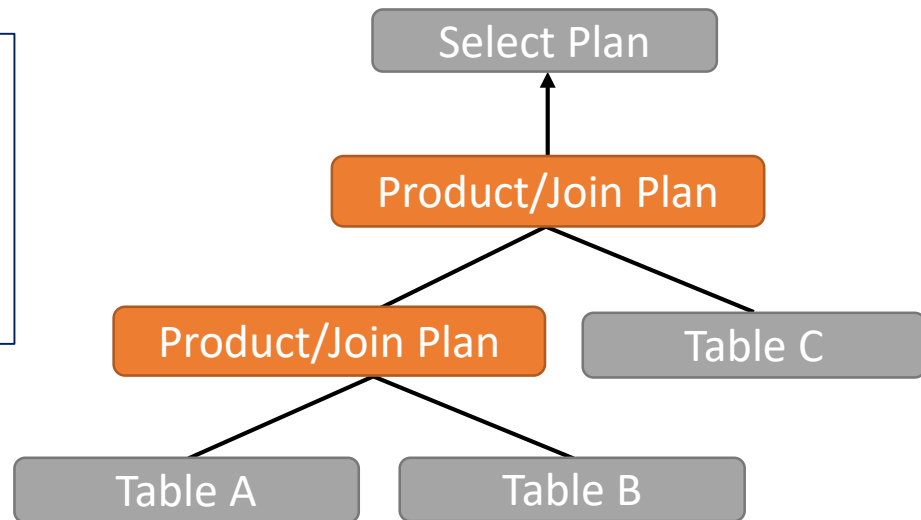
```
SELECT  A.c1, B.c2, C.c3
FROM    A, B, C
WHERE   A.aid = C.aid
AND     B.bid = C.bid
AND     A.c2 = xxx
```



- B(root) dominated by **#recs** of product/join ops
 - $B(\text{Product}(c1, c2)) = B(c1) + (R(c1) * B(c2))$
 - $B(\text{IndexJoin}(c1, c2)) = B(c1) + (R(c1) * \text{SearchCost}(\dots)) + \dots$
 - $B(\text{Select}(c)) = B(c)$

Optimizations

```
SELECT  A.c1, B.c2, C.c3
FROM    A, B, C
WHERE   A.aid = C.aid
AND     B.bid = C.bid
AND     A.c2 = xxx
```

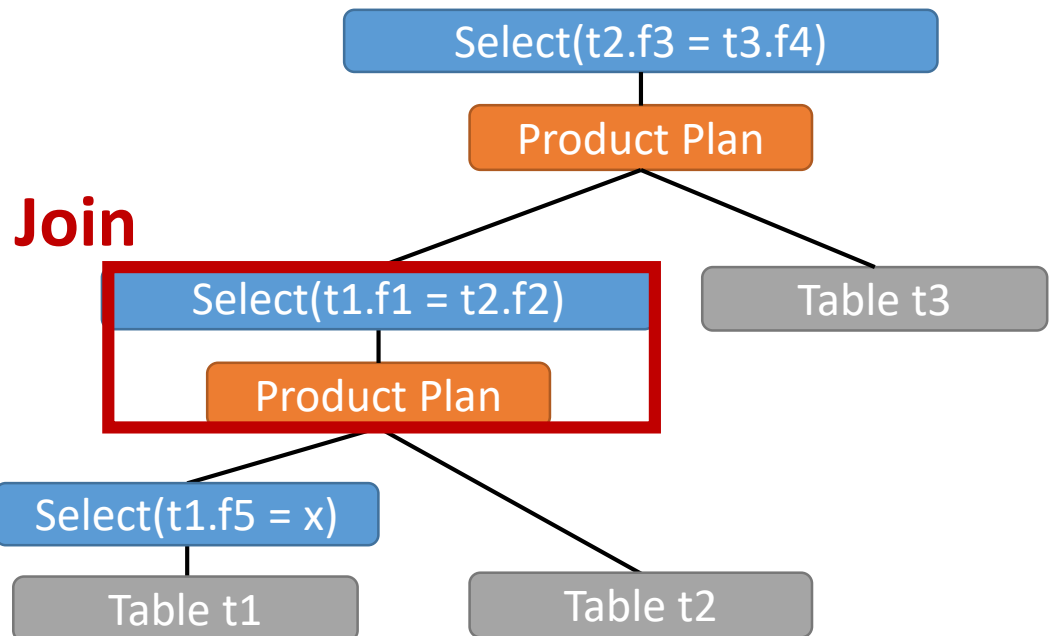


- Goal $\downarrow B(\text{root})$ reduced to $\downarrow R(c1)$ and $\downarrow R(c2)$
- Heuristics:
 - Pushing Select ops down
 - Greedy Join ordering

Pushing Select Ops Down


- Execute Select ops as early as possible
- $\downarrow R(c1)$ and $\downarrow R(c2)$ of each product/join op

SELECT	*
FROM	t1, t2, t3
WHERE	t1.f1 = t2.f2
AND	t2.f3 = t3.f4
AND	t1.f5 = x



Greedy Join Ordering

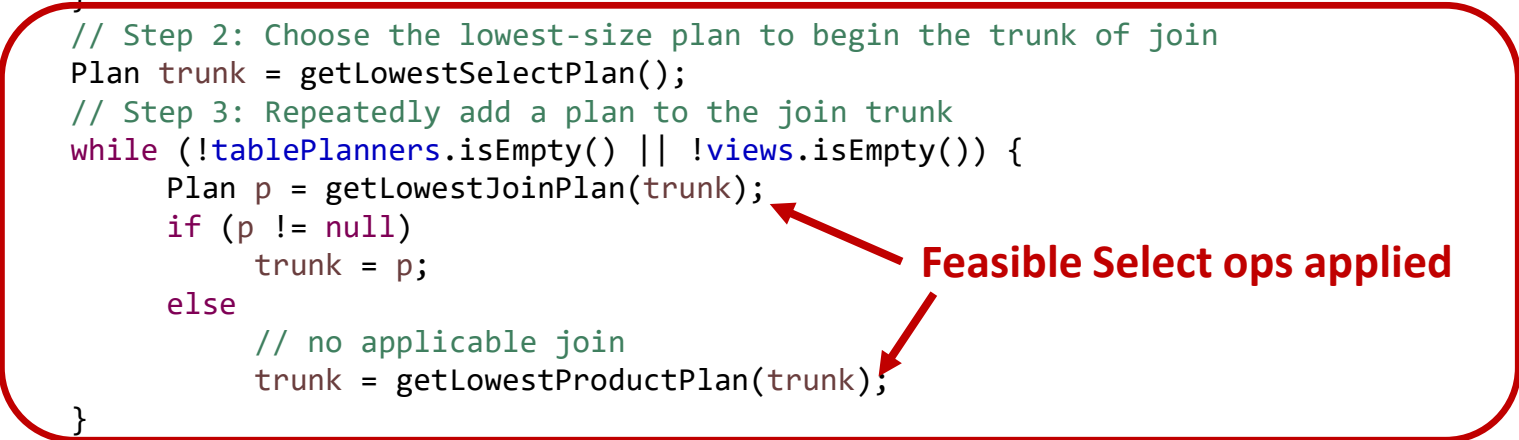
- $B(\text{root}) = B(p1) + (R(p1) * \dots) + \dots$
 - $\downarrow B(\text{root})$ implies $\downarrow(p1)$
- $B(p1) = B(c1) + (R(c1) * \dots) + \dots$
 - $\downarrow B(\text{root})$ also implies $\downarrow(c1)$
- ...
- $B(\text{root}) \propto R(p1) + R(c1) + \dots$



- **Greedy Join ordering**: repeatedly add table to the “trunk” that result in lowest $R(\text{trunk})$

HeuristicPlanner in VanillaCore

```
public Plan createPlan(QueryData data, Transaction tx) {  
    // Step 1: Create a TablePlanner object for each mentioned table/view  
    int id = 0;  
    for (String tbl : data.tables()) {  
        String viewdef = VanillaDb.catalogMgr().getViewDef(tbl, tx);  
        if (viewdef != null)  
            views.add(VanillaDb.newPlanner().createQueryPlan(viewdef, tx));  
        else {  
            TablePlanner tp = new TablePlanner(tbl, data.pred(), tx, id);  
            tablePlanners.add(tp);  
        }  
        id += 1;  
    }  
    // Step 2: Choose the lowest-size plan to begin the trunk of join  
    Plan trunk = getLowestSelectPlan();  
    // Step 3: Repeatedly add a plan to the join trunk  
    while (!tablePlanners.isEmpty() || !views.isEmpty()) {  
        Plan p = getLowestJoinPlan(trunk);  
        if (p != null)  
            trunk = p;  
        else  
            // no applicable join  
            trunk = getLowestProductPlan(trunk);  
    }  
    // Step 4: Add a group by plan if specified  
    // Step 5: Project on the field names  
    // Step 6: Add a sort plan if specified  
    // Step 7: Add a explain plan if the query is explain statement  
}
```

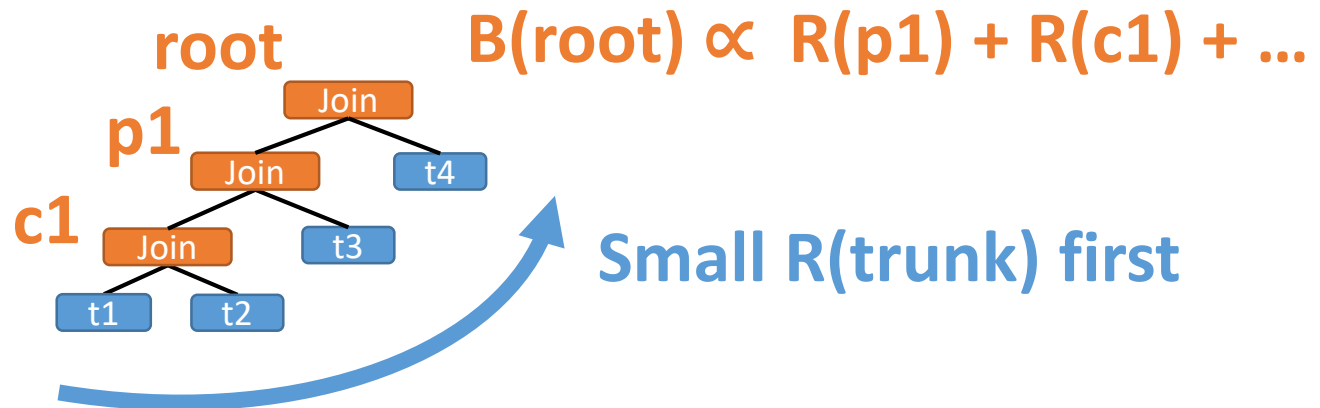


Feasible Select ops applied

Outline

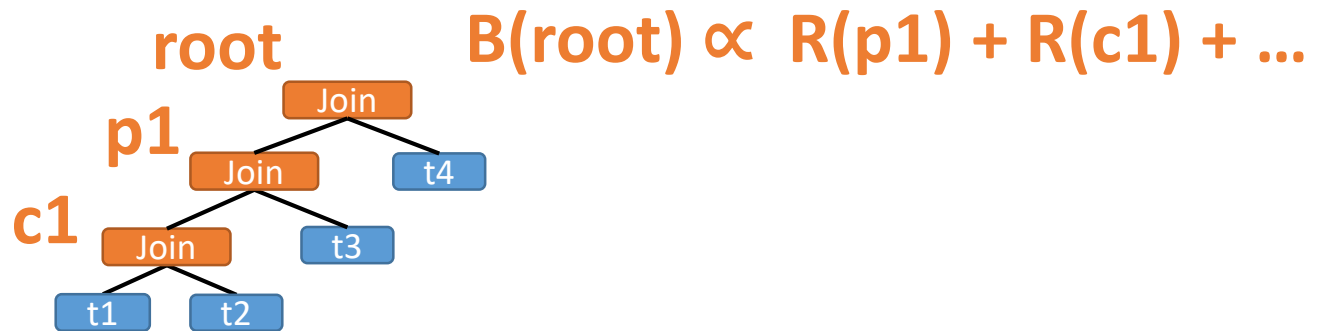
- Overview
- Cost Estimation
 - Cardinality Estimation
 - Histogram-based Estimation
 - Types of Histograms
- Heuristic Query Optimizer
 - Basic Planner
 - Pushing Select Down
 - Join Ordering
 - Heuristic Query Planner in VanillaCore
- Selinger-Style Query Optimizer

Why not HeuristicPlanner?



- Assumption: $\downarrow R(\text{c1})$ implies $\downarrow R(\text{p1})$
- May **not** be true: match rate matters
- Exhaustively searching the best join order?
 - Cost: **$O(n!)$** for n joins (e.g., $8! = 40320$)

Selinger-Style Optimizer



- Consider the best trees after 1, 2, 3, ... joins
- Observation: if $R(\mathbf{t3} \bowtie \mathbf{t1} \bowtie \mathbf{t2}) \leq R(\mathbf{t2} \bowtie \mathbf{t3} \bowtie \mathbf{t1})$, then $R(\mathbf{t3} \bowtie \mathbf{t1} \bowtie \mathbf{t2} \bowtie \mathbf{t4}) \leq R(\mathbf{t2} \bowtie \mathbf{t3} \bowtie \mathbf{t1} \bowtie \mathbf{t4})$
- We can use *dynamic programming* to avoid repeating computations

Selinger Optimizer Example (1/6)

- Here are 3 relations to join: A, B, C
- Step 1: compute the cost (R) of each relation's cheapest plan

1-Set	Best Plan	R
{A}	Index Select Plan	10
{B}	Table Plan	30
{C}	Select Plan	20

Selinger Optimizer Example (2/6)

- Step 2: compute the cost of 2-way joins reusing 1-way cost just cached
- $R(A \bowtie B) = R(B \bowtie A)$, so we just keep one

1-Set	Best Plan	R
{A}	Index Select Plan	10
{B}	Table Plan	30
{C}	Select Plan	20

2-Set	Best Plan	R
{A,B}	$A \bowtie B$	159
{A,C}	$A \bowtie C$	98
{B,C}	$B \bowtie C$	77

Selinger Optimizer Example (3/6)

- Here are 3 relations to join – A, B, C
- Step 2
 - Compute the cost of 2-way join by estimating all permutation using the single relation cost just cached
 - Ex. {A, B} =
 - Compare {A}B Cost: 159
 - Compare {B}A Cost: 189



Because the $R(AB)$, $R(BA)$ is the same, we can only keep the least cost one

Sub plan	Best Plan	Cost
{A}	Index Select Plan	10
{B}	Table Plan	30
{C}	Select Plan	20

Sub plan	Best Plan	Cost
{A, B}	$A \bowtie B$	159

Selinger Optimizer Example (4/6)

- Here are 3 relations to join – A, B, C
- Step 2
 - Compute the cost of 2-way join by estimating all permutation using the single relation cost just cached
 - Ex. {A, B} =
 - Compare {A}B Cost: 159
 - Compare {B}A Cost: 189

Sub plan	Best Plan	Cost
{A}	Index Select Plan	10
{B}	Table Plan	30
{C}	Select Plan	20

Sub plan	Best Plan	Cost
{A, B}	$A \bowtie B$	159
{A, C}	$C \bowtie A$	98
{B, C}	$C \bowtie B$	77

Selinger Optimizer Example (5/6)

- Here are 3 relations to join – A, B, C
- Step 3
 - Compute the cost of 3-way join by estimating all left-deep tree permutation using the step2's record
 - Ex. {A, B, C} =
 - Compare ({A, B})C Cost: 259
 - Compare ({B, C})A Cost: 111
 - Compare ({C, A})B Cost: 100



Sub plan	Best Plan	Cost
{A, B}	A ⋈ B	159
{A, C}	C ⋈ A	98
{B, C}	C ⋈ B	77

Sub plan	Best Plan	Cost

Selinger Optimizer Example (6/6)

- Here are 3 relations to join – A, B, C
- Step 3
 - Compute the cost of 3-way join by estimating all left-deep tree permutation using the step2's record
 - Ex. {A, B, C} =
 - Compare ({A, B})C Cost: 259
 - Compare ({B, C})A Cost: 111
 - Compare ({C, A})B Cost: 100

Sub plan	Best Plan	Cost
{A, B}	$A \bowtie B$	159
{A, C}	$C \bowtie A$	98
{B, C}	$C \bowtie B$	77

Sub plan	Best Plan	Cost
{A, B, C}	$C \bowtie A \bowtie B$	100

```

private Plan getAllCombination(Plan viewTrunk) {
    long finalKey = 0;

    // for layer = 1, use select down strategy to construct
    for (TablePlanner tp: tablePlanners) {
        Plan bestPlan = null;
        if (viewTrunk != null) {
            bestPlan = tp.makeJoinPlan(viewTrunk);
            if (bestPlan == null)
                bestPlan = tp.makeProductPlan(viewTrunk);
        }
        else
            bestPlan = tp.makeSelectPlan();

        AccessPath ap = new AccessPath(tp, bestPlan);
        lookupTbl.put(ap.getAPId(), ap);

        // compute final access path id
        finalKey += ap.getAPId();
    }

    .
    .
    .

}

```

```

// construct all combination layer by layer
for (int layer = 2; layer <= tablePlanners.size(); layer++) {
    Set<Long> keySet = new HashSet<Long>(lookupTbl.keySet());

    for (TablePlanner rightOne: tablePlanners) {
        for (Long key: keySet) {
            AccessPath leftTrunk = lookupTbl.get(key);

            // cannot join with table which (layer-1) combination already included
            if (leftTrunk.isUsed(rightOne.getId()))
                continue;

            // do join
            Plan bestPlan = rightOne.makeJoinPlan(leftTrunk.getPlan());
            if (bestPlan == null)
                bestPlan = rightOne.makeProductPlan(leftTrunk.getPlan());

            AccessPath candidate = new AccessPath(leftTrunk, rightOne, bestPlan);
            AccessPath ap = lookupTbl.get(candidate.getAPId());

            // there is no access path contains this combination
            if (ap == null) {
                lookupTbl.put(candidate.getAPId(), candidate);
            }
            // check whether new access path is better than previous
            else {
                if (candidate.getCost() < ap.getCost())
                    lookupTbl.put(candidate.getAPId(), candidate);
            }
        }
    }

    // remove the elements belong to layer-1
    // because in the next layer we only need this layer's combination
    for (Long key: keySet)
        lookupTbl.remove(key);
}

return lookupTbl.get(finalKey).getPlan();

```

- Iterate all table planners to join with all existing (layer-1) combination to construct this layer

```
public class AccessPath {
    private Plan p;
    private AccessPathId apId;
    private long cost = 0;
    private ArrayList<Integer> tblUsed = new ArrayList<Integer>();
```

```
public class AccessPathId {
    long id;

    AccessPathId(TablePlanner tp) {
        this.id = (long) Math.pow(2, tp.getId());
    }

    AccessPathId(AccessPath ap, TablePlanner tp) {
        this.id = ap.getAPId() + (long) Math.pow(2, tp.getId());
    }
    public long getID() {
        return id;
    }
}
```

- Using **sum of $\text{pow}(2, \text{tp.id})$** to represent the combination of tables in this access path
- Using **$\text{pow}(2, \text{tp.id})$** to avoid problems with different combinations but with the same apID
- Then we can use apID as the key of the lookup table

```
public AccessPath (TablePlanner newTp, Plan p) {
    this.p = p;
    this.tblUsed.add(newTp.getId());
    this.apId = new AccessPathId(newTp);
    this.cost = p.recordsOutput();
}

public AccessPath (AccessPath preAp, TablePlanner newTp, Plan p) {
    this.p = p;
    this.tblUsed.addAll(preAp.getTblUsed());
    this.tblUsed.add(newTp.getId());
    this.apId = new AccessPathId(preAp, newTp);

    // approximate cost = previous cost + new cost
    this.cost = preAp.getCost() + p.recordsOutput();
}
}
```



```

public class AccessPath {
    private Plan p;
    private AccessPathId apId;
    private long cost = 0;
    private ArrayList<Integer> tblUsed = new ArrayList<Integer>();

    public class AccessPathId {
        long id;

        AccessPathId(TablePlanner tp) {
            this.id = (long) Math.pow(2, tp.getId());
        }

        AccessPathId(AccessPath ap, TablePlanner tp) {
            this.id = ap.getAPId() + (long) Math.pow(2, tp.getId());
        }

        public long getID() {
            return id;
        }
    }
}

public AccessPath (TablePlanner newTp, Plan p) {
    this.p = p;
    this.tblUsed.add(newTp.getId());
    this.apId = new AccessPathId(newTp);
    this.cost = p.recordsOutput();
}

public AccessPath (AccessPath preAp, TablePlanner newTp, Plan p) {
    this.p = p;
    this.tblUsed.addAll(preAp.getTblUsed());
    this.tblUsed.add(newTp.getId());
    this.apId = new AccessPathId(preAp, newTp);

    // approximate cost = previous cost + new cost
    this.cost = preAp.getCost() + p.recordsOutput();
}
}

```

- Using **sum of pow(2, tp.id)** to represent the combination of tables in this access path
- Using **pow(2, tp.id)** to avoid problems with different combinations but with the same apID
- Then we can use apID as the key of the lookup table
- Approximate B(root) using $R(p1) + R(c1) \dots$

Reference

- <https://db.inf.uni-tuebingen.de/staticfiles/teaching/ws1011/db2/db2-selectivity.pdf>
- <https://www.cise.ufl.edu/~adobra/approxqp/histograms2>
- <https://pdfs.semanticscholar.org/b024/0a44105fa0a0967d96d109aac9f021902ebb.pdf>